



PostgreSQL

---

**Database Design and Management**

**Surendra Panpaliya**

# Agenda

Database Design Principles

Normalization and denormalization

Entity-Relationship (ER) modeling

# Agenda

01

**Creating and  
Managing  
Tables**

02

**Creating  
tables with  
SQL**

03

**Altering and  
dropping  
tables**

04

**Indexing for  
performance**

# Database Design Principles

Surendra Panpaliya

# Database Design Principles



Effective database design



Crucial for creating a database



Robust



Scalable



Efficient

# Understand the Requirements



Data to be stored



Relationships between different data entities

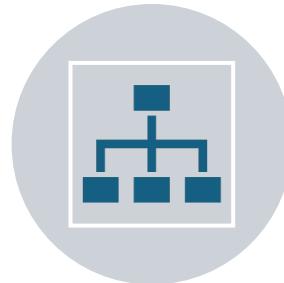


Queries that will be run

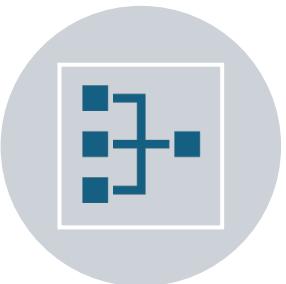
# Data Modeling



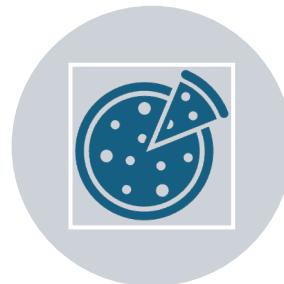
Use data modeling  
techniques



to represent the data  
structure.



Entity-Relationship (ER)  
modeling



is a popular method.

# Normalization

Organize data to minimize redundancy

dependency by dividing a database into

two or more tables

defining relationships between the tables.

# Denormalization

Necessary for performance optimization

Especially for read-heavy applications.

Involves combining tables

To reduce the number of joins

# Consistency and Integrity



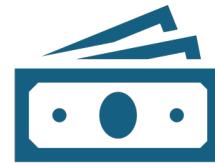
Ensure data  
integrity



through



Constraints



Transactions

# Consistency and Integrity



Use Primary  
keys,



Foreign keys



Unique  
constraints



To enforce data  
integrity

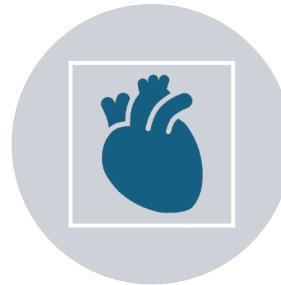
# Scalability



Design the database



vertically (increasing resources)



to scale horizontally (sharding) or



based on anticipated growth.

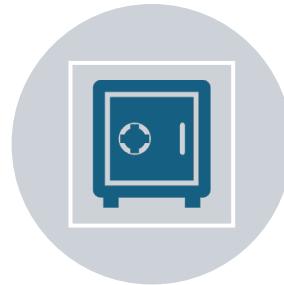
# Security



Incorporate security measures



including Access controls,



To protect sensitive data,



Encryption, and Auditing

# Performance



OPTIMIZE DATABASE  
PERFORMANCE



WITH INDEXING



QUERY  
OPTIMIZATION



CACHING  
STRATEGIES

# Database Design Principles for Healthcare



Designing a robust  
and efficient  
database



for the healthcare  
domain requires



careful  
consideration of  
several principles



to ensure data  
integrity,



security, and

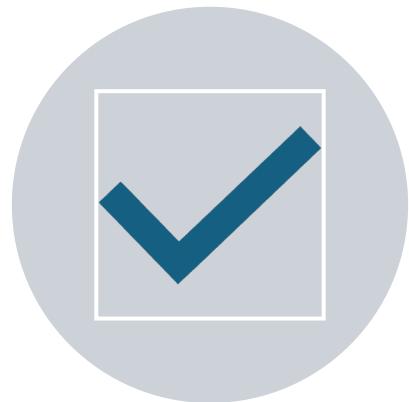


performance.

# Understand the Requirements



STAKEHOLDER  
ANALYSIS

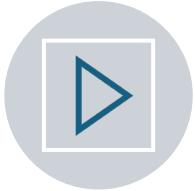


DATA REQUIREMENTS

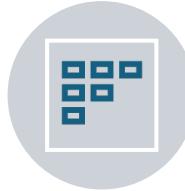


REGULATORY  
COMPLIANCE

# Stakeholder Analysis



Identify and  
Understand



Needs of all  
stakeholders



doctors, nurses



administrative  
staff



IT personnel

# Data Requirements



Determine the types of data



to be stored.



patient records



appointment schedules



medical histories,  
prescriptions



billing information

# Regulatory Compliance



Ensure compliance  
with



Healthcare regulations.



HIPAA (Health  
Insurance Portability  
and Accountability Act)



for data privacy and  
security

# Data Modeling

Entities and  
Relationships

Normalization

Denormalization

# Entities and Relationships

Identify key entities

Patients,  
Doctors,

Appointments,  
MedicalRecords,

Prescriptions,  
Billing

Define their  
relationships.

# Normalization



Apply normalization  
rules



Up to the 3rd Normal  
Form.



To eliminate data  
redundancy



Ensure data integrity

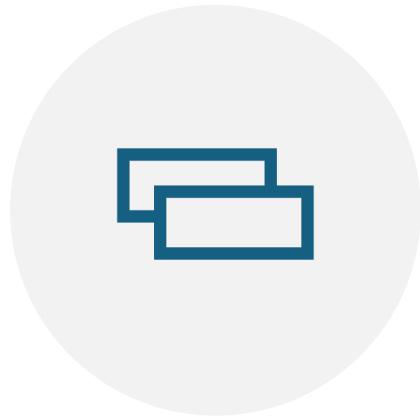
# Denormalization

Denormalize  
necessary

To improve  
performance

For read-  
heavy  
operations

# Schema Design



**TABLES AND  
COLUMNS**



**PRIMARY KEYS**



**FOREIGN KEYS**

# Tables and Columns

Create tables for each entity

with appropriate columns and data types

Use meaningful column names

Document their purpose

# Primary Keys

Define primary keys

for all tables

to uniquely identify

Records

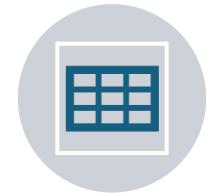
# Foreign Keys



Use foreign  
keys



to establish  
relationships



between tables,



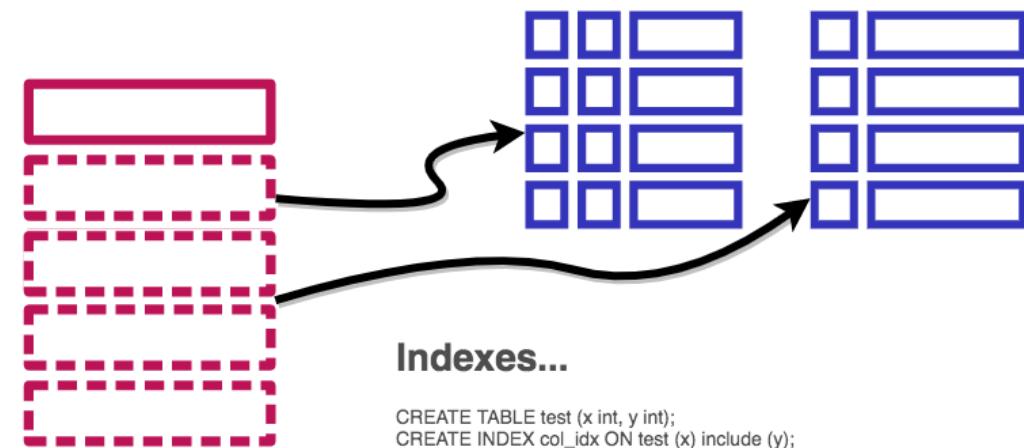
ensuring



referential  
integrity

# Constraints and Indexing

Constraints  
Indexes



# Constraints

Apply appropriate constraints

UNIQUE,  
CHECK, NOT  
NULL

to enforce  
data validity

# Indexes



Create indexes on



frequently queried  
columns



to enhance  
performance

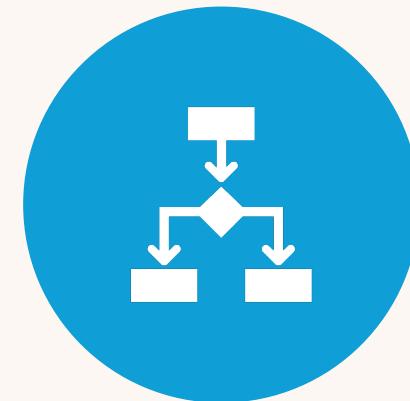
# Indexes



USE COMPOSITE  
INDEXES

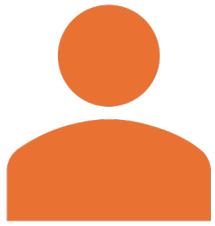


FOR QUERIES  
INVOLVING



MULTIPLE COLUMNS.

# Security and Access Control



Roles and Permissions



Encryption



Audit Logs

# Roles and Permissions

Define roles and

grant appropriate permissions

to ensure users have access

only to the data they need

# Encryption



Use encryption



for sensitive data



both at rest



in transit.

# Audit Logs



Implement



audit logging



to track



data access



modifications.

# Data Integrity and Consistency

---



**TRANSACTIONS**



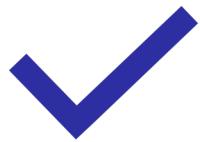
**TRIGGERS AND STORED  
PROCEDURES**

# Transactions

---



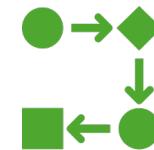
Use  
transactions



to ensure



atomicity of  
operations



for multi-step  
processes

# Triggers and Stored Procedures

---



IMPLEMENT  
TRIGGERS



STORED  
PROCEDURES



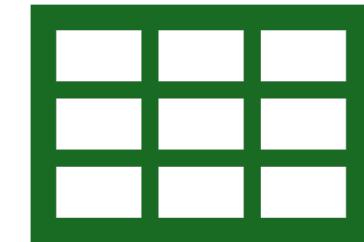
FOR COMPLEX  
VALIDATION



BUSINESS LOGIC  
ENFORCEMENT

# Performance Optimization

---

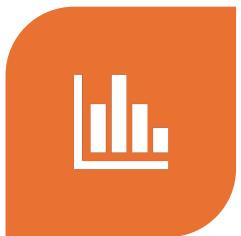


**Query Optimization**

**Partitioning**

# Query Optimization

---



ANALYZE AND



OPTIMIZE  
QUERIES



FOR BETTER  
PERFORMANCE



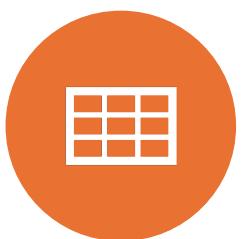
USING EXPLAIN



ANALYZE.

# Partitioning

---



USE TABLE  
PARTITIONING



FOR LARGE  
DATASETS



TO IMPROVE



QUERY  
PERFORMANCE



MAINTENANCE

# Scalability and Availability

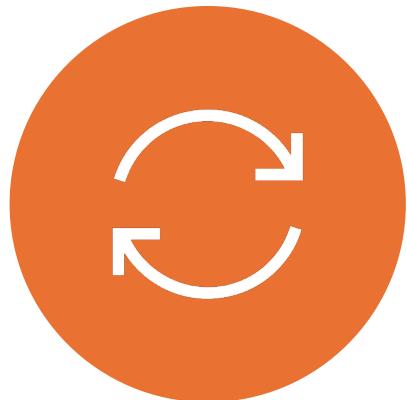
---

Replication

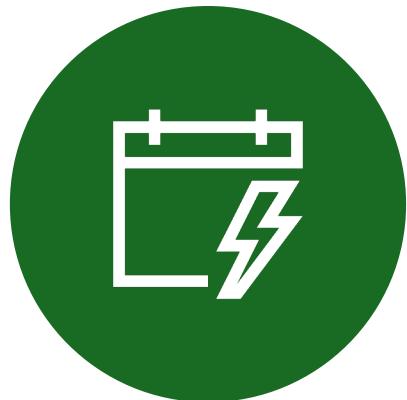
Backup and Restore

# Replication

---



SET UP REPLICATION



FOR HIGH  
AVAILABILITY



DISASTER RECOVERY.

# Backup and Restore

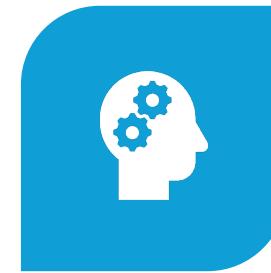
---



IMPLEMENT



REGULAR BACKUP  
STRATEGIES



ENSURE THE  
ABILITY



TO RESTORE DATA  
QUICKLY

# Compliance and Auditing

---



**REGULATORY  
COMPLIANCE**



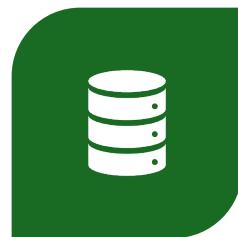
**AUDIT TRAILS**

# Regulatory Compliance

---



ENSURE



THE DATABASE  
DESIGN



OPERATIONS  
COMPLY



WITH RELEVANT  
HEALTHCARE



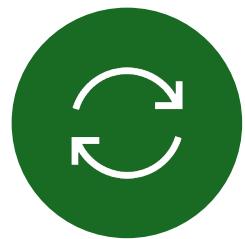
REGULATIONS

# Audit Trails

---



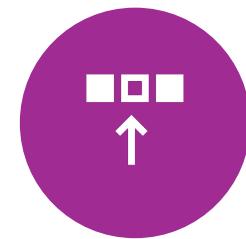
MAINTAIN  
AUDIT TRAILS



FOR CRITICAL  
DATA CHANGES



TO SUPPORT  
REGULATORY



COMPLIANCE



FORENSIC  
ANALYSIS

# Example PostgreSQL Schema for Healthcare Domain

Surendra Panpaliya

# Patients table

```
CREATE TABLE Patients (
    patient_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50) NOT NULL,
    date_of_birth DATE NOT NULL,
    gender VARCHAR(10),
    contact_number VARCHAR(15),
    email VARCHAR(100) UNIQUE,
    address TEXT,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

# Doctors table

```
CREATE TABLE Doctors (
    doctor_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50) NOT NULL,
    specialty VARCHAR(100),
    contact_number VARCHAR(15),
    email VARCHAR(100) UNIQUE,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

# Appointments table

```
CREATE TABLE Appointments (
    appointment_id SERIAL PRIMARY KEY,
    patient_id INT NOT NULL,
    doctor_id INT NOT NULL,
    appointment_date TIMESTAMP NOT NULL,
    status VARCHAR(20) CHECK (status IN ('Scheduled', 'Completed', 'Cancelled')),
    notes TEXT,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    FOREIGN KEY (patient_id) REFERENCES Patients(patient_id),
    FOREIGN KEY (doctor_id) REFERENCES Doctors(doctor_id)
);
```

# MedicalRecords table

```
CREATE TABLE MedicalRecords (
    record_id SERIAL PRIMARY KEY,
    patient_id INT NOT NULL,
    doctor_id INT,
    record_date TIMESTAMP NOT NULL,
    description TEXT,
    treatment TEXT,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    FOREIGN KEY (patient_id) REFERENCES Patients(patient_id),
    FOREIGN KEY (doctor_id) REFERENCES Doctors(doctor_id)
);
```

# Prescriptions table

```
CREATE TABLE Prescriptions (
    prescription_id SERIAL PRIMARY KEY,
    patient_id INT NOT NULL,
    doctor_id INT NOT NULL,
    prescription_date TIMESTAMP NOT NULL,
    medication TEXT NOT NULL,
    dosage TEXT NOT NULL,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    FOREIGN KEY (patient_id) REFERENCES Patients(patient_id),
    FOREIGN KEY (doctor_id) REFERENCES Doctors(doctor_id)
);
```

# Billing table

```
CREATE TABLE Billing (
    bill_id SERIAL PRIMARY KEY,
    patient_id INT NOT NULL,
    appointment_id INT,
    amount DECIMAL(10, 2) NOT NULL,
    payment_status VARCHAR(20) CHECK (payment_status IN ('Pending', 'Paid',
'Cancelled')),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    FOREIGN KEY (patient_id) REFERENCES Patients(patient_id),
    FOREIGN KEY (appointment_id) REFERENCES Appointments(appointment_id)
);
```

# PostgreSQL Schema for Healthcare Domain

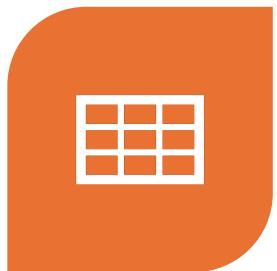
This schema design example showcases how to structure a PostgreSQL database for a healthcare application, ensuring data integrity, security, and performance while following best practices in database design.

# Normalization and Denormalization

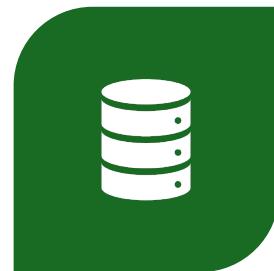
Surendra Panpaliya  
GKTCS Innovations

# Normalization

---



PROCESS OF  
ORGANIZING DATA



IN A DATABASE



TO REDUCE  
REDUNDANCY AND



IMPROVE DATA  
INTEGRITY

# Normalization

Involves dividing

large tables into smaller

related tables

using relationships

between them.

# Normalization Steps

---



**FIRST NORMAL FORM  
(1NF)**



**SECOND NORMAL  
FORM (2NF)**



**THIRD NORMAL  
FORM (3NF)**

# First Normal Form (1NF)

Ensure that each table

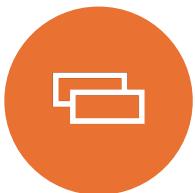
has a primary key.

Eliminate duplicate columns

from the same table.

# First Normal Form (1NF)

---



Create separate  
tables



for each group of  
related data and



identify each row  
with



a unique column



or set of columns

# Second Normal Form (2NF)

---

Ensure the table is in 1NF.

Remove subsets of data

that apply to multiple rows

place them in separate tables.

# Second Normal Form (2NF)

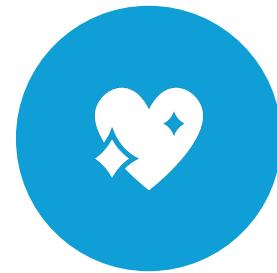
---



CREATE  
RELATIONSHIPS  
BETWEEN



THESE NEW TABLES  
AND



THEIR  
PREDECESSORS



USING FOREIGN  
KEYS.

# Third Normal Form (3NF)

---

Ensure the table is in 2NF.

Remove columns

that are not dependent

on the primary key

# Unnormalized Table

PatientID	PatientName	DoctorID	DoctorName	Appointment Date	Medication	Dosage
1	John Doe	10	Dr. Smith	2024-07-01	Aspirin	100mg
1	John Doe	10	Dr. Smith	2024-07-15	Ibuprofen	200mg
2	Jane Doe	20	Dr. Brown	2024-07-05	Paracetamol	500mg

# Normalized Tables

## Patients Table

PatientID	PatientName
1	John Doe
2	Jane Doe

# Doctors Table

DoctorID	DoctorName
10	Dr. Smith
20	Dr. Brown

# Appointments Table

AppointmentID	PatientID	DoctorID	AppointmentDate
1	1	10	2024-07-01
2	1	10	2024-07-15
3	2	20	2024-07-05

# Medications Table

MedicationID	AppointmentID	Medication	Dosage
1	1	Aspirin	100mg
2	2	Ibuprofen	200mg
3	3	Paracetamol	500mg

# Denormalization

Process of combining

normalized tables

to improve read performance

by reducing

the number of joins.

# Denormalization



PROCESS MAY  
INTRODUCE



REDUNDANCY



BUT CAN BE  
USEFUL



FOR  
OPTIMIZING



READ-HEAVY  
OPERATIONS.

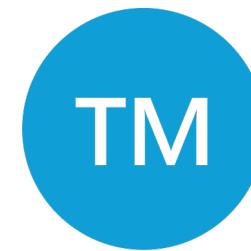
# Denormalization Considerations



Query Performance



Data Redundancy



Use Cases

# Query Performance

Combine tables

to reduce the number of joins

required in queries.

Useful for read-heavy operations

where speed is crucial.

# Data Redundancy

Introduce redundancy

to improve performance,

but manage it carefully

to maintain data integrity

# Use Cases

Reports and dashboards

where denormalized data

can be queried quickly.

Applications with

high read-to-write ratio.

# Normalized Patients Table

PatientID	PatientName
1	John Doe
2	Jane Doe

# Normalized Appointments Table

AppointmentID	PatientID	DoctorID	AppointmentDate
1	1	10	2024-07-01
2	1	10	2024-07-15
3	2	20	2024-07-05

# Denormalized Table

PatientID	PatientName	AppointmentID	DoctorID	AppointmentDate	Medication	Dosage
1	John Doe	1	10	2024-07-01	Aspirin	100mg
1	John Doe	2	10	2024-07-15	Ibuprofen	200mg
2	Jane Doe	3	20	2024-07-05	Paracetamol	500mg

# Query Performance



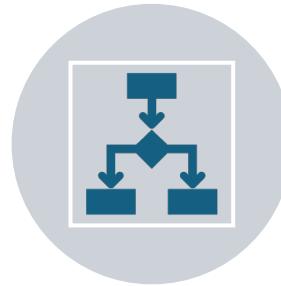
Denormalized data



allows for quicker read access



as it eliminates



the need for multiple joins.

# Query Performance



Ideal for applications



requiring fast read operations,



such as healthcare dashboards



showing patient data,



appointments, and



medications in a single view.

# Conclusion and Summary

Normalization

denormalization

have their place

in database design.

# Conclusion and Summary

Normalization  
ensures

data integrity  
and

reduces  
redundancy,

which is  
crucial for

write-heavy  
operations  
and

maintaining  
consistency.

# Conclusion and Summary

Denormalization

used to optimize  
read  
performance

in read-heavy  
applications,

such as  
Reporting and

dashboard  
applications.

# Entity-Relationship (ER) Modeling

Surendra Panpaliya

# Entity-Relationship (ER) Model

---



Visually  
represents



Data



Relationships  
between



different  
entities



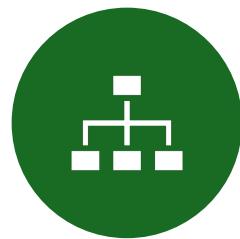
in a system.

# ER Model in Healthcare

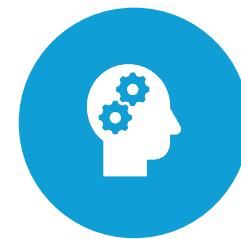
---



ER MODEL HELPS  
IN



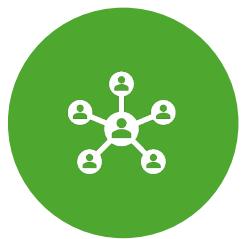
ORGANIZING



UNDERSTANDING



VARIOUS DATA  
COMPONENTS



INTERACTIONS.

# **Key Entities and Relationships**

Surendra Panpaliya  
GKTCS Innovations

# Patient

---

Attributes:

---

PatientID (PK),

---

FirstName, LastName,

---

DateOfBirth, Gender,

---

ContactNumber, Email,

---

Address

# Doctor

---

Attributes:

---

DoctorID (PK),

---

FirstName, LastName,

---

Specialty,

---

ContactNumber,

---

Email

# Appointment

---

Attributes:

---

AppointmentID (PK),

---

PatientID (FK),

---

DoctorID (FK),

---

AppointmentDate,

---

Status, Notes

# Appointment

---

Relationships:

---

Each appointment is

---

associated with

---

one patient and

---

one doctor.

# **MedicalRecord**

---

**Attributes:**

---

**RecordID (PK), PatientID (FK),**

---

**DoctorID (FK), RecordDate,**

---

**Description, Treatment**

# MedicalRecord

---

Relationships:

---

Each medical record is

---

associated with one patient and

---

optionally one doctor.

# Prescription

---

Attributes:

---

PrescriptionID (PK),

---

PatientID (FK),

---

DoctorID (FK),

---

PrescriptionDate,

---

Medication, Dosage

# Prescription

---

Relationships:

---

Each prescription is

---

associated with

---

one patient and

---

one doctor.

# Billing

---

Attributes:

---

BillID (PK),

---

PatientID (FK),

---

AppointmentID (FK),

---

Amount,

---

PaymentStatus

# Billing

---

Relationships:

---

Each bill is associated with

---

one patient and

---

one appointment.

# User



Attributes:



UserID (PK),



Username,



PasswordHash,



Role (Admin, Doctor, Patient)

# User



Relationships:



a patient,



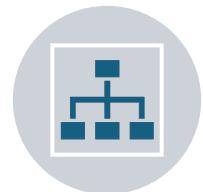
Each user can  
be



doctor, or



associated with



administrative  
role.

# Insurance

Attributes:

InsuranceID (PK),

ProviderName,

PlanType,

PatientID (FK)

# Insurance



RELATIONSHIPS:



EACH INSURANCE  
PLAN IS



ASSOCIATED  
WITH



ONE PATIENT.

Patient	Doctor	User
PatientID (PK)	DoctorID (PK)	UserID (PK)
FirstName	FirstName	Username
LastName	LastName	PasswordHash
DateOfBirth	Specialty	Role
Gender	ContactNumber	
ContactNumber	Email	
Email		
Address		

Appointment	
AppointmentID (PK)	
PatientID (FK)	
DoctorID (FK)	
AppointmentDate	
Status	
Notes	

MedicalRecord	
RecordID (PK)	
PatientID (FK)	
DoctorID (FK)	
RecordDate	
Description	
Treatment	

Prescription		Billing	
PrescriptionID (PK)		BillID (PK)	
PatientID (FK)		PatientID (FK)	
DoctorID (FK)		AppointmentID (FK)	
PrescriptionDate		Amount	
Medication		PaymentStatus	
Dosage			

Insurance	
InsuranceID (PK)	
ProviderName	
PlanType	
PatientID (FK)	

# Managing Tables

Surendra Panpaliya  
GKTCS Innovations

# Altering Tables

Allows to modify

the structure

of existing tables

Adding, Modifying

Dropping columns

# Altering Tables

---

Adding a New Column to a Table

```
ALTER TABLE Patients
```

```
ADD COLUMN middle_name VARCHAR(50);
```

Add a new column middle\_name to the Patients table

# Modifying an Existing Column

---

Change the data type of the contact\_number column in the Doctors table

```
ALTER TABLE Doctors  
ALTER COLUMN contact_number TYPE VARCHAR(20);
```

## **Set a default value for the status column in the Appointments table**

---

```
ALTER TABLE Appointments
```

```
ALTER COLUMN status SET DEFAULT 'Scheduled';
```

# Dropping a Column from a Table

---

- Remove the specialty column from the Doctors table:
- ALTER TABLE Doctors
- DROP COLUMN specialty;

# Renaming a Column

---

**Rename the contact\_number column to phone\_number in the Patients table**

```
ALTER TABLE Patients
```

```
RENAME COLUMN contact_number TO phone_number;
```

# Adding a Foreign Key Constraint

---

- Add a **foreign key constraint** to the **MedicalRecords** table to reference the **Doctors** table:
  - ALTER TABLE MedicalRecords
  - ADD CONSTRAINT fk\_doctor
  - FOREIGN KEY (doctor\_id) REFERENCES Doctors(doctor\_id);

# Dropping a Foreign Key Constraint

---

**Drop the foreign key constraint  
from the MedicalRecords table**

```
ALTER TABLE MedicalRecords  
DROP CONSTRAINT fk_doctor;
```

# Dropping Tables

---

- Removes Tables
- Entirely from
- Database

# Drop the Insurance table

---

```
DROP TABLE Insurance;
```

# Dropping a Table Only If It Exists

**Drop the Billing table only if it exists**

```
DROP TABLE IF EXISTS Billing;
```

# Dropping Multiple Tables

**Drop the Appointments and Prescriptions tables**

```
DROP TABLE Appointments, Prescriptions;
```

# Dropping a Table and Its Dependent Objects

**Drop the Users table and any objects that depend on it (e.g., foreign key constraints)**

```
DROP TABLE Users CASCADE;
```

# Example Workflow

**Scenario:**  
**Adding a Column and**  
**Then Dropping a Table**

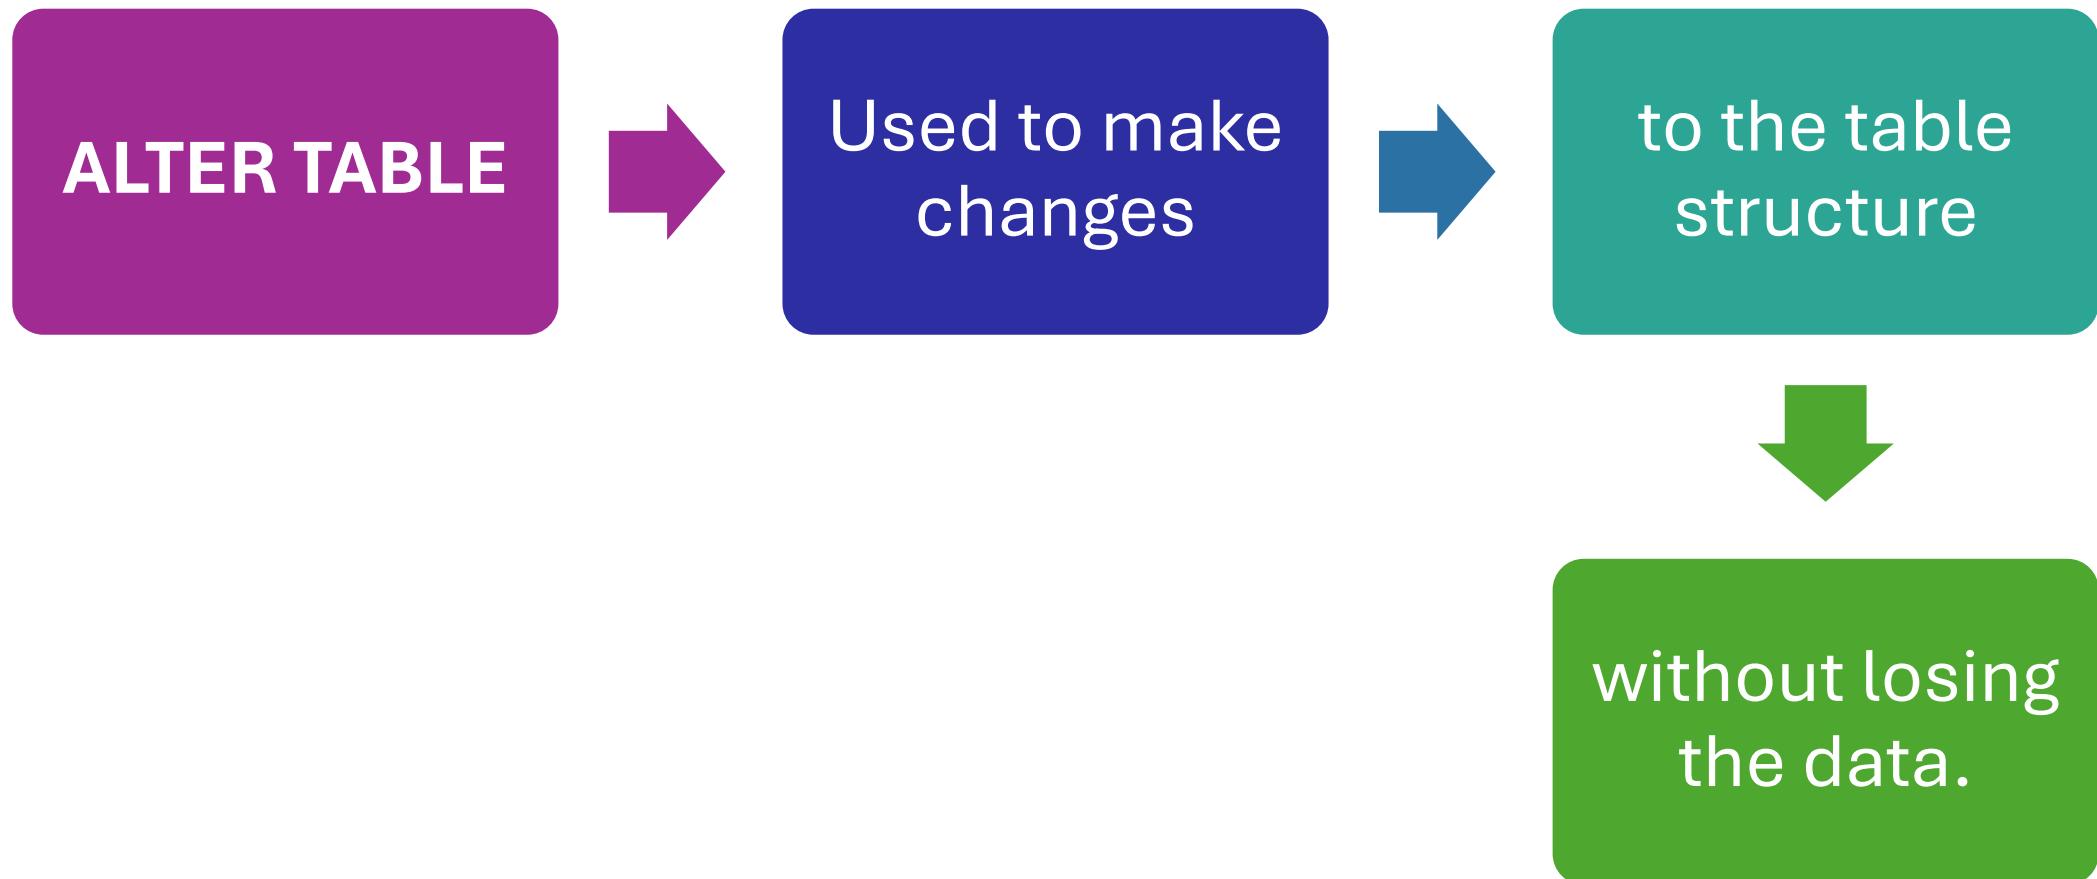
# **Step1: Add a new column emergency\_contact to the Patients table**

```
ALTER TABLE Patients  
ADD COLUMN emergency_contact VARCHAR(50);
```

# **Step 2: Drop the MedicalRecords table**

```
DROP TABLE MedicalRecords;
```

# Important Notes



# Important Notes

## DROP TABLE

remove the table and

all its data from

the database permanently.

# Important Notes

---



USING  
CASCADE



WITH **DROP**  
TABLE



ENSURES THAT



ALL DEPENDENT  
OBJECTS



ARE ALSO  
DROPPED.

# Important Notes

---

**IF EXISTS**  
is useful

To avoid  
errors

When  
attempting

To drop  
tables

That may  
not exist

# **Indexing for Performance**

**Surendra Panpaliya**  
**GKTCS Innovations**

# Indexing for Performance



Indexing is



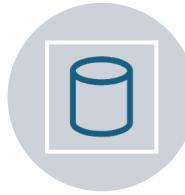
a powerful  
technique



to improve



the  
performance of



database  
queries.

# **Indexing for Performance**

Proper indexing

can significantly

speed up

data retrieval operations.

# Types of Indexes

B-tree Index

Hash Index

GIN (Generalized Inverted Index)

GiST (Generalized Search Tree)

BRIN (Block Range INdex)

# B-tree Index

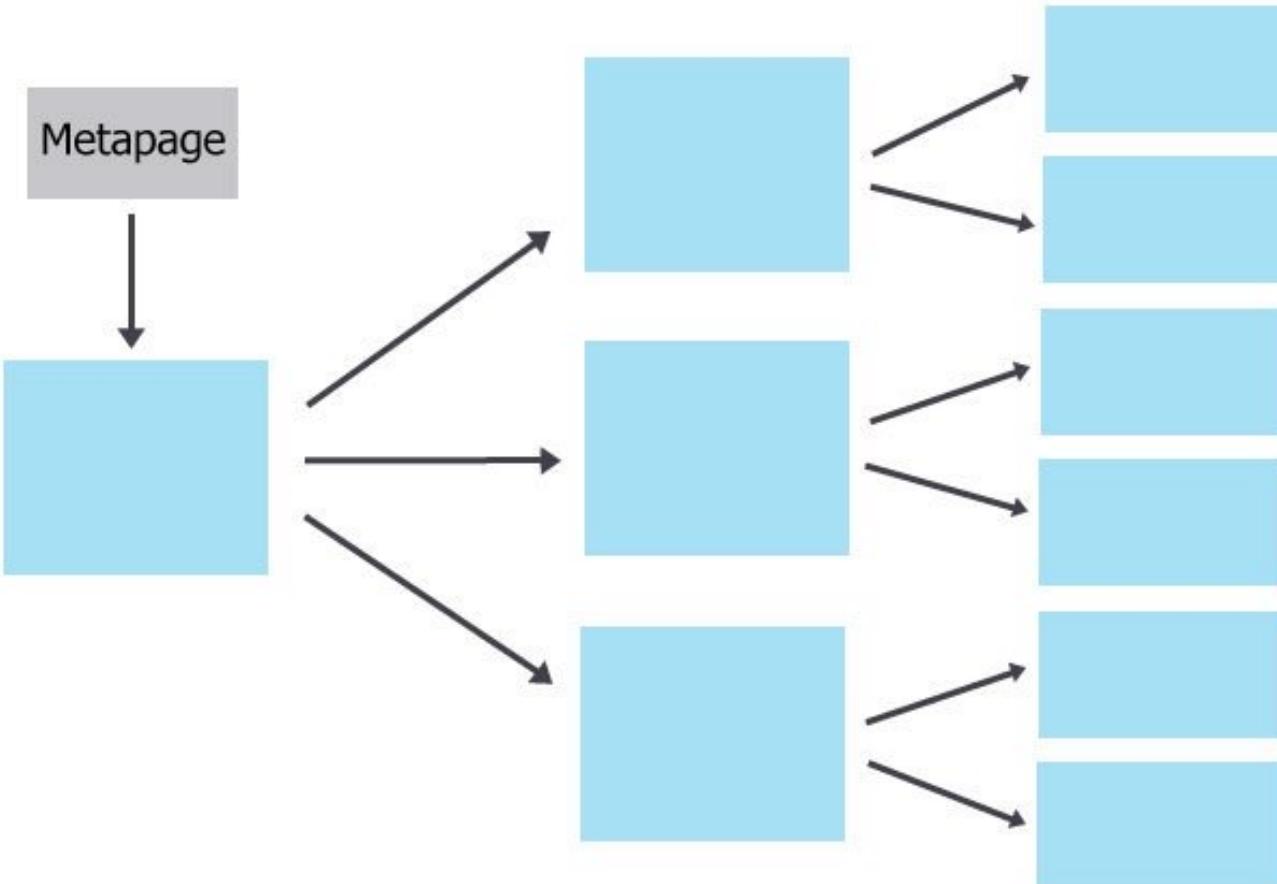
Default

Most common type of index

Ideal for equality

Range queries

# PostgreSQL complete B-tree



# Creating a B-tree Index



**Index on the email column of the Patients table**



```
CREATE INDEX idx_patients_email ON Patients(email);
```

# Creating a Unique Index



**Ensure email addresses are unique in the Doctors table:**



```
CREATE UNIQUE INDEX idx_doctors_email_unique ON  
Doctors(email);
```

# **Creating a Composite Index**

**Index on the `first_name` and `last_name` columns of the Patients table**

```
CREATE INDEX idx_patients_name ON Patients(first_name,  
last_name);
```

# Creating a Partial Index



**Index on active appointments only in the Appointments table**



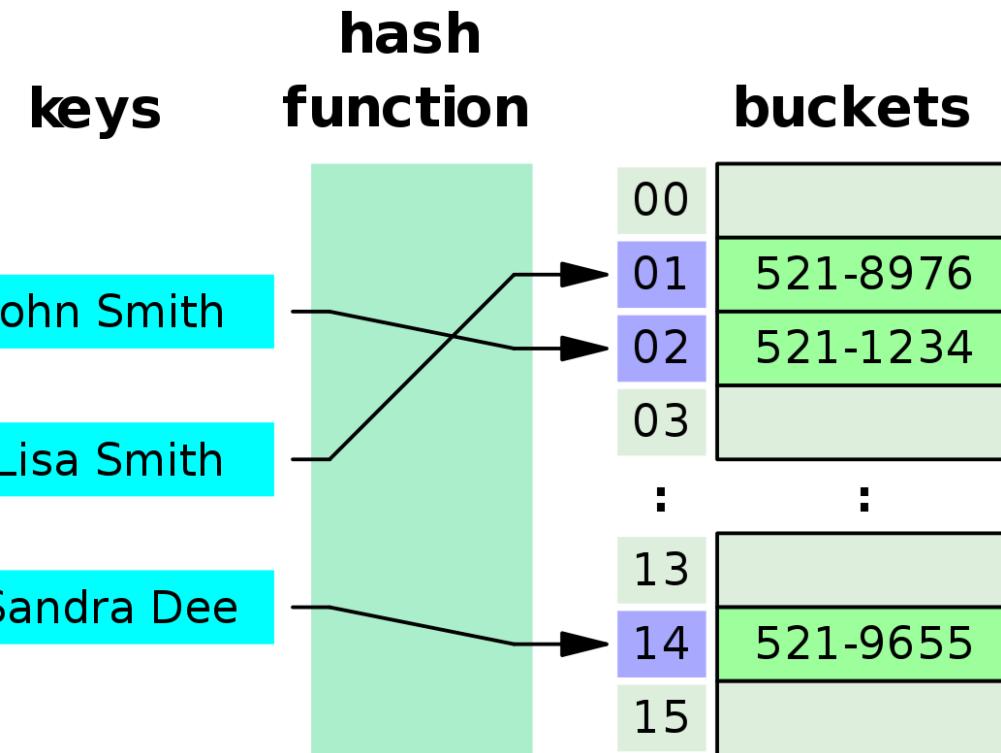
```
CREATE INDEX idx_active_appointments ON  
Appointments(appointment_date)
```



```
WHERE status = 'Scheduled';
```

# Hash Index

Useful for equality comparisons.



# **GIN (Generalized Inverted Index)**

---

Ideal for array and full-text search operations.

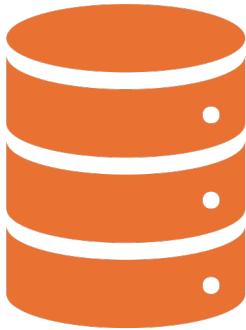
# Creating a GIN Index

---

**Index for full-text search on the description column of the MedicalRecords table:**

```
CREATE INDEX idx_medicalrecords_description_gin ON MedicalRecords  
USING gin(to_tsvector('english', description));
```

# GiST (Generalized Search Tree)

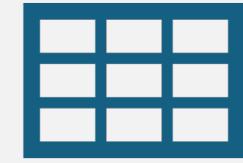


Useful for complex data types



like geometric data.

# BRIN (Block Range INdex)



EFFICIENT FOR VERY  
LARGE TABLES

WHERE DATA IS  
NATURALLY CLUSTERED.

# Creating a BRIN Index



**Index on the `appointment_date` column of the Appointments table for large datasets**



`CREATE INDEX idx_appointments_date_brin ON Appointments`



`USING brin(appointment_date);`

# Managing Indexes

Surendra Panpaliya

# Dropping an Index

**Drop the index on the email column of the Patients table**

```
DROP INDEX idx_patients_email;
```

# Reindexing a Table

**Rebuild all indexes on the Patients table**

```
REINDEX TABLE Patients;
```

# Viewing Indexes

**List all indexes in the current database**

```
postgres#\di
```

# Analyzing Index Usage



**Analyze the Appointments table**



**to update statistics for query planner**



**ANALYZE Appointments;**

# Best Practices for Indexing

Surendra Panpaliya  
GKTCS Innovations

# Index Selective Columns

Index columns that are frequently

used in

WHERE clauses,

JOIN conditions

ORDER BY clauses.

# Avoid Over-Indexing



Too many indexes



can slow down



INSERT, UPDATE, and DELETE operations



due to the overhead of



maintaining the indexes.

# Use Composite Indexes



For queries



that filter on multiple columns,



use composite indexes



to improve performance.

# Consider Index Types

Choose the appropriate index type

based on

the query patterns

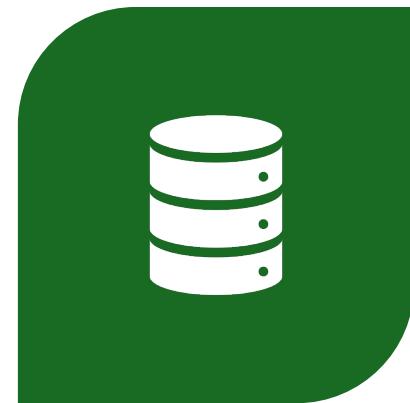
data types.

# Consider Index Types

---



FOR EXAMPLE,



USE GIN INDEXES



FOR FULL-TEXT  
SEARCH.

# Monitor Index Usage

Regularly monitor

Analyze index usage

to identify and remove

unused or redundant indexes.

# Use Partial Indexes

Create  
partial  
indexes

to index only  
a subset of  
data

that meets  
specific  
criteria.

# Maintain Indexes



REGULARLY REINDEX  
TABLES AND



UPDATE STATISTICS



TO ENSURE OPTIMAL  
PERFORMANCE.

# Scenario

**Optimizing a Query on the Appointments Table**

---

# Identify the Query to Optimize

```
SELECT * FROM Appointments  
WHERE doctor_id = 1 AND status = 'Scheduled'  
ORDER BY appointment_date DESC;
```

# Create a Composite Index

```
CREATE INDEX idx_appointments_doctor_status_date ON  
Appointments(doctor_id, status, appointment_date);
```

# Analyze the Table

ANALYZE Appointments;

# Check the Execution Plan

```
EXPLAIN ANALYZE SELECT * FROM Appointments  
WHERE doctor_id = 1 AND status = 'Scheduled'  
ORDER BY appointment_date DESC;
```

# Summary



Following indexing  
techniques



Best practices



Can significantly  
improve

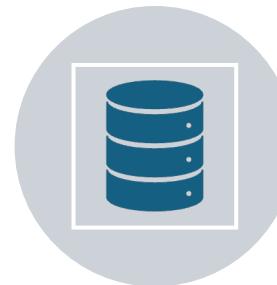


Performance of queries

# Summary



Ensuring efficient



data retrieval



better overall database



performance

# Thank You

- Surendra Panpaliya
- Founder and CEO
- GKTC Innovations
- <https://www.gktcs.com>