



PostgreSQL

PostgreSQL Architecture

Surendra Panpaliya

PostgreSQL Architecture



POSTGRESQL IS



POWERFUL OPEN-
SOURCE



RELATIONAL DATABASE
MANAGEMENT SYSTEM

PostgreSQL Architecture

Known for

Robustness,

Extensibility

Standards compliance

PostgreSQL Architecture



Understanding



its
architecture



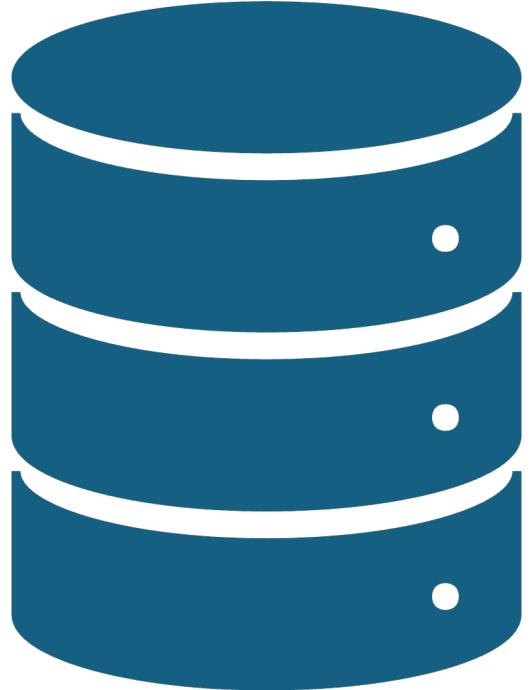
is crucial for
effective



Database
management



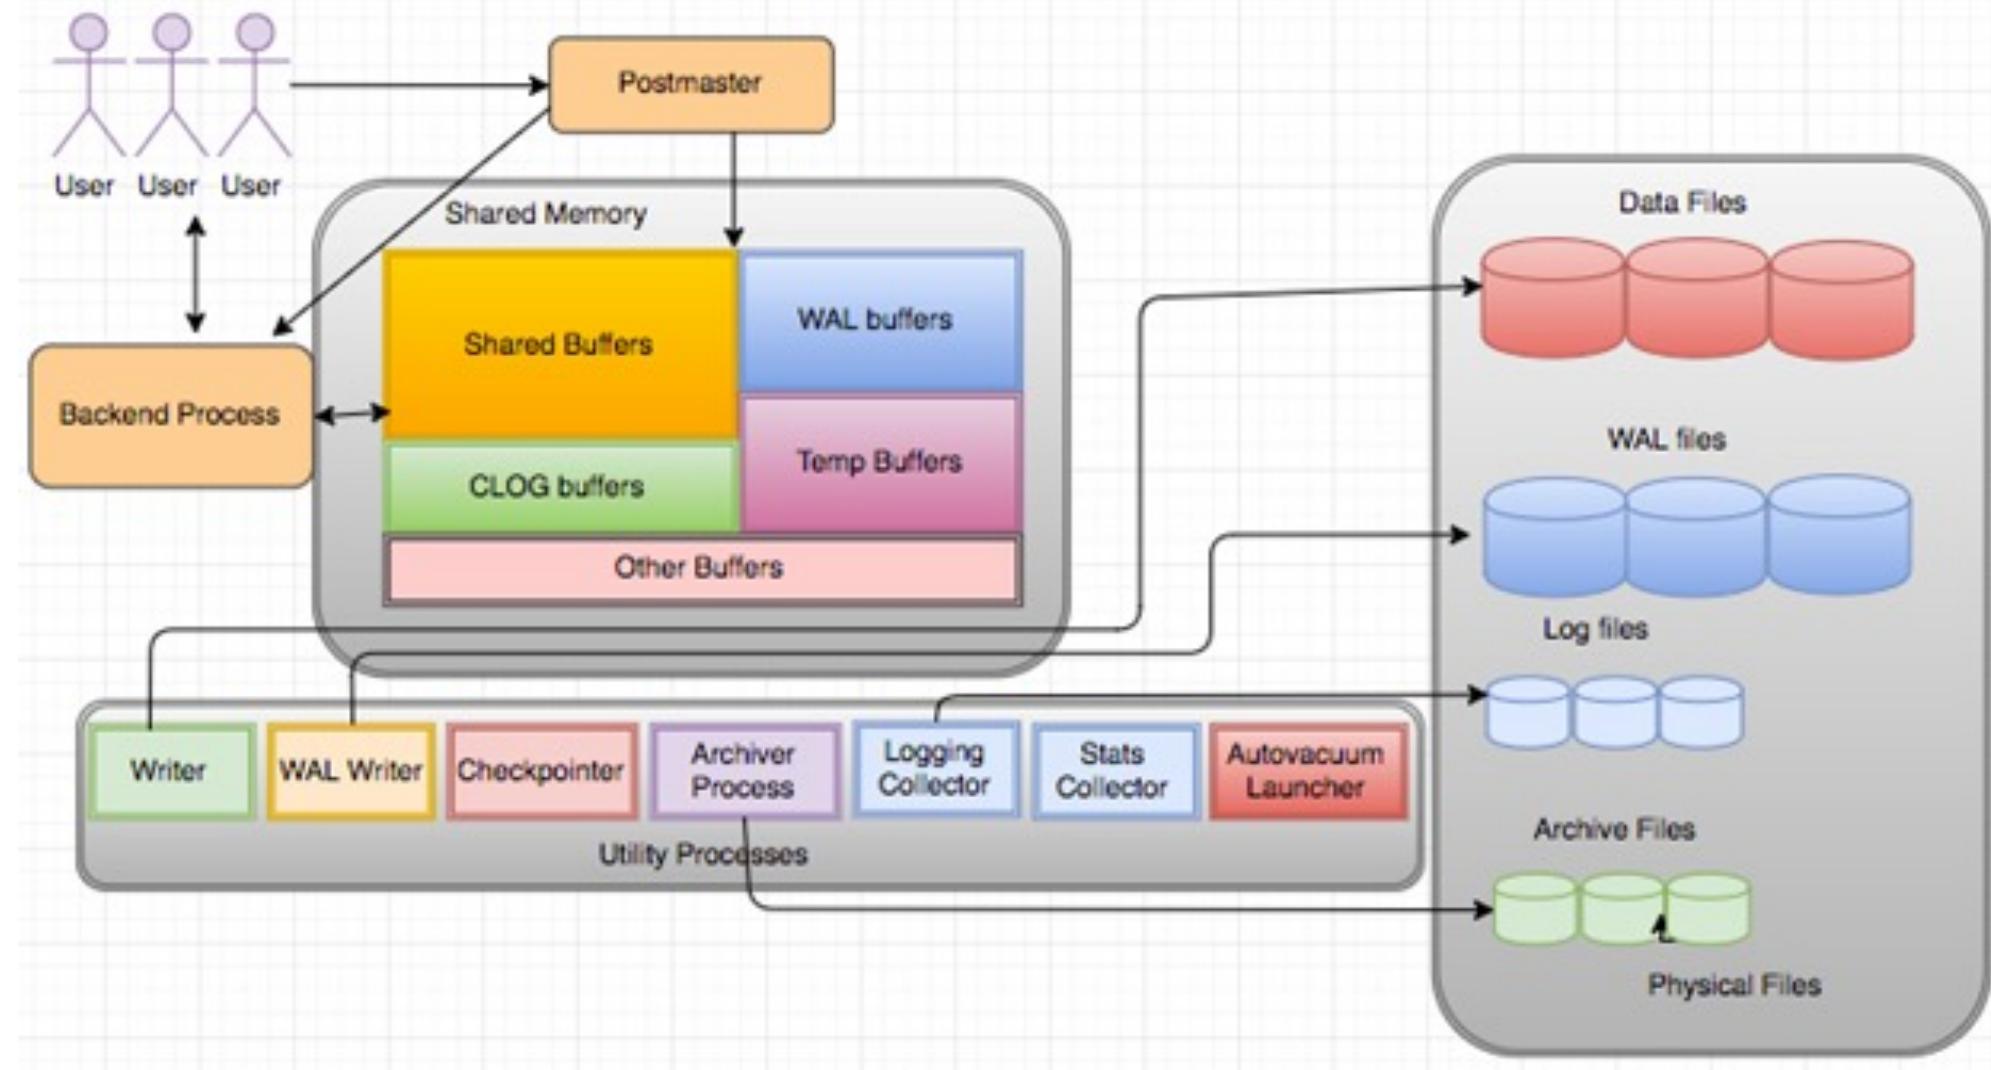
Optimization



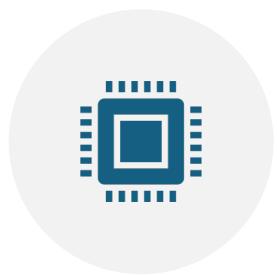
PostgreSQL Architecture Components

Surendra Panpaliya

PostgreSQL Basic Architecture



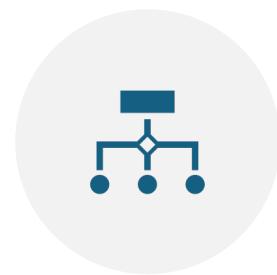
PostgreSQL Instance



CONSISTS OF SET OF
PROCESS AND
MEMORY.



POSTGRESQL USES A
SIMPLE



“PROCESS PER-
USER”



CLIENT/SERVER
MODEL.

Postmaster Process



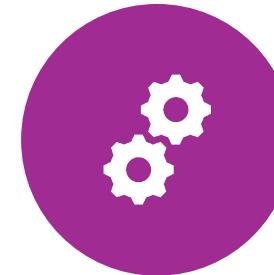
Main PostgreSQL
server process



Responsible for
Initializing the
database



Handling connections



Starting other
background processes

Postmaster Process



LISTENS FOR INCOMING
CONNECTIONS



SPAUNS NEW SERVER
PROCESSES



FOR EACH
CONNECTION.

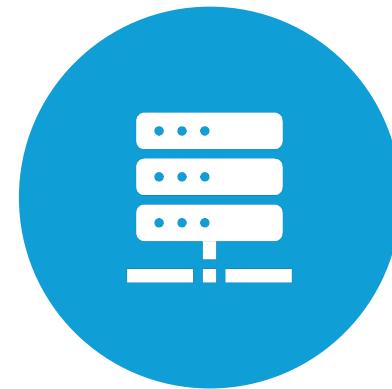
Database Cluster



A COLLECTION OF
DATABASES



MANAGED BY A
SINGLE



POSTGRESQL
SERVER INSTANCE.

Database Cluster



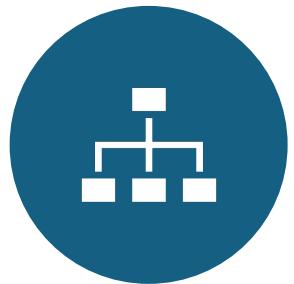
Shared Memory



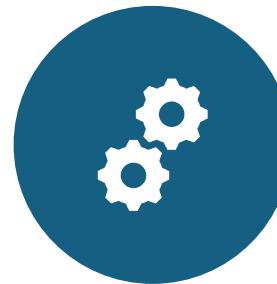
Shared Buffers



WAL Buffers

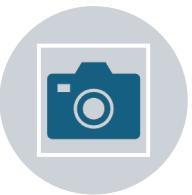


Work_mem

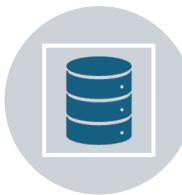


Maintenance_work_mem

Shared Buffers



Main memory area



PostgreSQL Caches



Data Blocks Read from the disk.



Improves performance



by reducing disk I/O

WAL Buffers

Stores Write-Ahead Log

Entries before Written to disk

Ensures data durability and crash recovery

Work_mem

Memory used for

internal operations

like sorting
and hash
tables

during query
execution.

Maintenance_work_mem



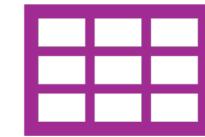
Memory
allocated



Maintenance
tasks



VACUUM,
CREATE INDEX



ANALYZE
operations.

Background Processes

Background Writer (bgwriter)

WAL Writer (walwriter)

Autovacuum Daemon

Stats Collector

Background Writer (bgwriter)



Periodically writes dirty pages



from the shared buffers to the disk



To ensure the buffer cache



has space for new pages.

WAL Writer (walwriter)



WRITES WAL DATA



FROM WAL BUFFERS



TO THE DISK.

Autovacuum Daemon



MANAGES AUTOMATIC
VACUUMING OF



DATABASE TO RECLAIM
STORAGE



MAINTAIN TABLE
STATISTICS.

Stats Collector

Collects and aggregates

Database activity

Performance statistics.

Process Architecture



POSTMASTER
PROCESS



BACKEND
PROCESSES



BACKGROUND
WORKER PROCESSES

Postmaster Process

The parent process that

initializes the database system and

handles incoming connections.

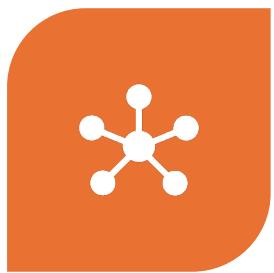
Backend Processes

Individual server processes

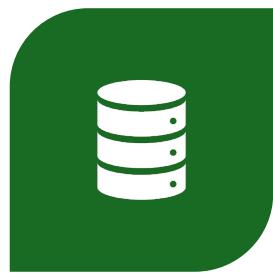
Spawned by the Postmaster

To handle client connections.

Backend Processes



EACH
CONNECTION



TO THE DATABASE



HAS A
CORRESPONDING



BACKEND
PROCESS

Background Worker Processes

Custom or
built-in

background
jobs

perform
specific
tasks,

such as
parallel
queries.

Storage

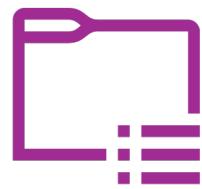
Data Directory

Data Files

WAL Files

Configuration Files

Data Directory



File system
directory



Stores database
files



Configuration
files



Other metadata

Data Files

Store
actual

Database
tables

Indexes

WAL Files



Logs of all
changes



Made to the
database



Used for



Crash Recovery
and Replication

Configuration Files



Key configuration
files include



postgresql.conf
(main configuration)



pg_hba.conf (client
authentication)



pg_ident.conf (user
identity mapping)

Storage



Data Directory: The file system directory where PostgreSQL stores its database files, configuration files, and other metadata.



Data Files: Files that store actual database tables and indexes.



WAL Files: Logs of all changes made to the database. Used for crash recovery and replication.



Configuration Files: Key configuration files include `postgresql.conf` (main configuration), `pg_hba.conf` (client authentication), and `pg_ident.conf` (user identity mapping).

Logical Storage Structures

Tablespaces

Databases

Schemas

Tablespaces

Logical storage units

Map to file system locations

Allow for the distribution of data

Across different storage locations

Databases



LOGICAL
COLLECTIONS OF



SCHEMAS,
TABLES, INDEXES



FUNCTIONS, AND

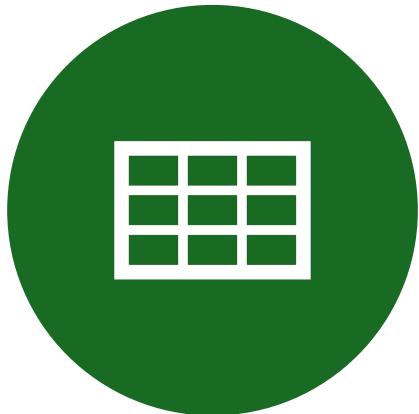


OTHER OBJECTS.

Schemas



NAMESPACE WITHIN A
DATABASE

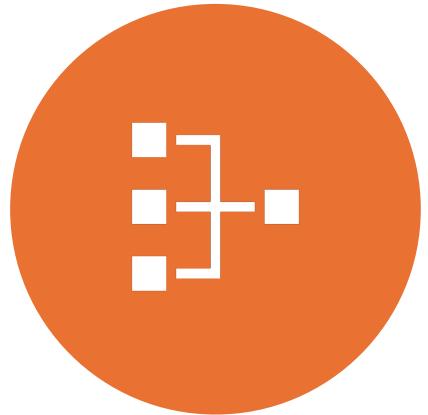


CONTAINS TABLES, VIEWS,
FUNCTIONS



SCHEMAS HELP ORGANIZE
OBJECTS LOGICALLY

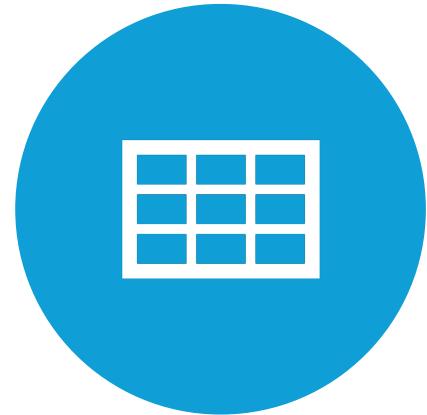
Concurrency Control



MVCC (MULTI-VERSION
CONCURRENCY CONTROL)



ENSURES DATA CONSISTENCY
AND ISOLATION



BY MAINTAINING MULTIPLE
VERSIONS OF DATA ROWS

MVCC

Allows concurrent transactions

to read and write

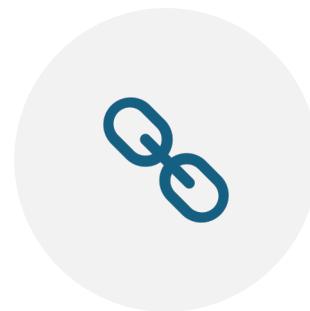
without blocking each other

using a mechanism called snapshots.

Indexing



VARIOUS INDEXING
METHODS ARE
AVAILABLE



B-TREE, HASH, GIST,



SP-GIST, GIN, AND BRIN,



HELP OPTIMIZE QUERY
PERFORMANCE

Replication and High Availability



Streaming Replication



Logical Replication

Streaming Replication



Provides real-time



data replication



to standby servers



for high availability and



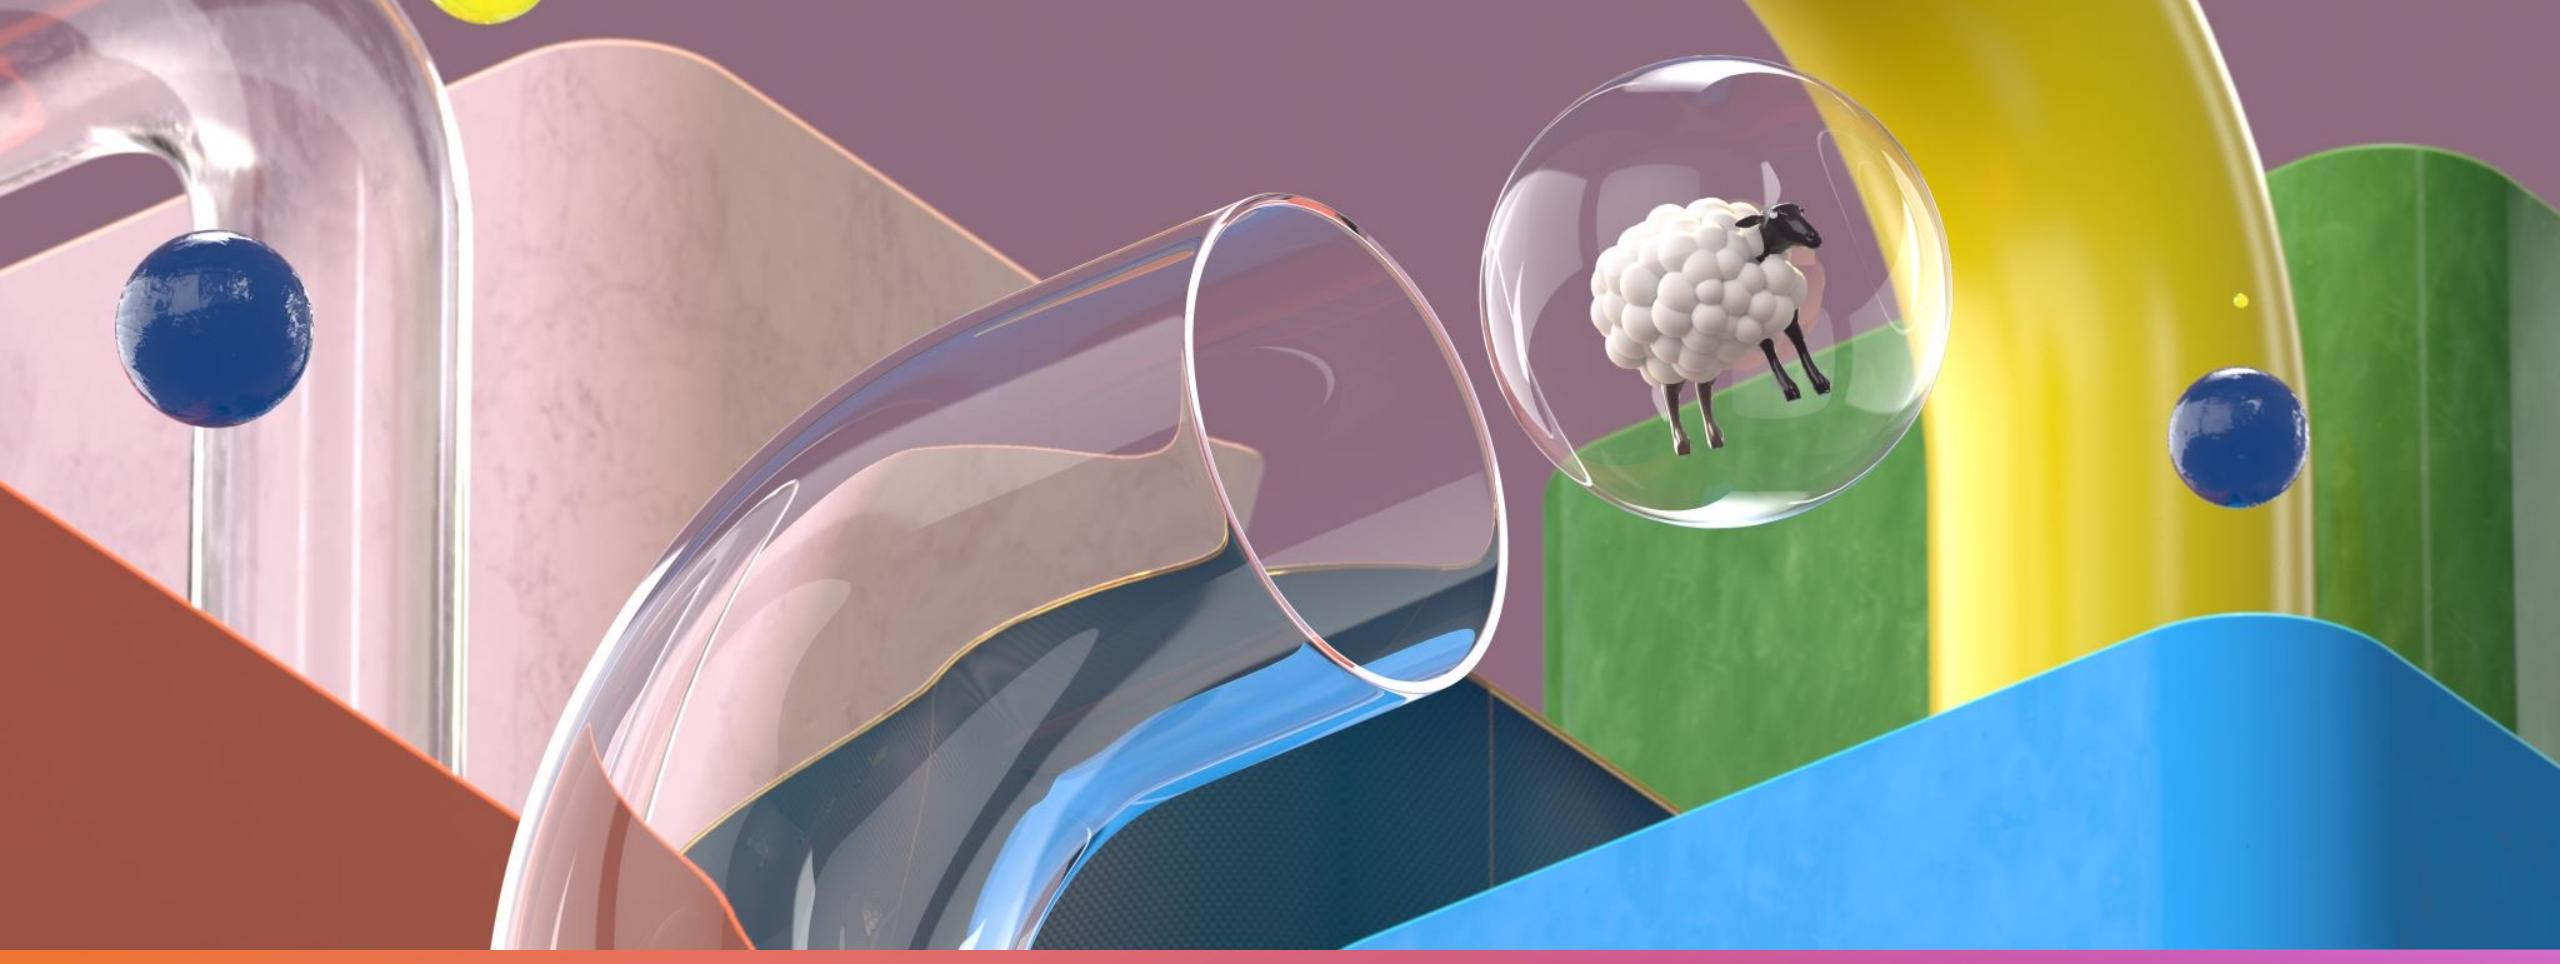
disaster recovery.

Logical Replication

Allows selective
replication of

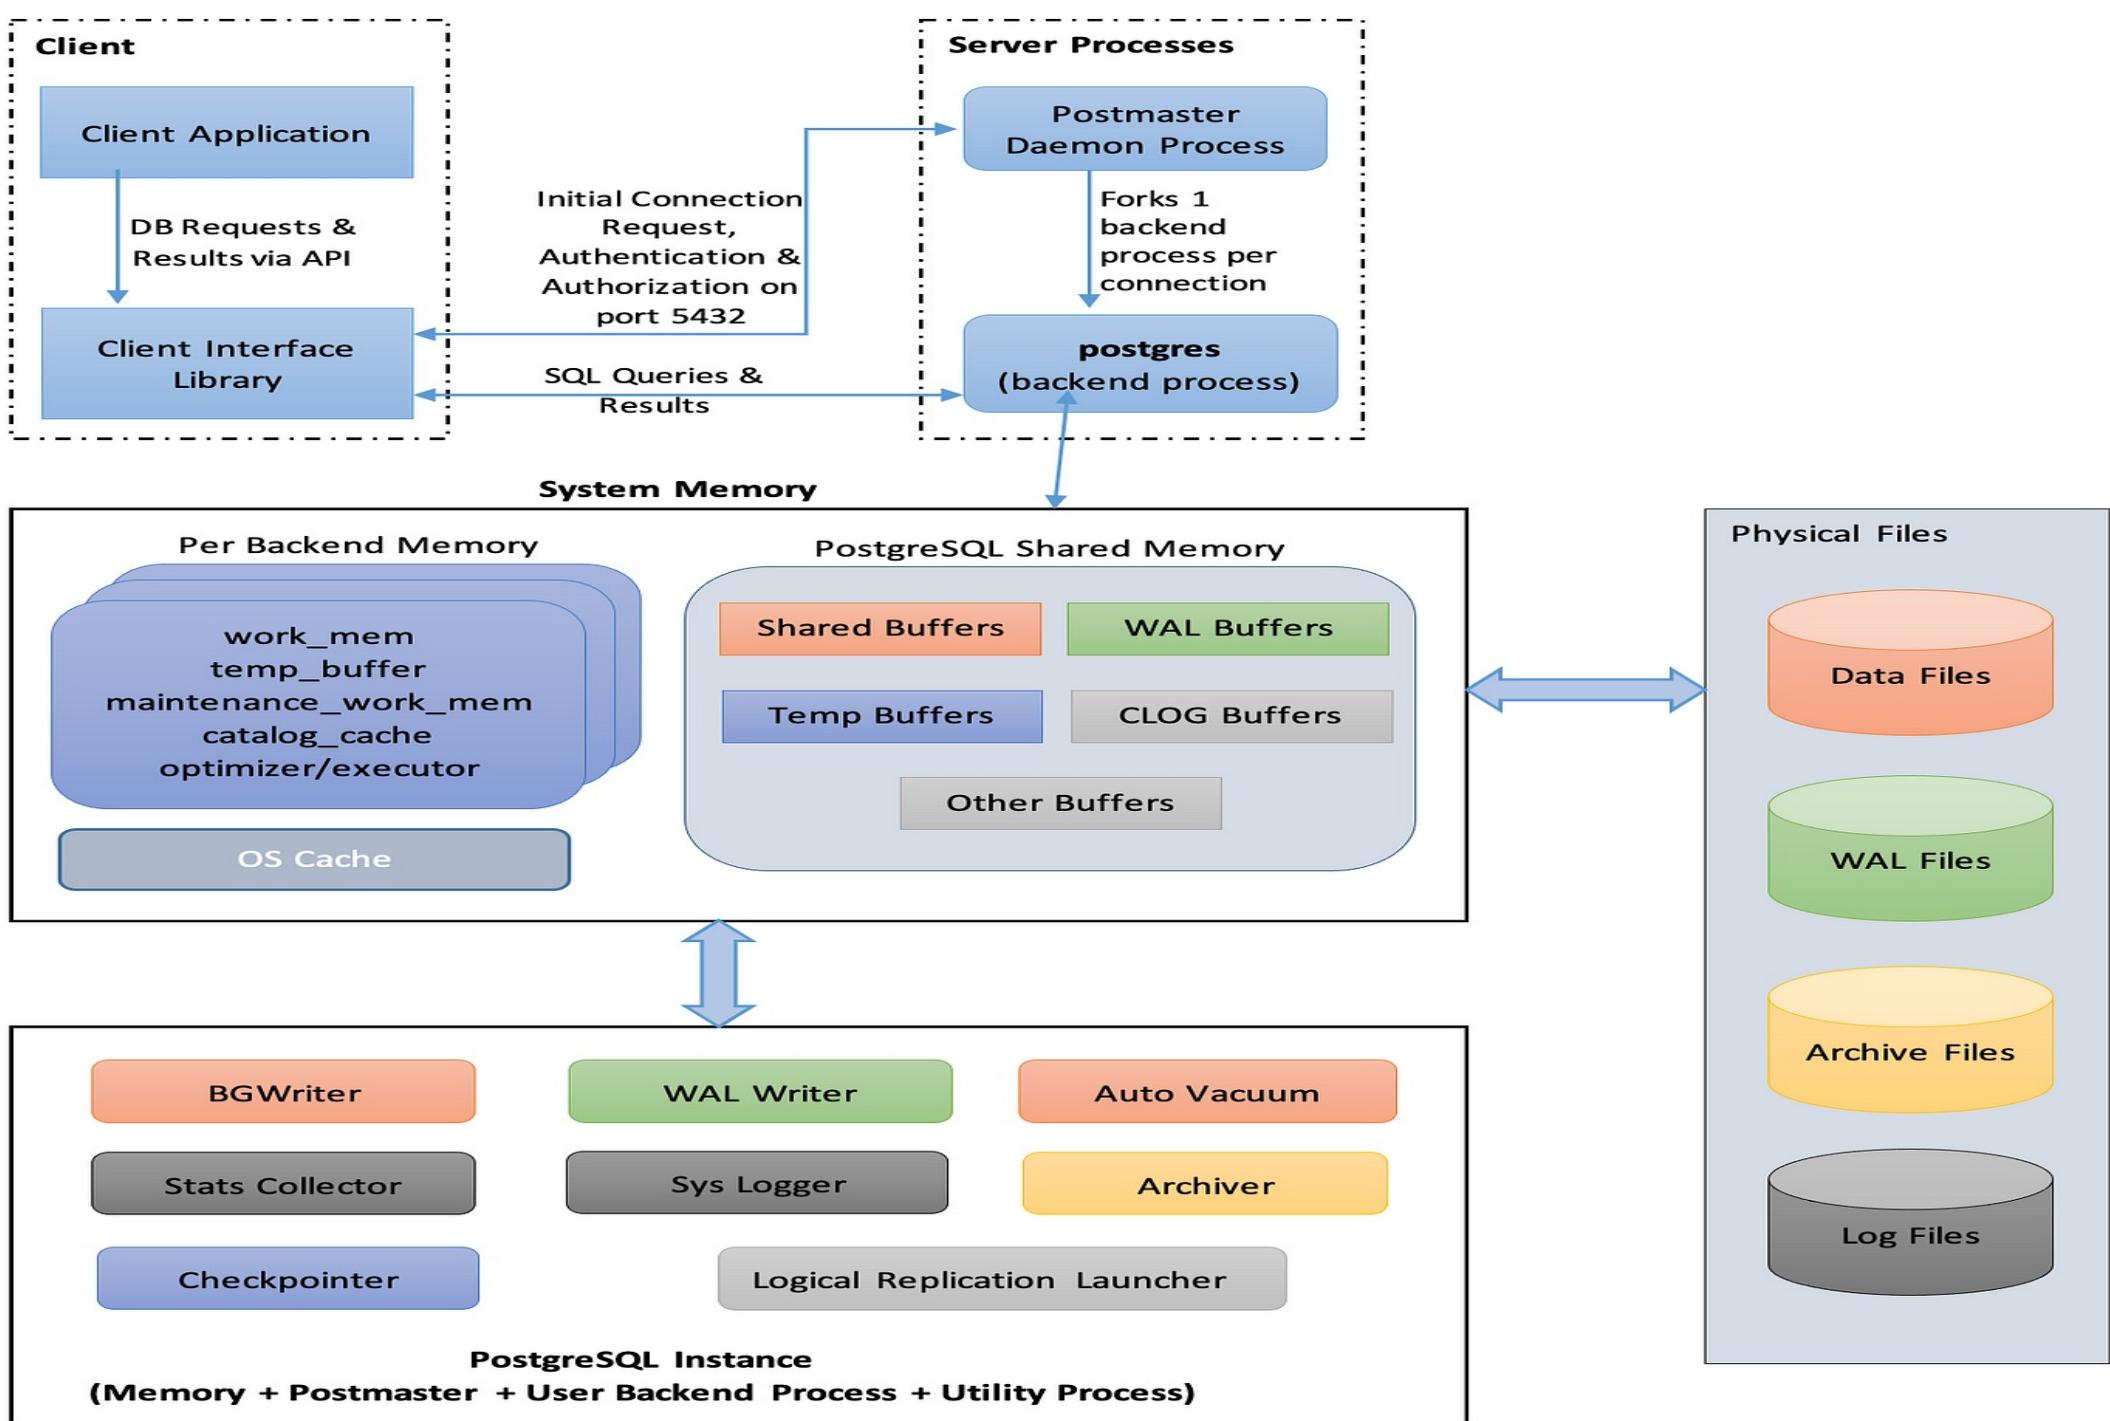
data changes
based on

user-defined
configurations



PostgreSQL Process Flow

Surendra Panpaliya



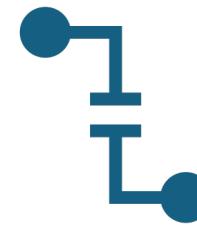
Client Connection



When a client connects
to PostgreSQL



Postmaster process
accepts the connection



Spawns a new backend
process to handle it

Query Processing



Parser



Planner/Optimizer

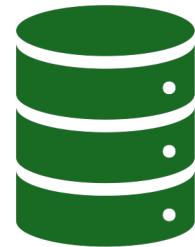


Executor

Parser



Converts



SQL queries into



Parse tree

Planner/Optimizer

Converts the
parse tree

into an
execution plan

optimizing for
performance

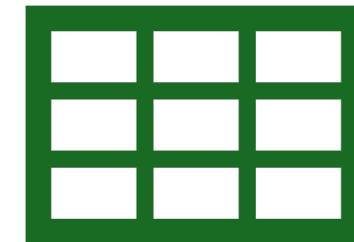
Executor

Executes the
plan

By retrieving

Manipulating
data as needed.

Buffer Management



Data is fetched from the disk

into the shared buffers.

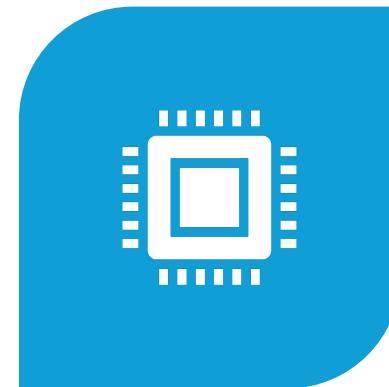
Buffer Management



ENSURES



FREQUENTLY
ACCESSED DATA



CACHED IN MEMORY

Transaction Management



ENSURES



ACID
PROPERTIES



THROUGH



MVCC AND
WAL

Background Tasks



PROCESSES



LIKE



BACKGROUND
WRITER.

Background Tasks

Autovacuum
daemon

Ensure data
integrity

Optimal
performance

by Managing

Memory and
storage

Summary

PostgreSQL Architecture

designed for reliability,

performance, and

extensibility.

Summary

Uses a combination of

shared memory and process-based management

to handle client connections,

perform query execution

manage data storage.

Summary

Its robust concurrency control,

indexing methods, and replication mechanisms

Make it suitable for

various applications,

including large-scale enterprise systems.



**Thank you for
your support and
patience**

Surendra Panpaliya
Founder and CEO
GKTCS Innovations
<https://www.gktcs.com>