



PostgreSQL

**Advanced Postgres**

Surendra Panpaliya

# Agenda

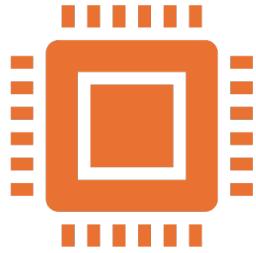
Introduction to PostgreSQL

History and development of PostgreSQL

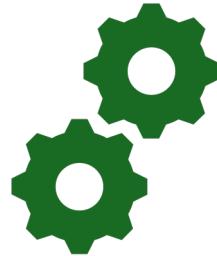
Key features and benefits of PostgreSQL

PostgreSQL vs. Oracle database systems

# Agenda



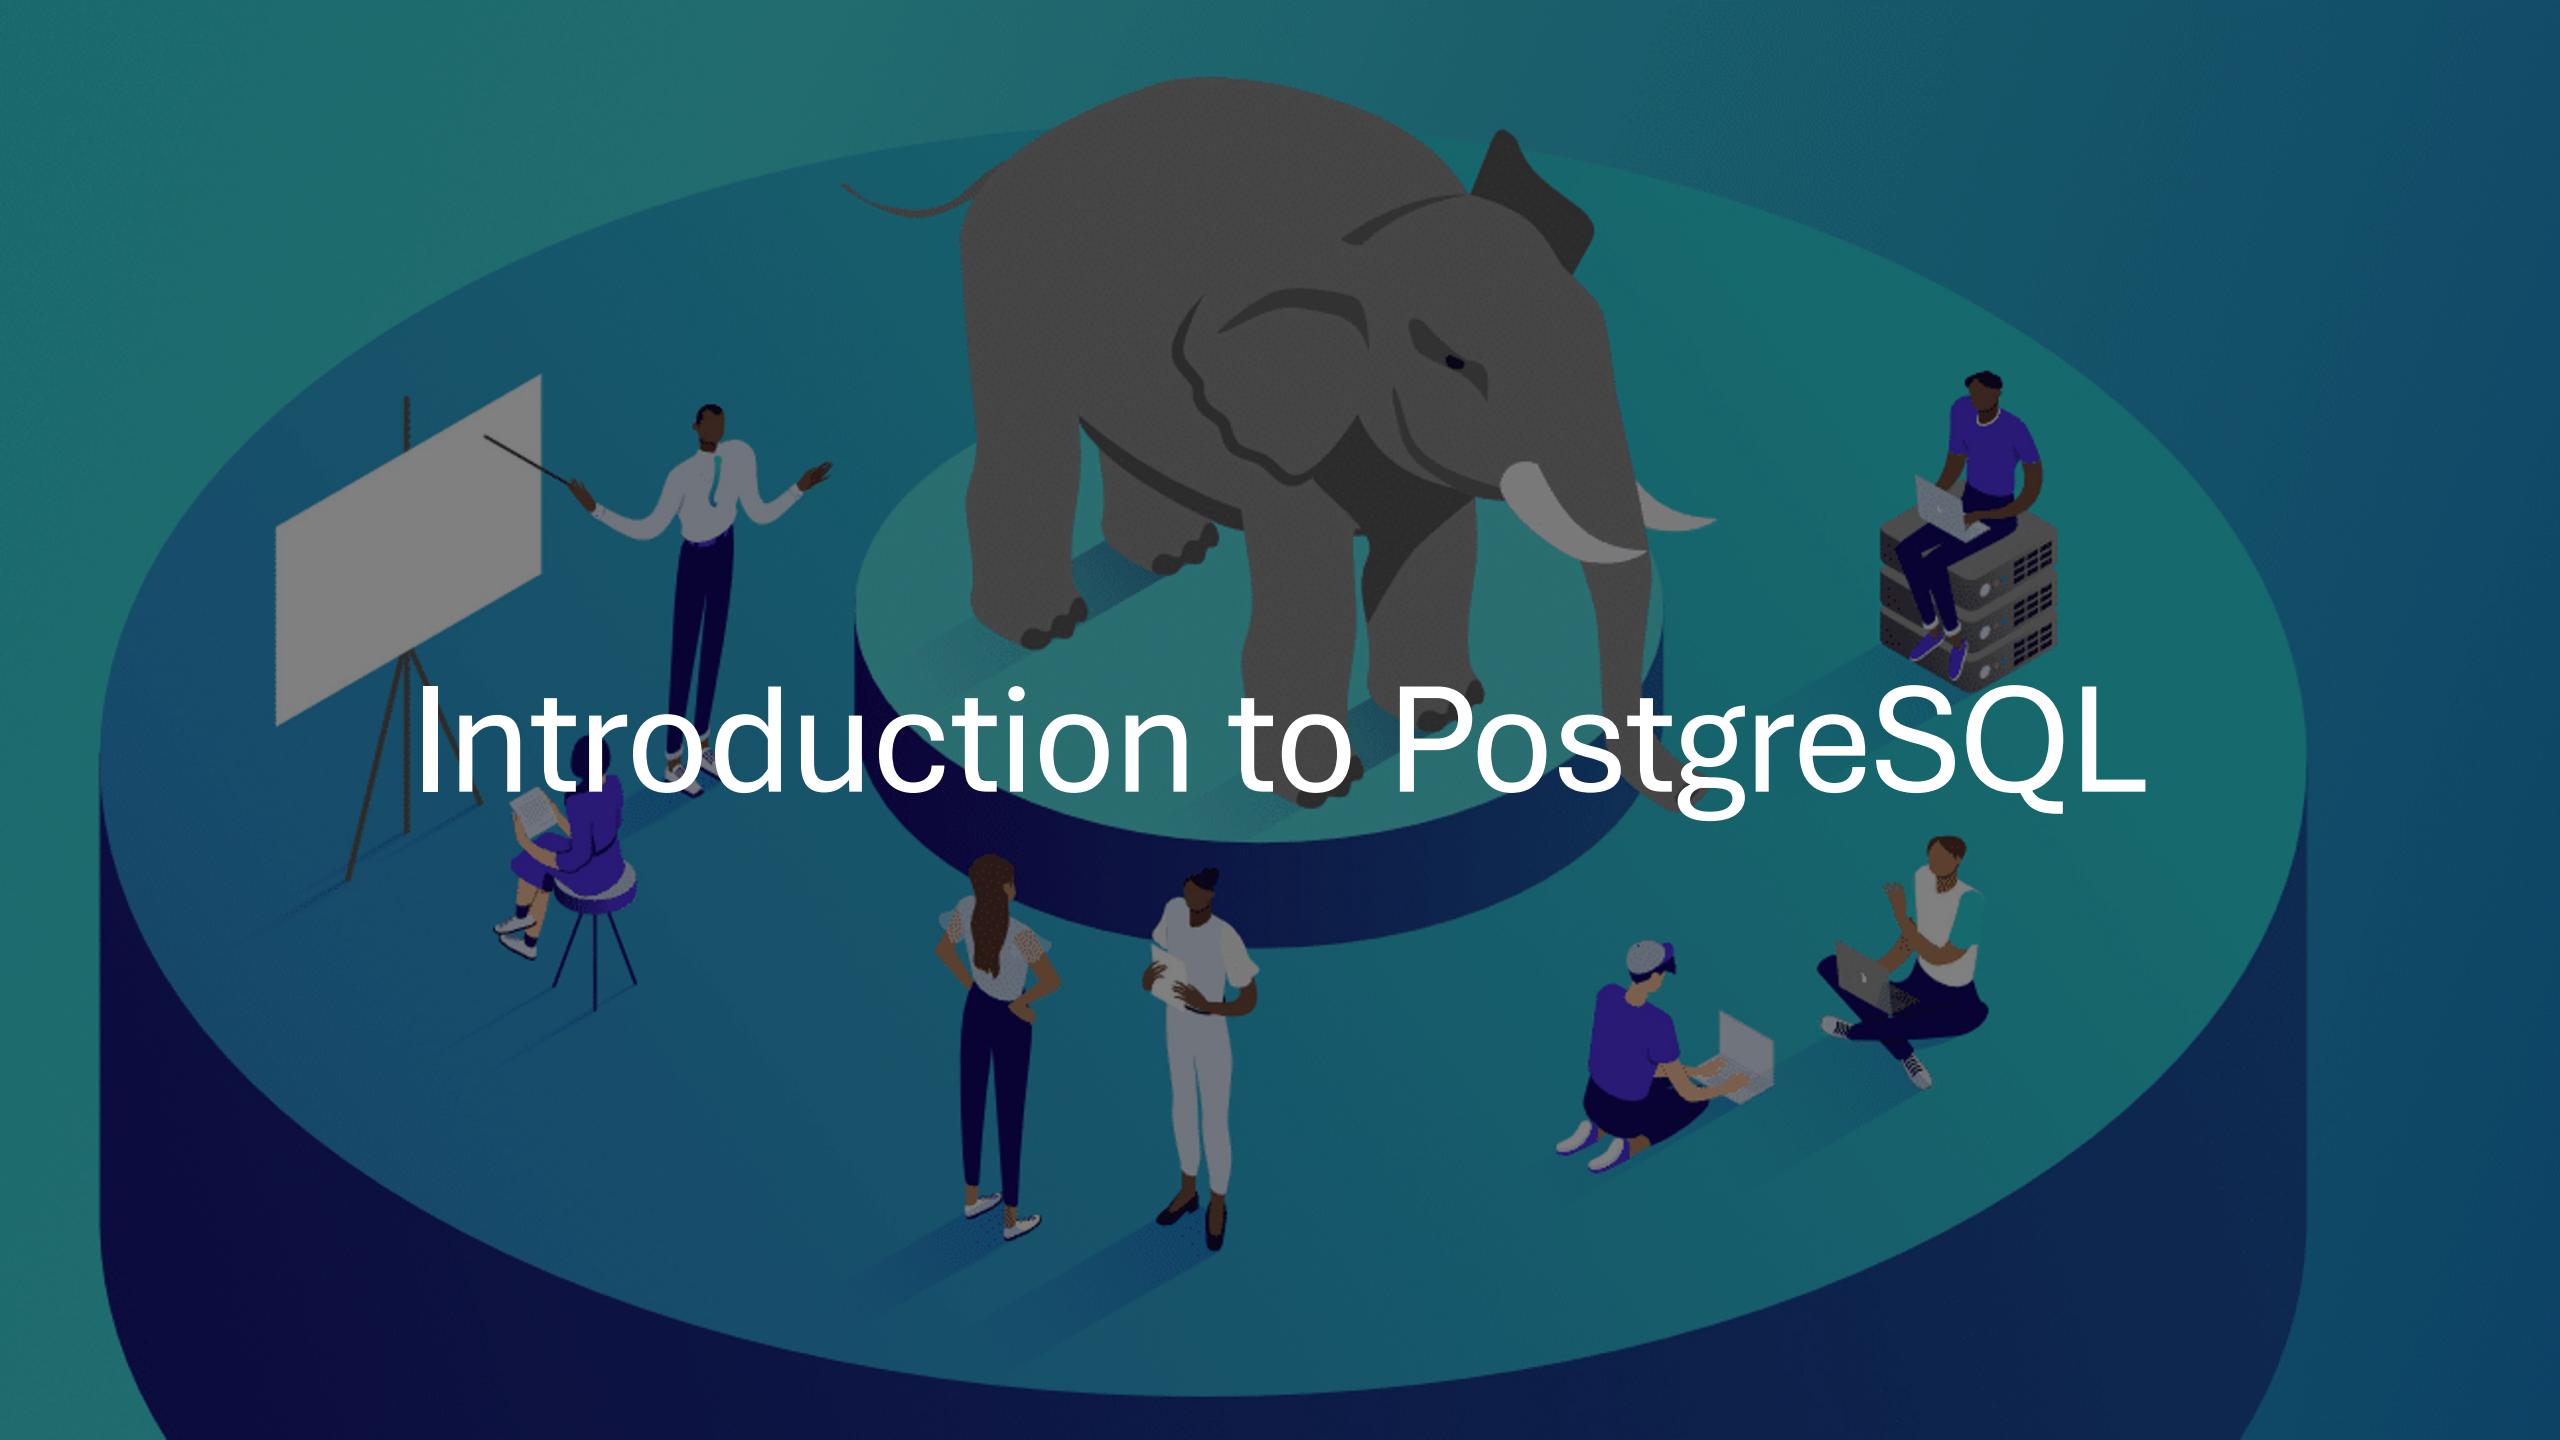
Installation (Windows,  
macOS, Linux)



Configuration basics



Connecting to the  
database



# Introduction to PostgreSQL

# What is PostgreSQL?



PostgreSQL (Postgres)



An advanced open-source



Relational Database Management System (RDBMS).

# What is PostgreSQL?

Known for

Robustness

Extensibility

Standards compliance

# What is PostgreSQL?



Designed to  
handle



Single-machine  
applications



Large-scale data  
warehousing



Web services

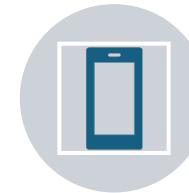
# What is PostgreSQL?



Used for



Web



Mobile



Geospatial



Analytics  
applications

## History of PostgreSQL

1986-  
1985

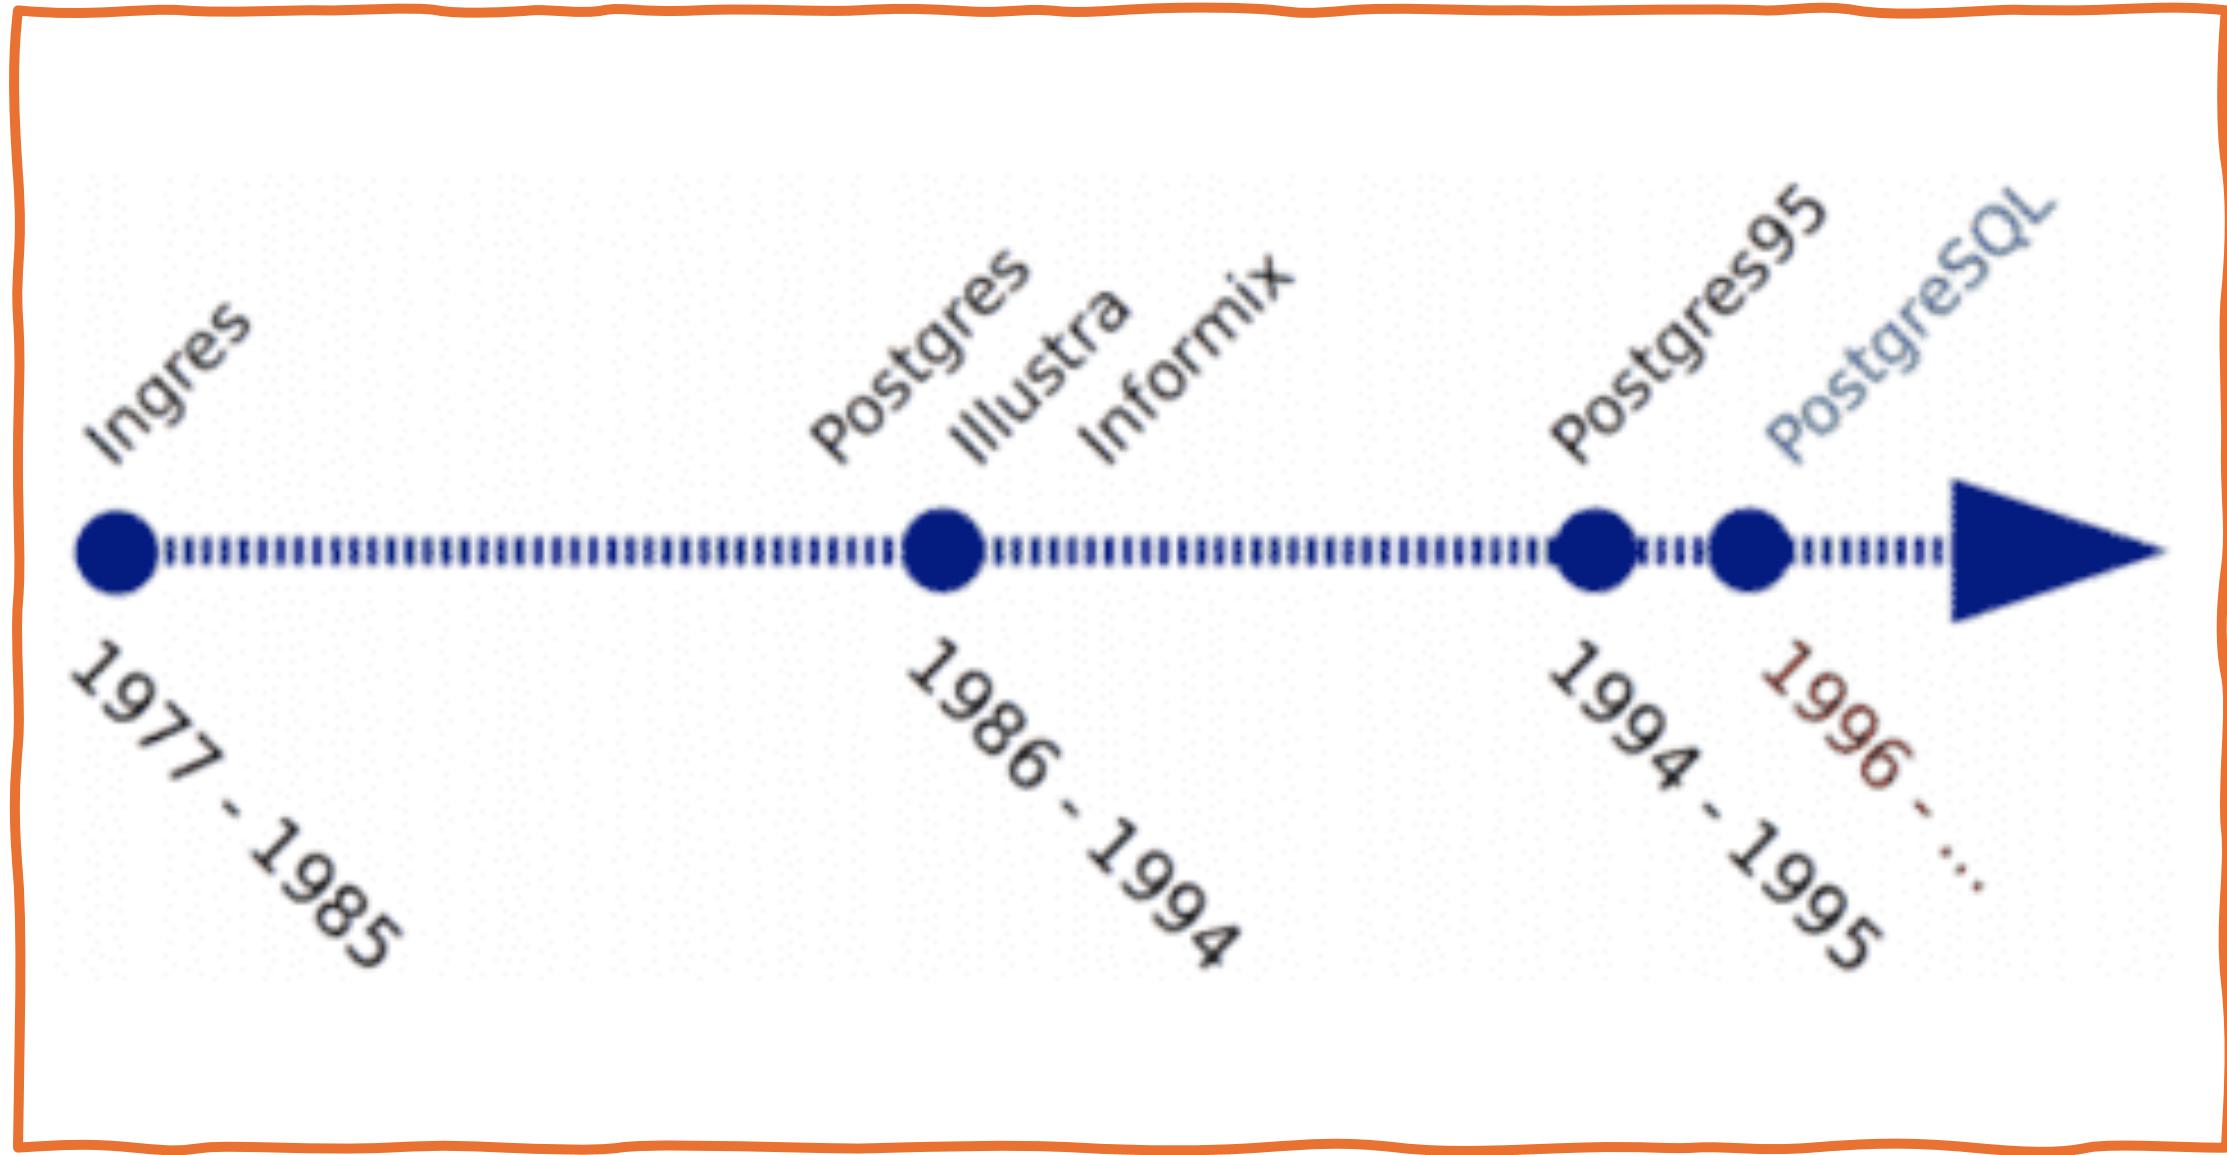
1986-  
1994

1994-  
1995

INGRES

POSTGRES

POSTGRES95



# Origins of PostgreSQL



Ingres project



University of  
California, Berkeley

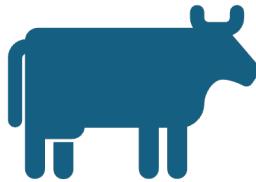


Professor Michael  
Stonebraker

# Origins of PostgreSQL



Initiated in the 1970s



Laid the groundwork



Relational Database  
Systems

# **Postgres Project (1986-1995)**



1986

Michael Stonebraker

Successor

To Ingres database



# 1986

Postgres =  
Post + Ingres

Designed to  
overcome

Limitations  
of Ingres  
system

# 1989

Released first version of Postgres

Introducing Object-Relational databases

Allowing users to

Define types

Functions, and operators

# 1994

---



Postgres95, an enhanced version of Postgres



Replaced the PostQUEL query language with SQL



Standard database query language



Offered improved performance and usability

# **PostgreSQL (1996- Present)**

---

# 1996

Postgres95 officially

Renamed to PostgreSQL

Reflect its SQL compliance

First version PostgreSQL 6.0

Released in January 1997

# 1997

PostgreSQL 6.1 introduced

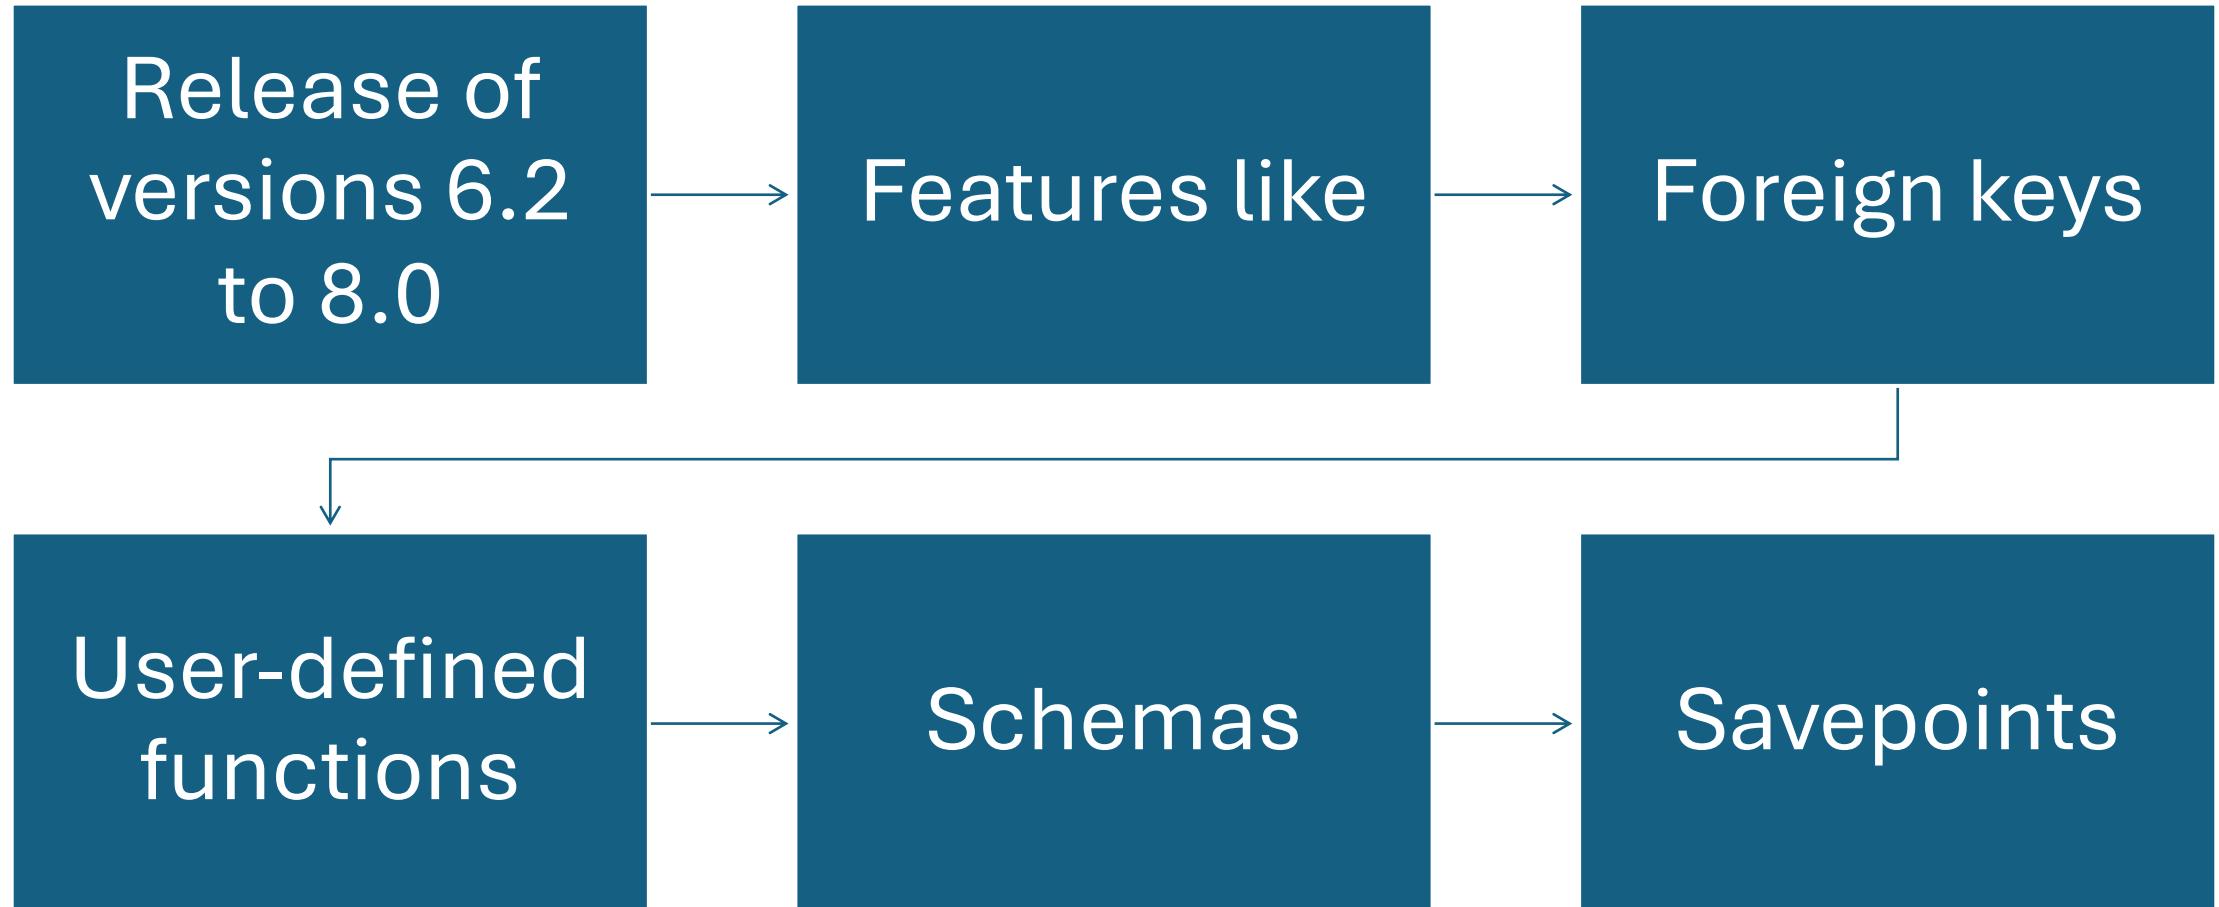
Multi-Version Concurrency Control (MVCC)

Major advancement

Allows for high concurrency

Without significant locking

# 1998-2005



# 1998-2005



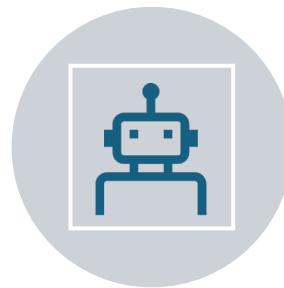
Version 8.0



Native Windows  
support

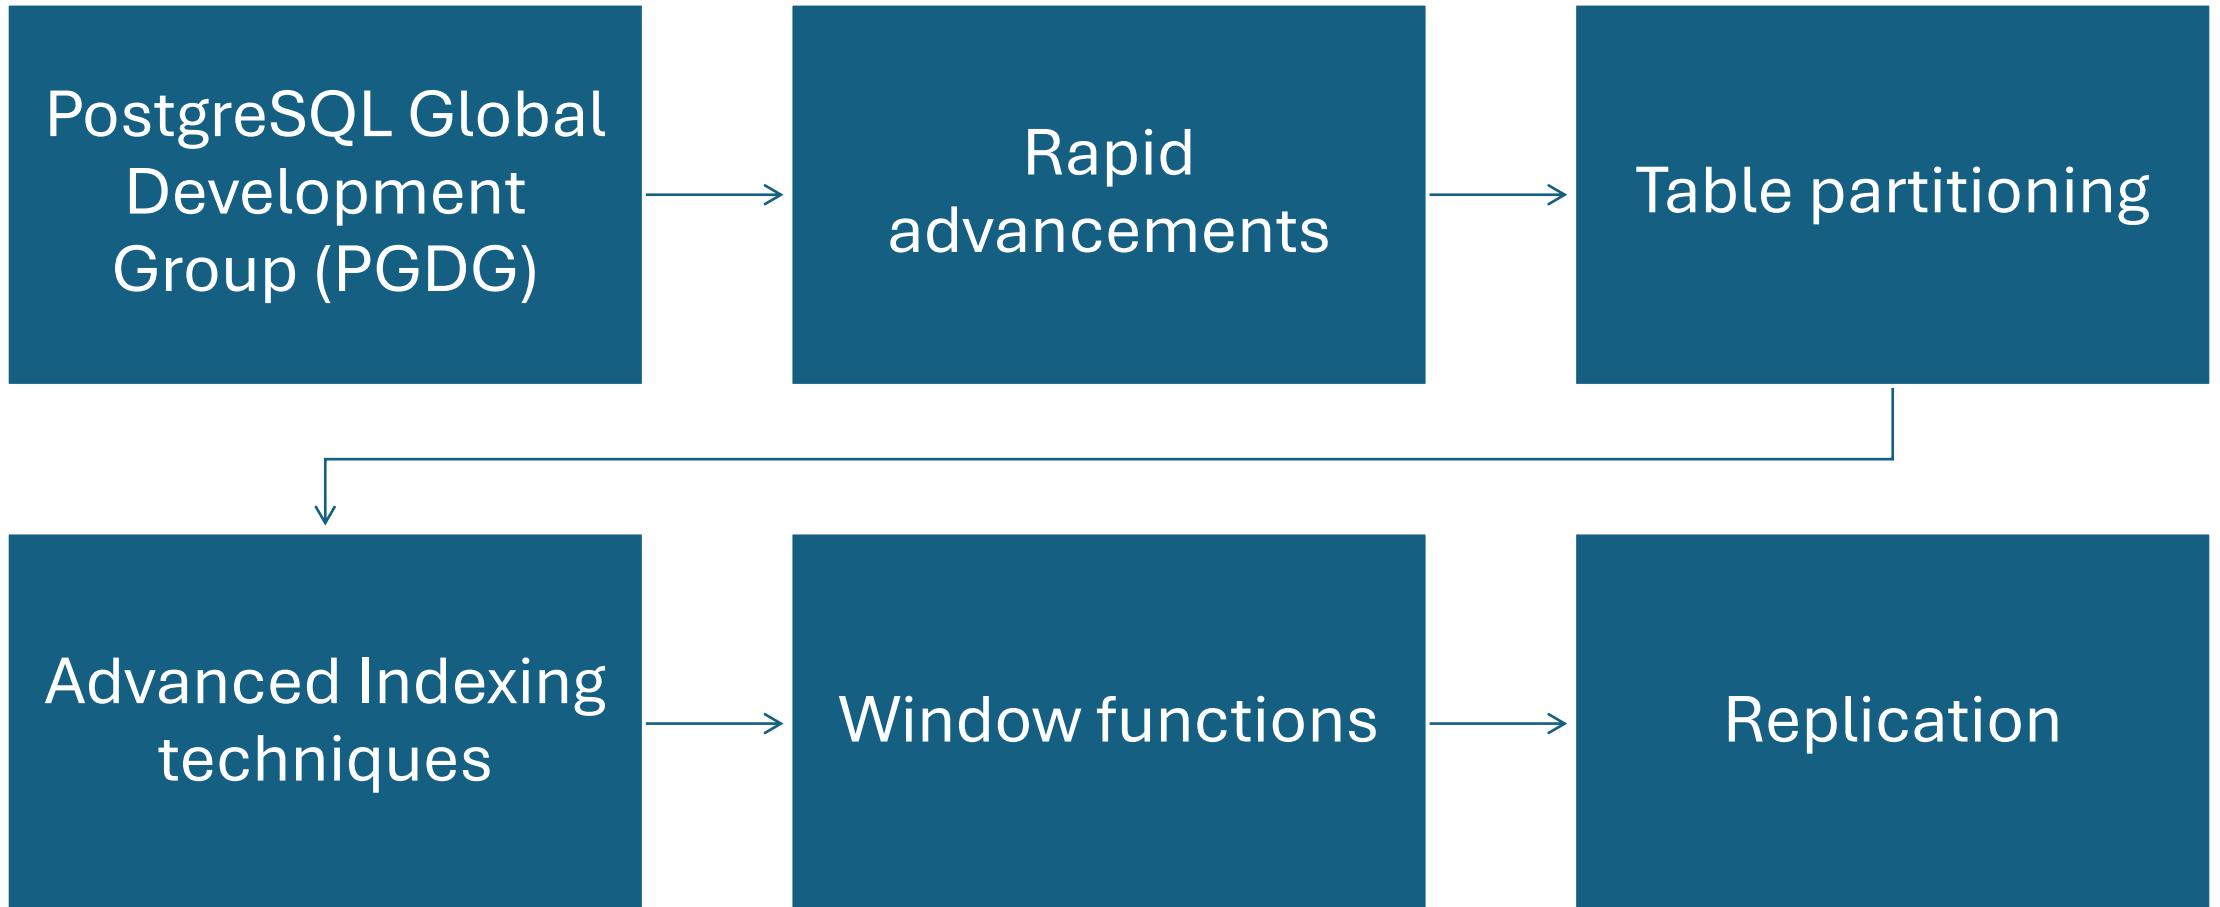


Released in 2005



Expanding its usability

# 2006-2015



# 2006-2015



Version 9.0



Released in  
2010



Streaming  
Replication



Hot standby

# 2016-2020

PostgreSQL  
9.6

Introduced

Parallel  
Query  
execution

# 2016-2020

---



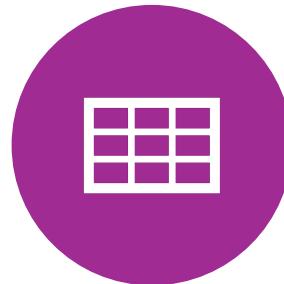
Version 10.0



Logical Replication



Released in 2017



Declarative Table  
Partitioning

# 2016-2020

---



PostgreSQL 11  
to 12



Focused on



Performance  
improvements



Partitioning



SQL  
compliance

# 2021-Present

---



POSTGRESQL  
13 AND 14



ENHANCED  
PARTITIONING



QUERY  
PERFORMANCE



SECURITY  
FEATURES

# 2021-Present



PostgreSQL 15, 16



Continues to build  
on these  
improvements



Enhanced  
Parallelism



Compression  
methods



JSON features

# PostgreSQL 16

---



VARIETY OF  
NEW FEATURES



IMPROVE  
PERFORMANCE



USABILITY



SECURITY

# Performance Improvements

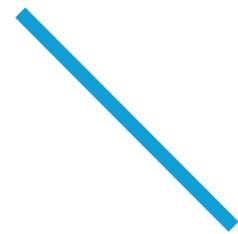
---



Query Performance  
Enhancements



Partitioning  
Improvements



Parallelism  
Enhancements

# Usability Enhancements

---

JSON and  
JSONB  
Enhancements

Enhanced  
Window  
Functions

Improved psql  
Command-  
Line Tool

# Security Enhancements



Enhanced  
Authentication and  
Authorization



Advanced Data  
Encryption

# Administrative Features



Enhanced Backup and Restore



Improved Monitoring and Logging

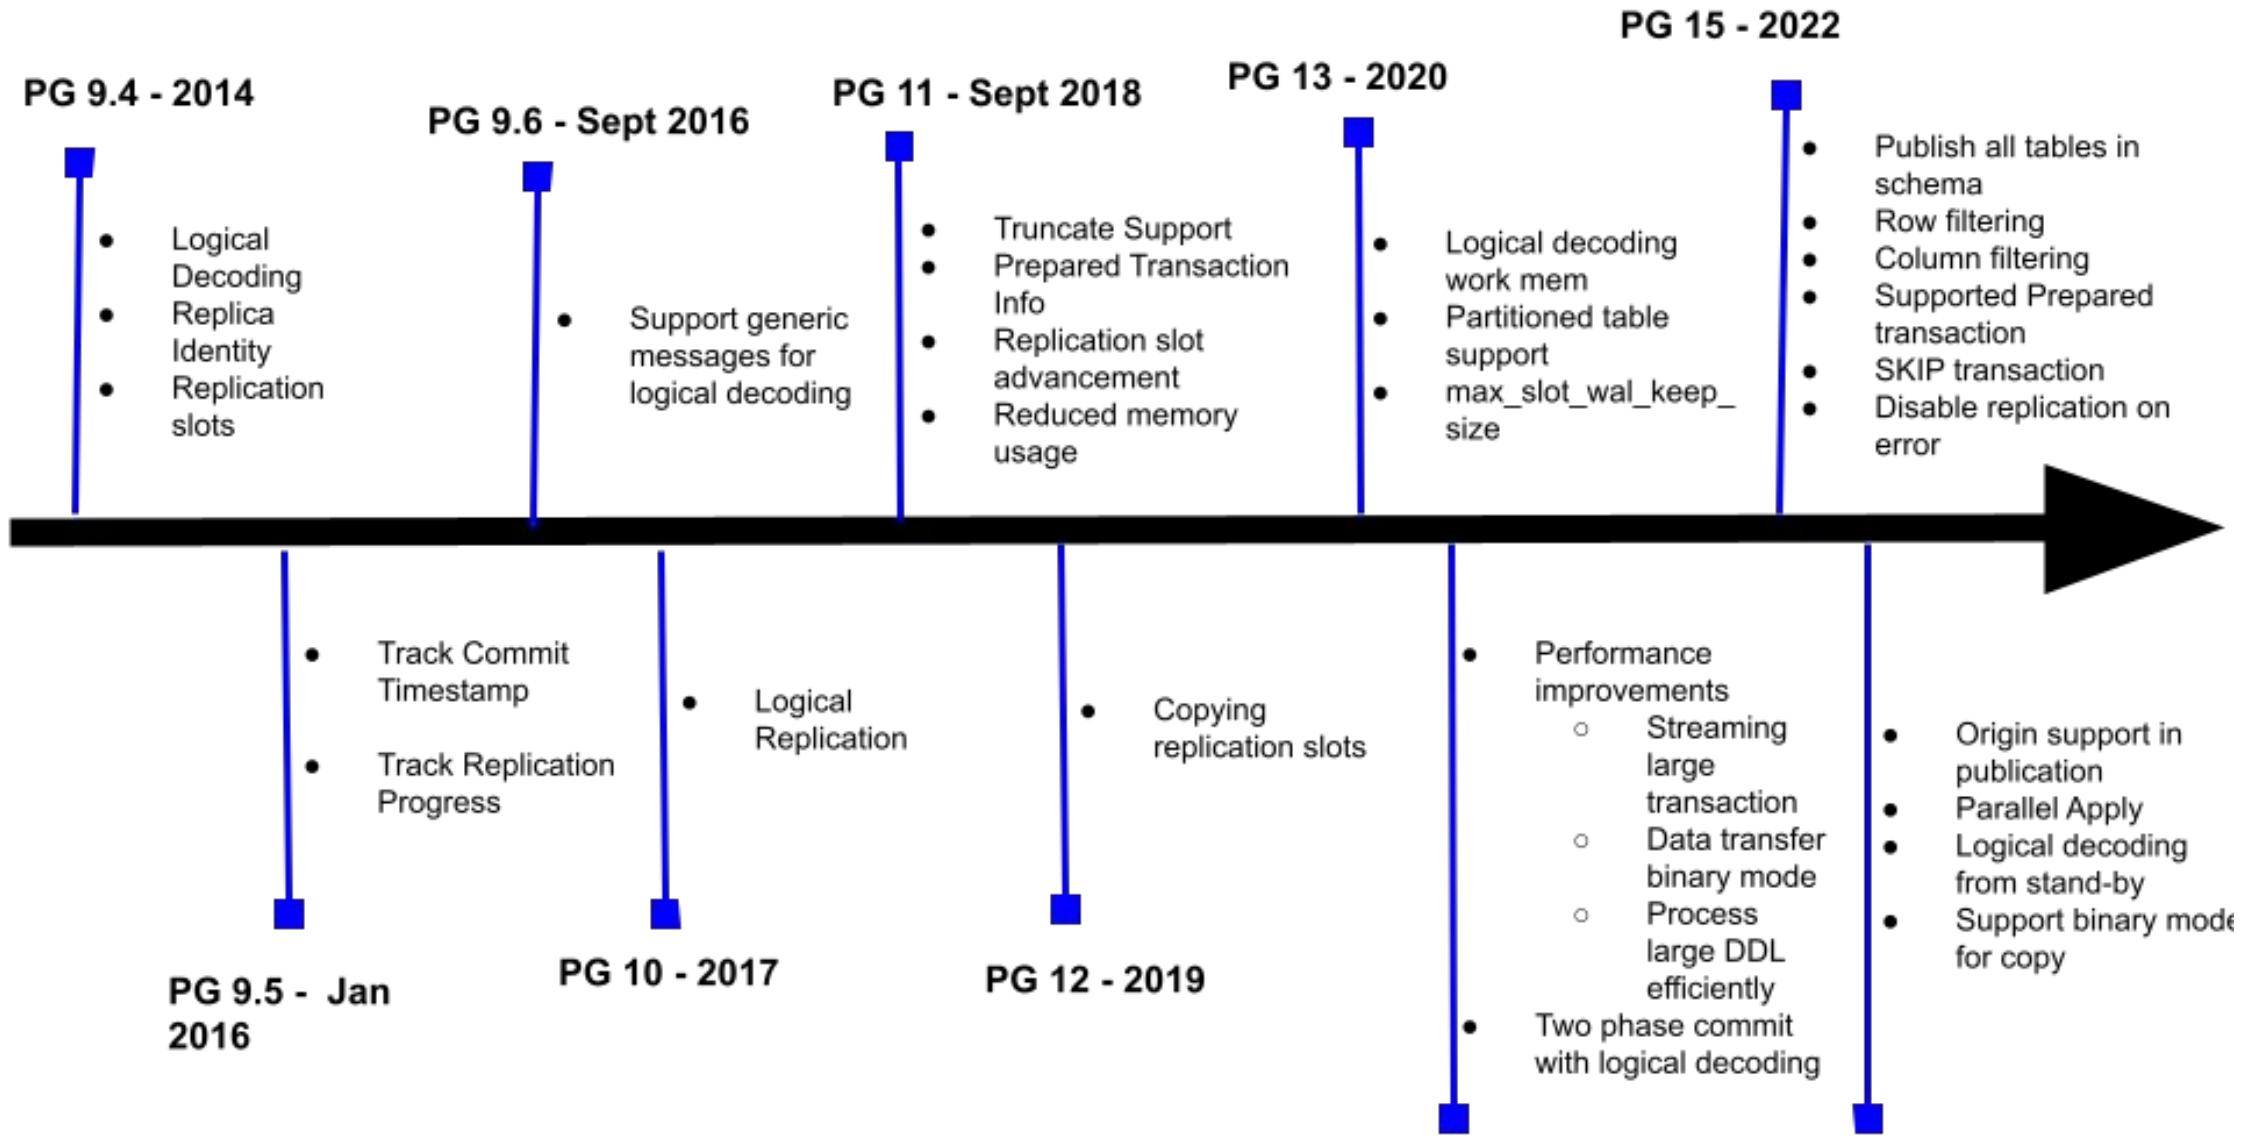
# Developer Features



Enhanced PL/pgSQL



Advanced Indexing





## Key Contributors

- **Michael Stonebraker**
- The primary Architect
- Original Postgres project
- Key figure in database research

# PostgreSQL Global Development Group (PGDG)



A group of  
volunteers and  
Companies



Oversee the  
development



Maintenance of  
PostgreSQL

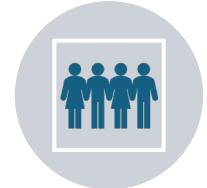
# Community



Community-driven



Contributions from  
developers  
worldwide



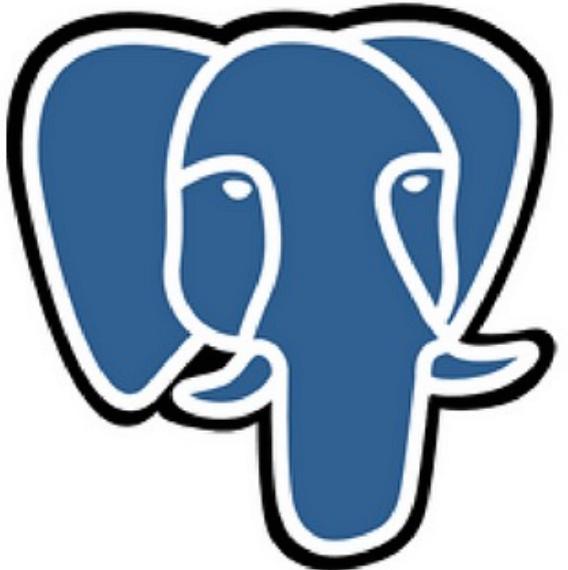
Known for its  
collaborative spirit



Extensive  
documentation



Active mailing lists



PostgreSQL

**Design Philosophy**

# Standards Compliance



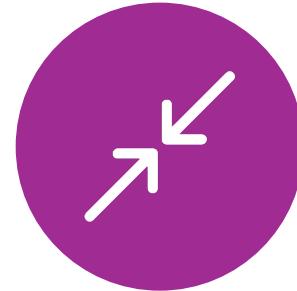
PostgreSQL adheres



To SQL standards



Ensuring compatibility



Ease of use

# Extensibility

Designed to  
be highly  
extensible

Allowing  
users to  
define

New data  
types

Functions

Operators

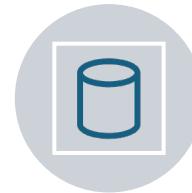
# Reliability and Stability



Focuses on delivering



Stable



Reliable database system



Suitable for



Mission-critical applications

# Open Source

Released under  
the PostgreSQL  
License

An open-source  
license

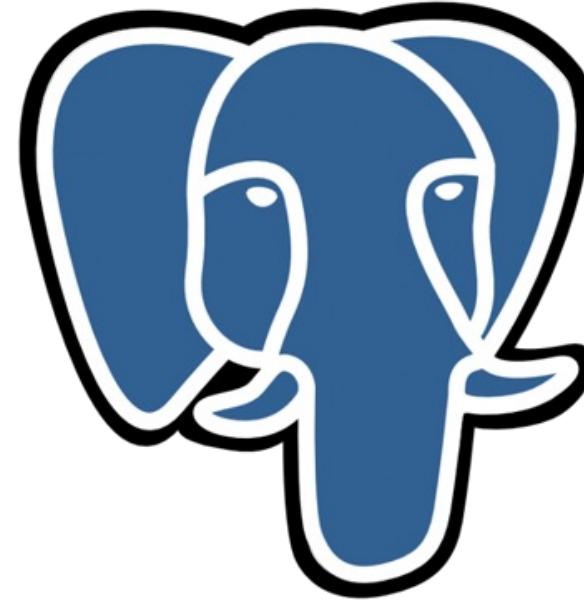
Allows for free  
use

Modification

Distribution



# Major Milestones



PostgreSQL

# MVCC (1997)

Multi-Version Concurrency Control

A crucial feature

for high Concurrency

Performance

# SQL Compliance (1996)



Transition from



PostQUEL to SQL



Aligning with



Industry  
standards

# Windows Support (2005)

Native support for

Windows

Broadening

PostgreSQL's user base

# Streaming Replication (2010)



Hot standby



Improving



Data



Redundancy



Availability

# Parallel Query Execution (2016)

Support for

Parallel  
Query  
Execution

Enhancing  
performance

for large  
queries.

# Logical Replication (2017)



Introduction of



Logical  
replication,



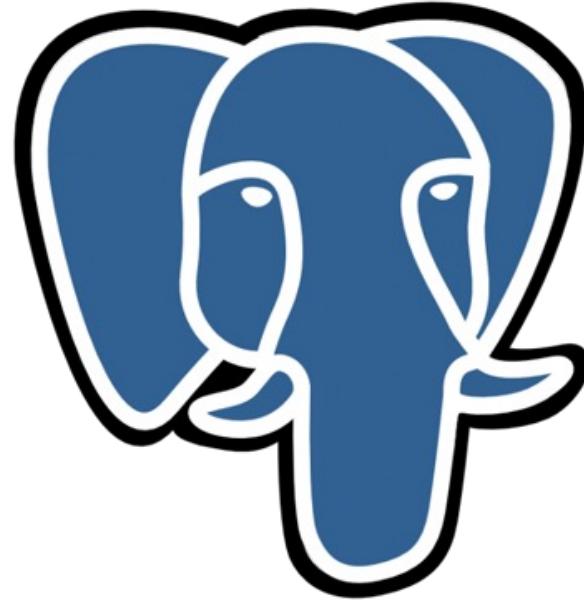
Allowing  
selective



Replication of  
data

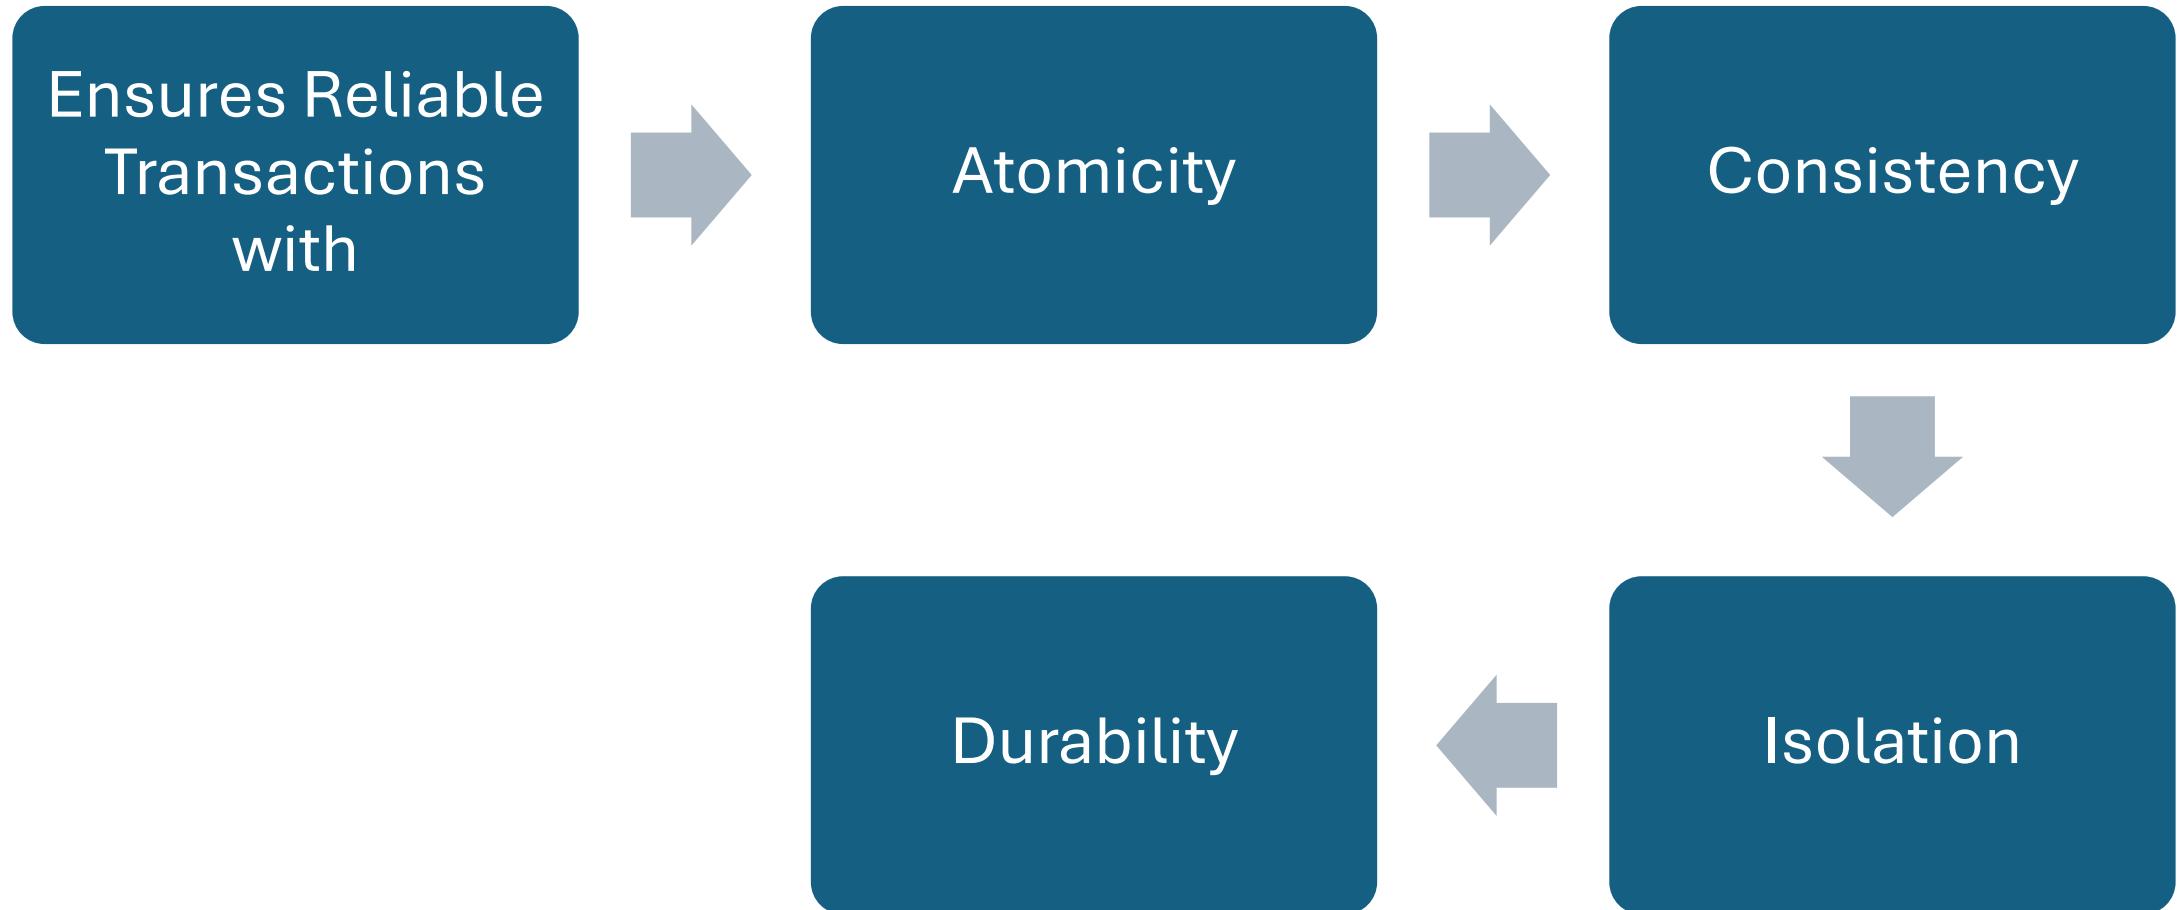


## Key features and Benefits



# Postgre<sup>SQL</sup>

# ACID Compliance



# Atomicity

Each transaction is treated

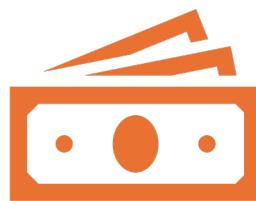
Single "Atomic" Unit

Either completely succeeds or

Completely fails.

# Atomicity

---



If any part of the transaction fails



Entire transaction is rolled back



Database remains unchanged.

# Example

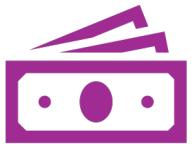
If a transaction involves

transferring money from

Account A to Account B.

# Example

---



Either both the  
debit from



Account A and



The credit to  
Account B



Happen or  
neither happens.

# Example

If the debit succeeds

but the credit fails,

the transaction is rolled back.

Account A's balance is restored

to its original state

# ACID Compliance



Provides robust  
support



For concurrent  
transactions



Without compromising  
data integrity

# Consistency

Ensures that a transaction

From one valid state

To another valid state,

Maintain all predefined rules

Constraints, and triggers

# Consistency

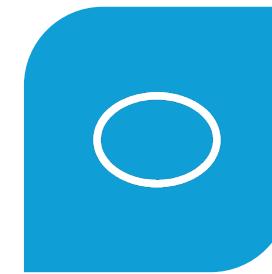
---



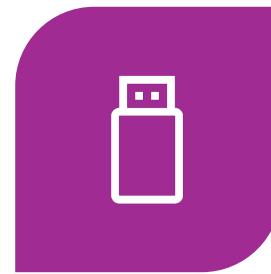
DATABASE  
MUST



REMAIN  
CONSISTENT

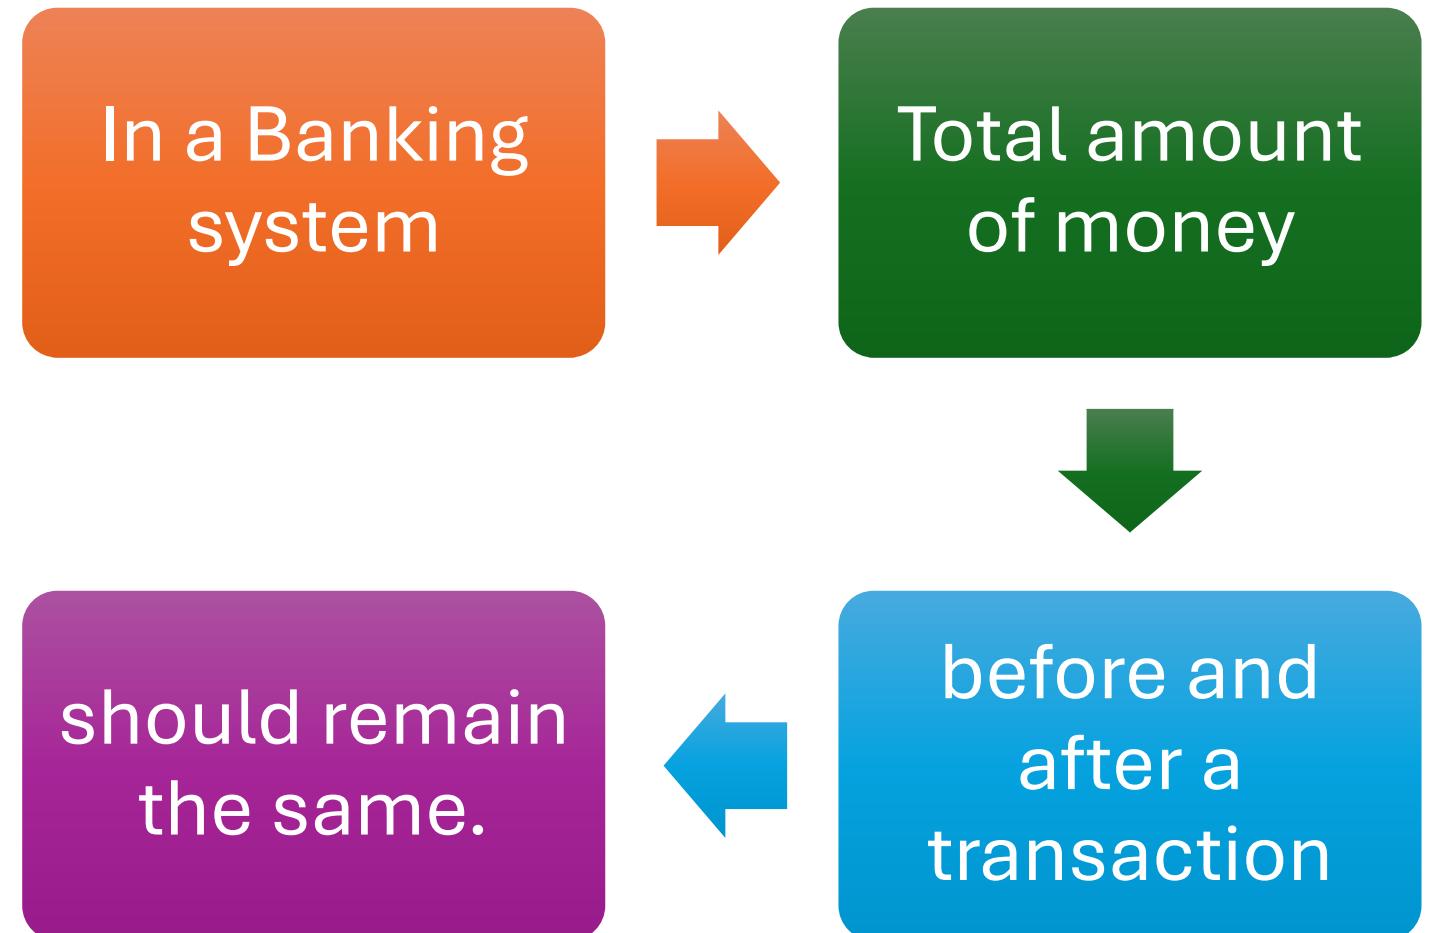


BEFORE AND  
AFTER



TRANSACTION

# Consistency



# Isolation

Transactions

Executed in

Isolation

From One another

# Isolation

If two transactions

Running Concurrently

One transferring money from

Account A to Account B

Another from Account C to Account D

# Isolation

Isolation ensures

Transactions

Do not see each other's

Intermediate states

# Durability

Once a transaction

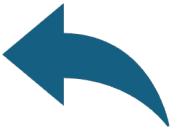
Has been committed

Will remain

Even in the event of

a system failure

# Durability



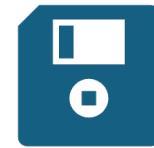
If a transaction  
to update



A bank balance  
is committed,



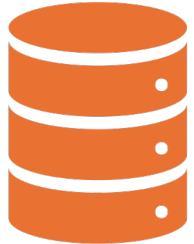
New balance is  
guaranteed



to be saved in  
the database.

# Advanced Data Types

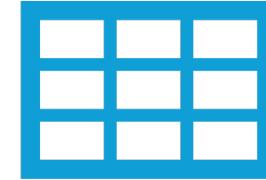
# JSON/JSONB



Efficient storage



Querying of



JSON data.

# Arrays



Support for

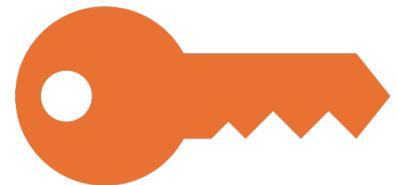


Multi-dimensional



Arrays

# Hstore



Key-value pairs

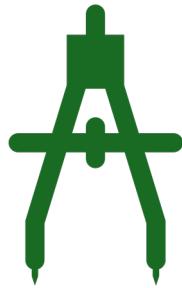


Storage

# Geometric Types



Points,

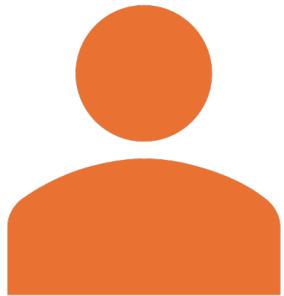


lines

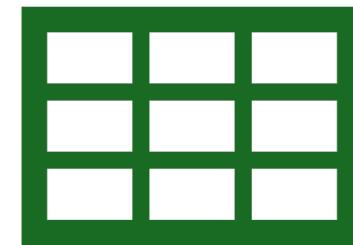


polygons

# Custom Types



Users can define



Own data types.

# Extensibility



Users can  
create



Custom  
functions



Operators,

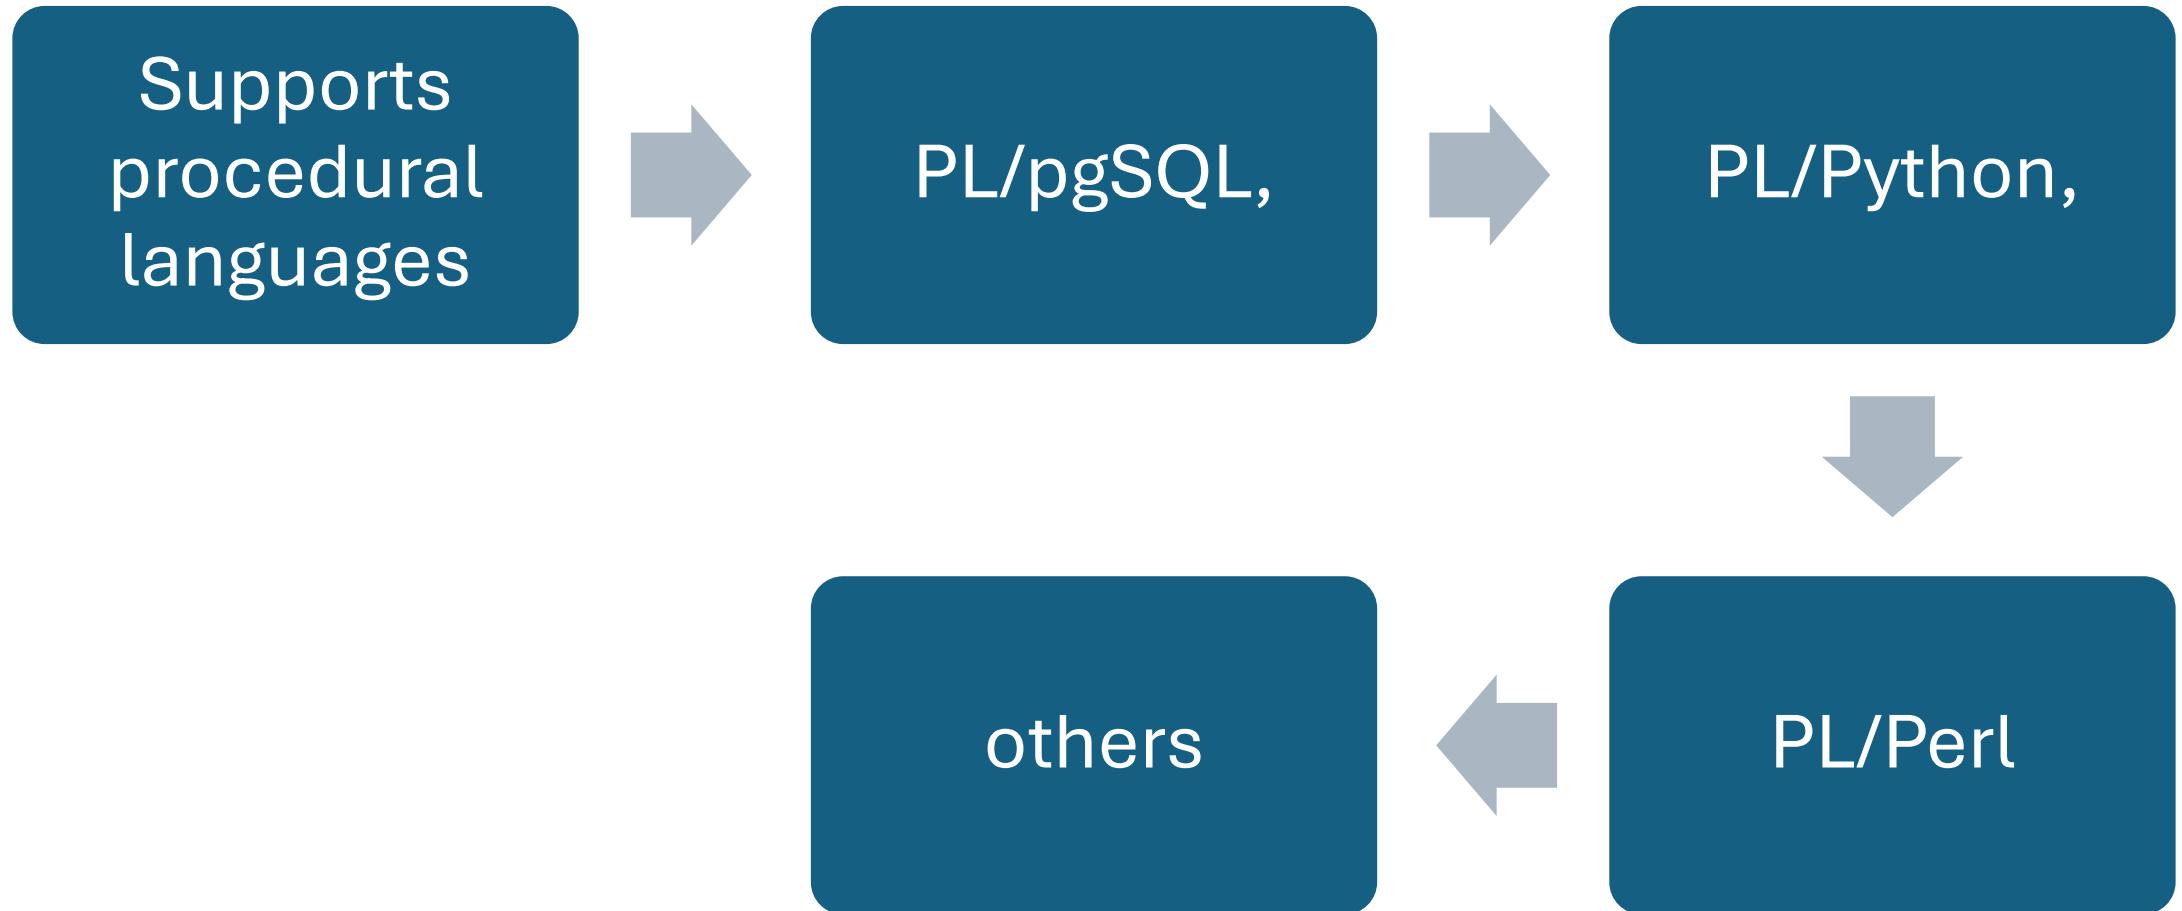


Data types



index types.

# Extensibility



# Extensibility



Extensions like



PostGIS (for  
geospatial data)

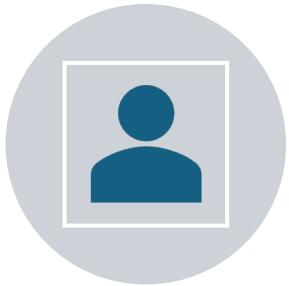


Full-text search



Easily integrated

# Multi-Version Concurrency Control (MVCC)



Allows multiple transactions



to occur simultaneously without locking.



Ensures data consistency



High performance

# Indexing



Supports various  
index types



B-tree,



Hash,



GiST,

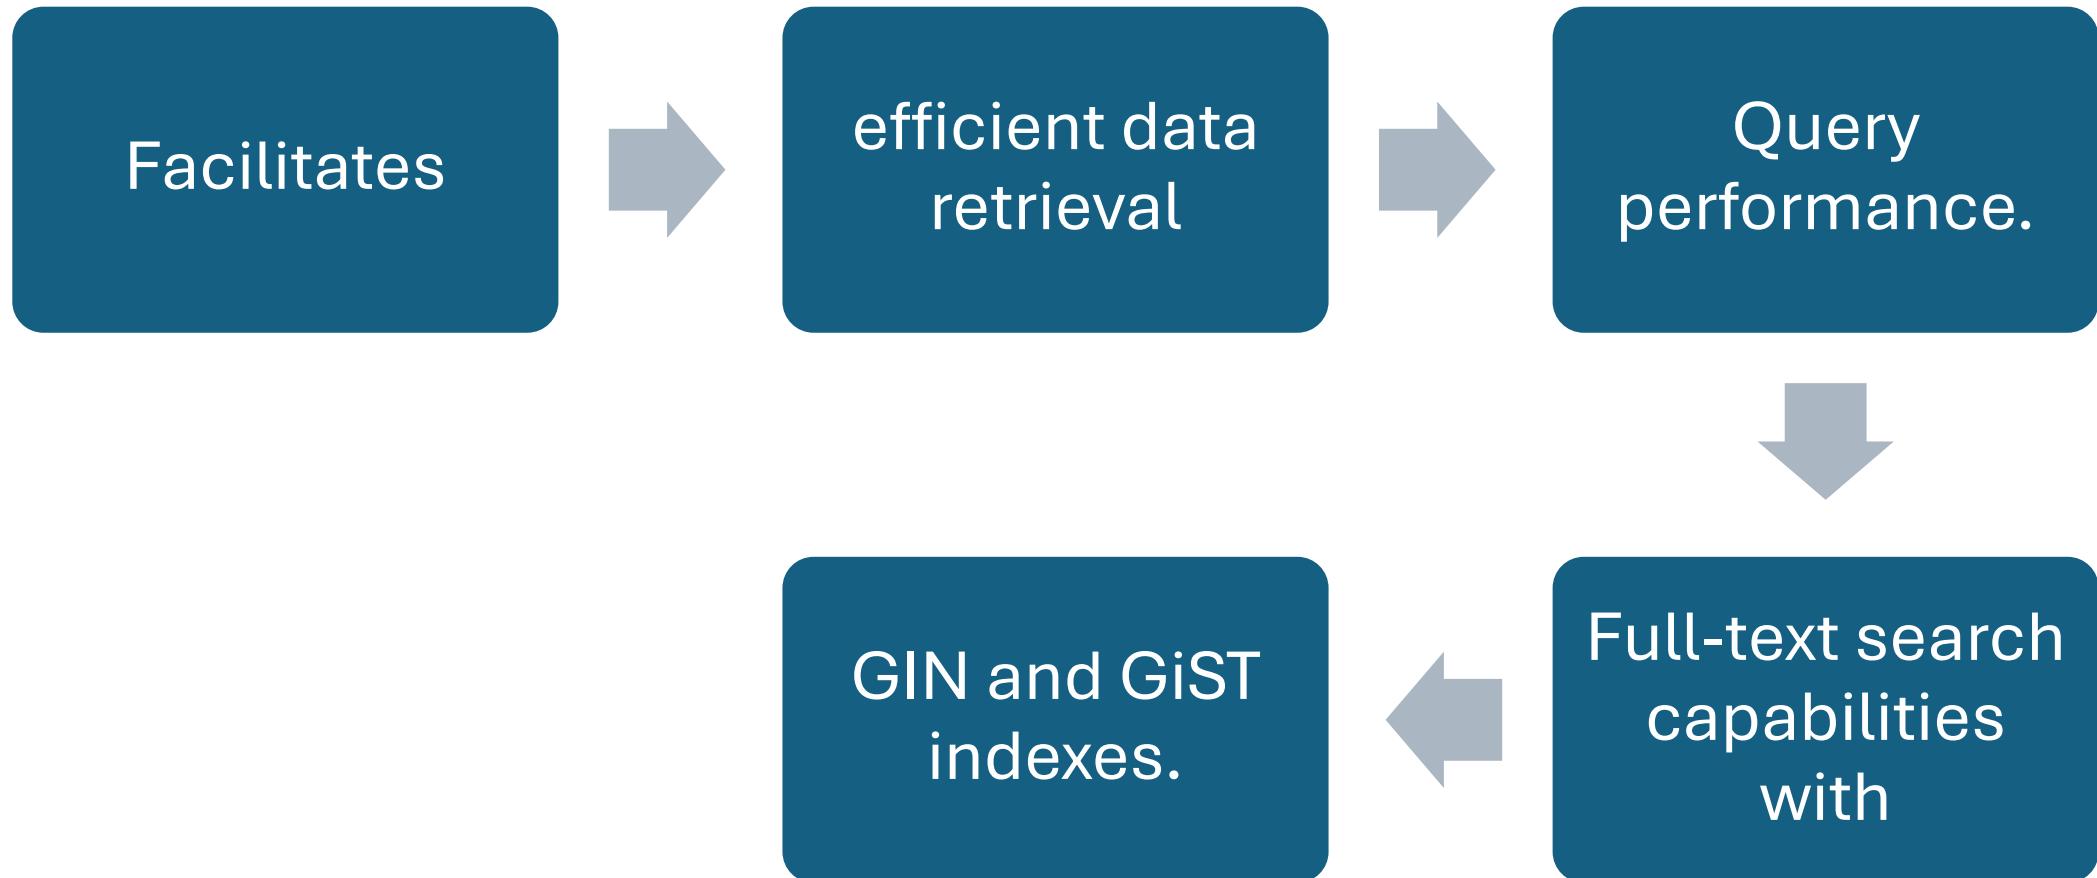


GIN



BRIN

# Indexing



# Foreign Data Wrappers (FDW)

Allows integration with

other databases and data sources.

Enables querying external data

as if it were a part of

the PostgreSQL database.

# Full-Text Search



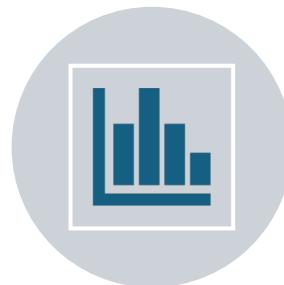
Built-in support for



full-text search  
capabilities.



Indexing and querying  
of



text data for fast search  
results

# Replication and High Availability

Supports both

Synchronous

Asynchronous

Replication

# Replication and High Availability

---



STREAMING  
REPLICATION



REAL-TIME DATA  
REDUNDANCY



LOGICAL  
REPLICATION



SELECTIVE DATA  
SYNCHRONIZATION

# Partitioning

---



Declarative  
partitioning



for efficient  
management



of large tables.

# Partitioning

Supports

Range,

List,

Hash

composite partitioning methods.

# Security



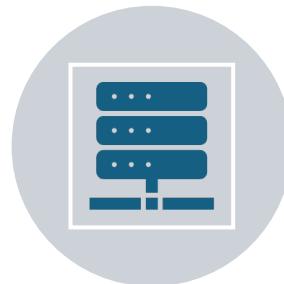
Robust access control  
mechanisms



with roles and  
permissions.



SSL support for



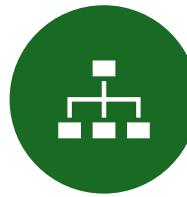
secure client-server  
communication.

# Security

---



Advanced  
features like



row-level  
security and



data  
encryption.

# Performance Optimization

Parallel Query  
execution

for faster  
processing

of large  
datasets.

# Performance Optimization

Query planner

Optimizer

for efficient

Query execution

# Performance Optimization



CACHING MECHANISMS



TO IMPROVE  
PERFORMANCE.

# Backup and Recovery

Comprehensive backup and

restore options with

tools like pg\_dump and

pg\_basebackup.

# Backup and Recovery

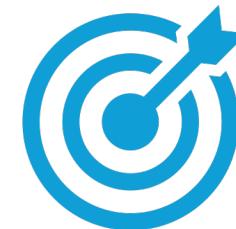
---



Point-in-time  
recovery (PITR)



for restoring  
databases



to a specific state.

# **Benefits of PostgreSQL**

---

# Open Source and Free

PostgreSQL is open source,  
allowing for free usage,  
modification, and distribution.  
No licensing costs,  
making it cost-effective  
for organizations

# Flexibility

Highly extensible and  
adaptable to various use cases.  
Suitable for a wide range of applications,  
from small-scale to  
large-scale enterprise systems.

# Community and Support

---



Strong, active community



providing extensive documentation,



support, and regular updates.



Numerous third-party tools and



extensions available.

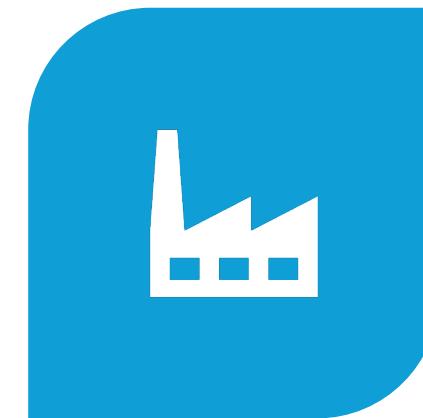
# Reliability and Stability



PROVEN TRACK  
RECORD OF



RELIABILITY AND  
STABILITY IN



PRODUCTION  
ENVIRONMENTS.

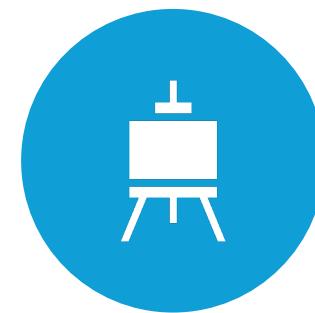
# Standards Compliance



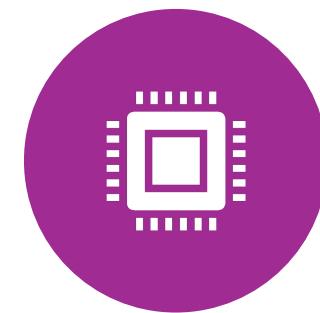
ADHERES TO SQL  
STANDARDS,



ENSURING  
COMPATIBILITY AND



EASE OF USE FOR  
DEVELOPERS



FAMILIAR WITH  
SQL.

# Standards Compliance



REGULAR UPDATES



TO MAINTAIN  
COMPLIANCE



WITH EVOLVING  
STANDARDS.

# Cross-Platform Compatibility

---



Runs on various operating systems



including Windows, macOS, and Linux.



Supports integration with



Various programming



languages and frameworks.

# Performance



Advanced indexing  
and



query optimization  
techniques



ensure high  
performance.



Scalability to  
handle



large volumes of  
data and



high transaction  
loads.

# Data Integrity

Strong support for data integrity

through constraints,

triggers, and foreign keys.

Ensures accurate and

consistent data storage and

retrieval.

# Geospatial Capabilities

PostGIS extension

provides robust support

for geographic information systems (GIS).

Ideal for applications

requiring spatial data management.

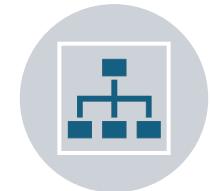
# Enterprise Features



Comprehensive  
feature set



that meets the  
needs of



enterprise-level  
applications.



Support for  
complex queries,



transactions



data analytics.

# Oracle vs PostgreSQL

---

# Overview

## Oracle

- A commercial, enterprise-level relational database management system.
- Developed by Oracle Corporation.
- Known for high performance, advanced features, and extensive support for complex, high-volume applications.
- Typically used in large-scale, mission-critical environments.

# Overview

## PostgreSQL

- An open-source relational database management system.
- Developed and maintained by a global community of developers.
- Known for its robustness, extensibility, and adherence to SQL standards.
- Suitable for a wide range of applications, from small projects to large enterprise systems.

# Licensing and Cost

## Oracle:

Proprietary software with licensing fees that can be quite high.

Offers various editions (Standard, Enterprise, etc.) with different feature sets and price points.

Licensing can include additional costs for support, features, and scalability options.

# Licensing and Cost

## PostgreSQL:

Open-source and free to use under the PostgreSQL License.

No licensing fees, making it cost-effective.

Community support is free; commercial support options are available but not mandatory.

# Features

## Oracle:

- Advanced features like Real Application Clusters (RAC), Data Guard, and Flashback Technology.
- Strong focus on security with features like Transparent Data Encryption (TDE) and Label Security.
- Comprehensive data warehousing and OLAP capabilities.
- Advanced indexing and partitioning options.
- Built-in support for PL/SQL for complex stored procedures and triggers.

# Features

## PostgreSQL

- Extensible with support for custom data types, operators, and functions.
- Advanced data types including JSON/JSONB, arrays, and hstore.
- Multi-Version Concurrency Control (MVCC) for high concurrency.
- Support for various procedural languages (PL/pgSQL, PL/Python, etc.).
- Extensions like PostGIS for geospatial data and full-text search.

# Performance and Scalability

- **Oracle**
- Known for high performance and scalability in large, complex environments.
- Supports clustering, parallel execution, and advanced optimization techniques.
- Can handle very large databases and high transaction volumes with ease.
- In-memory database capabilities for performance enhancement.

# Performance and Scalability

## PostgreSQL:

- Highly performant and scalable, particularly with recent improvements in parallel query execution and indexing.
- Suitable for medium to large-scale applications.
- Support for partitioning, sharding, and replication for scalability.
- Performance can be further enhanced through tuning and optimization

# Security

## Oracle:

Extensive security features including Advanced Security Option (ASO), Virtual Private Database (VPD), and fine-grained auditing.

Robust user and role management with detailed access controls.

# Security

## PostgreSQL:

Strong security features including roles and permissions, SSL support, and row-level security.

Supports data encryption, though some advanced features may require additional configuration or extensions.

# Support and Community

## Oracle:

Comprehensive commercial support options from Oracle Corporation.  
Extensive documentation, training, and certification programs.  
Large user community, but primarily focused on enterprise customers.

# Support and Community

## **PostgreSQL:**

Strong community support with active mailing lists, forums, and extensive documentation.

Commercial support options available through various companies.

Large and diverse user base ranging from small startups to large enterprises.

# Ease of Use and Management

- . **Oracle:**
- . Advanced management tools like Oracle Enterprise Manager for database administration.
- . Can be complex to set up and manage, particularly for advanced features.

# Ease of Use and Management

- **PostgreSQL:**
- User-friendly with tools like pgAdmin and a variety of third-party management tools.
- Generally easier to set up and manage, with a straightforward configuration.

# Integration and Compatibility

- **Oracle:**
  - Integrates well with other Oracle products and many third-party applications.
  - Extensive support for various programming languages and platforms.

# Integration and Compatibility

## PostgreSQL:

- Highly compatible with many programming languages, frameworks, and platforms.
- Foreign Data Wrappers (FDWs) for integrating with other databases and data sources.

# Conclusion

- **Oracle** and **PostgreSQL** are both powerful database management systems, but they cater to different needs and budgets

# Conclusion

**Oracle** is ideal for large enterprises requiring high performance, advanced features, and comprehensive support for mission-critical applications, albeit with higher costs.

# Conclusion

**PostgreSQL** offers a robust, cost-effective solution suitable for a wide range of applications, from small projects to large-scale enterprise systems, with strong community support and extensibility.

# Oracle vs PostgreSQL Queries

# Oracle

```
CREATE TABLE employees ( employee_id NUMBER GENERATED BY  
DEFAULT AS IDENTITY, first_name VARCHAR2(50), last_name  
VARCHAR2(50), email VARCHAR2(100), hire_date DATE, salary  
NUMBER(8, 2), PRIMARY KEY (employee_id));
```

# PostgreSQL

```
CREATE TABLE employees ( employee_id SERIAL PRIMARY KEY,  
first_name VARCHAR(50), last_name VARCHAR(50), email  
VARCHAR(100), hire_date DATE, salary NUMERIC(8, 2));
```

# *Inserting Data*

## **Oracle:**

```
INSERT INTO employees (first_name, last_name, email, hire_date,  
salary)    VALUES    ('John',    'Doe',    'john.doe@example.com',  
TO_DATE('2023-07-01', 'YYYY-MM-DD'), 50000);
```

# *Inserting Data*

## **PostgreSQL**

```
INSERT INTO employees (first_name, last_name, email, hire_date, salary)
VALUES ('John', 'Doe', 'john.doe@example.com', '2023-07-01', 50000);
```

# *Selecting Data*

## **Oracle:**

```
SELECT employee_id, first_name, last_name, email FROM  
employees WHERE salary > 40000;
```

## **PostgreSQL:**

```
SELECT employee_id, first_name, last_name, email FROM  
employees WHERE salary > 40000;
```

# *Updating Data*

## **Oracle:**

```
UPDATE employees SET salary = salary * 1.10 WHERE hire_date <  
TO_DATE('2020-01-01', 'YYYY-MM-DD');
```

## **PostgreSQL:**

```
UPDATE employees SET salary = salary * 1.10 WHERE hire_date <  
'2020-01-01';
```

# *Deleting Data*

## **Oracle:**

```
DELETE FROM employees WHERE last_name = 'Doe';
```

## **PostgreSQL:**

- `DELETE FROM employees WHERE last_name = 'Doe';`

# Functions and Procedural Languages

- *String Functions*

## Oracle:

```
SELECT CONCAT(first_name, ' ', last_name) AS full_name FROM employees;
```

## PostgreSQL:

- ```
SELECT first_name || ' ' || last_name AS full_name FROM employees;
```

# *Date Functions*

**Oracle:**

```
SELECT SYSDATE FROM dual;
```

**PostgreSQL:**

```
SELECT CURRENT_DATE;
```

# *Conditional Logic*

## **Oracle:**

```
SELECT employee_id, first_name, last_name, CASE WHEN  
    salary > 50000 THEN 'High' WHEN salary BETWEEN 30000 AND  
    50000 THEN 'Medium' ELSE 'Low' END AS salary_levelFROM  
employees;
```

# *Conditional Logic*

## PostgreSQL:

- `SELECT employee_id, first_name, last_name, CASE WHEN salary > 50000 THEN 'High' WHEN salary BETWEEN 30000 AND 50000 THEN 'Medium' ELSE 'Low' END AS salary_level FROM employees;`

# *PL/SQL vs PL/pgSQL*

## **Oracle (PL/SQL):**

```
CREATE OR REPLACE PROCEDURE raise_salary (emp_id NUMBER)
ASBEGIN UPDATE employees SET salary = salary * 1.05 WHERE
employee_id = emp_id;END;
```

# *PL/SQL vs PL/pgSQL*

## PostgreSQL (PL/pgSQL):

```
CREATE OR REPLACE FUNCTION raise_salary(emp_id INT) RETURNS VOID AS $$BEGIN UPDATE employees SET salary = salary * 1.05 WHERE employee_id = emp_id;END;$$ LANGUAGE plpgsql;
```

# *JSON Support*

## **Oracle:**

```
SELECT      json_value(employee_data,      '$.firstName')      AS
first_nameFROM employees_json;
```

## **PostgreSQL:**

```
SELECT employee_data->>'firstName' AS first_nameFROM
employees_json;
```

# *Recursive Queries*

## **Oracle:**

```
WITH employee_hierarchy AS ( SELECT employee_id,
manager_id, first_name, last_name FROM employees
WHERE manager_id IS NULL UNION ALL SELECT
e.employee_id, e.manager_id, e.first_name,
e.last_name FROM employees e
JOIN employee_hierarchy eh ON
e.manager_id = eh.employee_id)
SELECT * FROM employee_hierarchy;
```

# *Recursive Queries*

**PostgreSQL:**

```
WITH RECURSIVE employee_hierarchy AS
( SELECT employee_id, manager_id, first_name, last_name
  FROM employees WHERE manager_id IS NULL UNION ALL
  SELECT e.employee_id, e.manager_id, e.first_name, e.last_name
    FROM employees e JOIN employee_hierarchy eh
   ON e.manager_id = eh.employee_id)
SELECT * FROM employee_hierarchy;
```

# Summary

## Syntax Similarities:

Both Oracle and PostgreSQL share SQL syntax for basic operations, making it relatively easy to switch between the two for common tasks.

# Summary

**Syntax Similarities:** Both Oracle and PostgreSQL share SQL syntax for basic operations, making it relatively easy to switch between the two for common tasks.

**Differences in Functions:** Some functions and operators differ between the two systems, particularly with string concatenation and date functions.

# Summary

**Procedural Languages:** Both databases support procedural languages, but the syntax and available features can vary.

**JSON Support:** PostgreSQL offers more straightforward syntax for JSON operations, while Oracle requires specific functions.

**Licensing and Cost:** PostgreSQL is open-source and free, whereas Oracle requires licensing fees, which can be significant.