



PostgreSQL

Advanced Postgres

Surendra Panpaliya

# Agenda

Backup and Recovery

Importance of backups

Backup strategies (pg\_dump, pg\_basebackup)

Restoring data from backups

# Agenda

Advanced SQL Queries

Joins (inner, outer, cross)

Subqueries and common table expressions (CTEs)

Window functions

# Backup and Recovery

Surendra Panpaliya

# Backup and Recovery



Backups are critical



for ensuring



data integrity



availability.

# Backup and Recovery

Provide a safeguard

against data loss

due to Hardware failures

Software bugs

Human errors

# Backup and Recovery

Regular  
backups

Allow to

Restore  
Database

To a Previous  
state

Minimize  
downtime

Data loss.

# Importance of backups



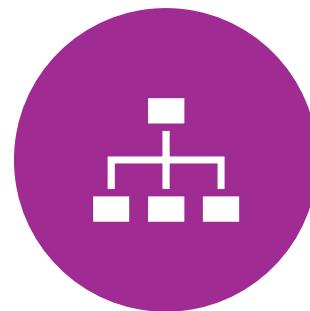
DATA IS CRITICAL AND  
OFTEN SENSITIVE,



INCLUDING PATIENT  
RECORDS,



TREATMENT HISTORIES,  
DIAGNOSTIC RESULTS



ADMINISTRATIVE  
INFORMATION.

# Patient Safety



ENSURING THE  
AVAILABILITY AND



INTEGRITY OF PATIENT  
RECORDS IS



CRUCIAL FOR  
PROVIDING  
CONTINUOUS AND



ACCURATE CARE.

# Data Integrity and Compliance



HEALTHCARE  
PROVIDERS MUST  
COMPLY



WITH REGULATIONS  
SUCH AS HIPAA IN THE  
U.S.



WHICH MANDATE  
SECURE



RELIABLE HANDLING OF  
PATIENT INFORMATION.

# Disaster Recovery

Backups provide a means

to restore data in the event of

Hardware failure

Cyber-attacks

Accidental deletions

Other unforeseen events

# Operational Continuity



Ensuring



Administrative  
functions



Appointment  
scheduling



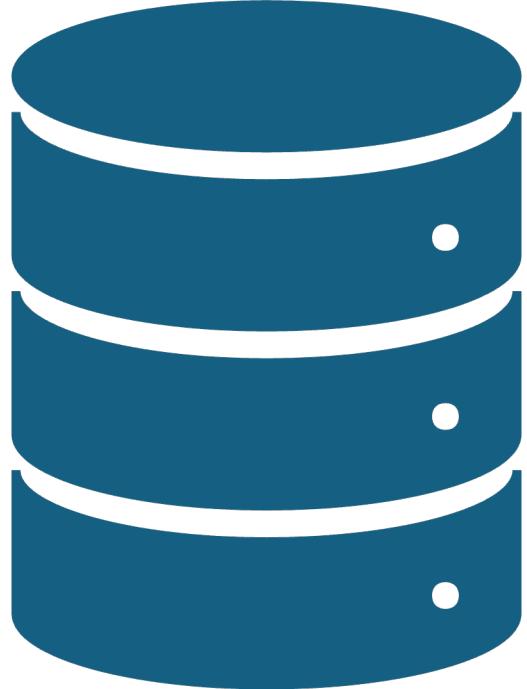
Billing



Continue  
smoothly



without data  
loss



# Backup Examples

---

Surendra Panpaliya

# Backup with pg\_dump



Step 1: Backup Individual Tables



Backup the patients table

# Backup with pg\_dump

-- Backup the patients table

```
pg_dump -U postgres -d healthcare_db -t patients -F c -f  
/path/to/backup/patients_table.backup
```

# **Backup the appointments table**

```
pg_dump -U postgres -d healthcare_db -t appointments -F c -f  
/path/to/backup/appointments_table.backup
```

# Backup the appointments table

By backing up

Individual tables

Ensure critical patient

Appointment data is Secure

Can be quickly restored if needed

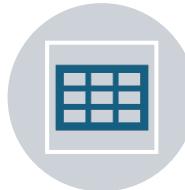
# **Step 2: Backup the Entire Database**

```
pg_dump -U postgres -d healthcare_db -F c -f  
/path/to/backup/healthcare_db.backup
```

# Step 2: Backup the Entire Database



A full database backup



captures all  
data, schemas



functions,  
allowing



a complete  
restoration

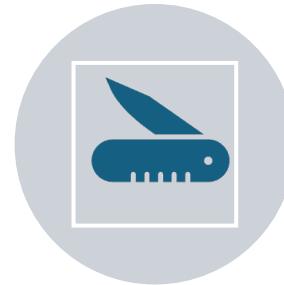


in case of a  
disaster

# Backup with pg\_basebackup



For full-cluster  
backups,



Useful for point-in-time  
recovery



Ensuring complete  
Data Integrity



Across multiple  
databases

# Step 1: Create the Backup

```
pg_basebackup -U postgres -D  
/path/to/backup/healthcare_cluster -Fp -Xs -P
```

# Step 1: Create the Backup

Captures  
Entire cluster,

Including  
configuration  
files

Making it  
suitable for  
recovering

from severe  
failures

where entire  
instances

might need to  
be restored

# Example Scenario: Disaster Recovery

---

Suppose there is

---

Some Attack

---

on the healthcare system

---

the data becomes

---

Inaccessible

# Backup strategies

Provides several methods

for backing up data,

Suitable for

different Scenarios

Requirements

# pg\_dump

---



A UTILITY



FOR  
PERFORMING



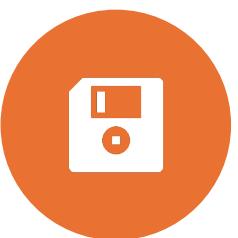
LOGICAL



BACKUPS

# pg\_dump

---



DUMPS A  
DATABASE



INTO A SCRIPT OR



ARCHIVE FILE



CONTAINING SQL  
COMMANDS



TO RECREATE  
THE DATABASE

# Use Cases

---



SUITABLE FOR



SMALLER DATABASES



LOGICAL BACKUPS

# Syntax

---

`pg_dump dbname > backupfile.sql`

---

**Example:**

---

`pg_dump mydatabase > mydatabase_backup.sql`

# pg\_basebackup

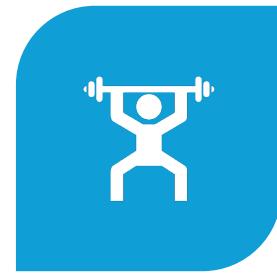
---



A UTILITY



FOR  
PERFORMING



PHYSICAL



BACKUPS

# pg\_basebackup

---



CREATES A



BINARY COPY



OF THE  
DATABASE



CLUSTER'S  
FILES

# pg\_basebackup

Useful for

creating base backups

for streaming replication or

continuous archiving

# Use Cases

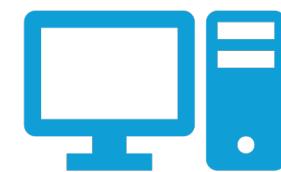
---



Suitable for



larger databases



physical backups

# Use Cases

---

You need

to restore

the entire

database cluster

# Syntax



```
pg_basebackup -D /path/to/backupdir -Fp -Xs -P
```



**Example:**



```
pg_basebackup -D /var/lib/postgresql/data/backup -Fp -Xs -P
```

# Continuous Archiving and Point-in-Time Recovery (PITR):

---



USES



PG\_BASEBACKUP



FOR BASE BACKUPS

# Continuous Archiving and Point-in-Time Recovery (PITR):

---



WITH CONTINUOUS



WAL (WRITE-AHEAD  
LOGGING)  
ARCHIVING



ENABLE RECOVERY



ANY POINT IN TIME

# Use Cases

Suitable for environments

Requiring high availability

Minimal data loss.

# Configuration



Enable WAL archiving in postgresql.conf:



`archive_mode = on`



`archive_command = 'cp %p /path/to/archive/%f'`

# Restoring data from backups

---

Surendra Panpaliya

# Restoring from pg\_dump

---



**Restore a Database**



`psql dbname < backupfile.sql`

# Restore a Database

---

```
psql mydatabase < mydatabase_backup.sql
```

# Restoring from pg\_basebackup

---

## **Steps:**

1. Stop the PostgreSQL server.
2. Clear the existing data directory.
3. Restore the backup files.
4. Start the PostgreSQL server.

# Restoring from pg\_basebackup:

---

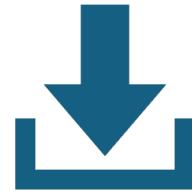
## **Example:**

```
pg_ctl stop -D /var/lib/postgresql/data  
rm -rf /var/lib/postgresql/data/*  
cp -r /path/to/backup/* /var/lib/postgresql/data/  
pg_ctl start -D /var/lib/postgresql/data
```

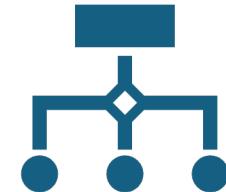
# Point-in-Time Recovery (PITR)



Restore the base backup

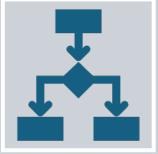


Restore the archived  
WAL files



Start the PostgreSQL  
server in recovery mode.

# Example Configuration



Create a recovery command in recovery.conf



```
restore_command = 'cp /path/to/archive/%f %p'
```

# Example Configuration



Start the server in recovery mode



```
pg_ctl start -D /var/lib/postgresql/data -m recovery
```

# Best Practices

---

Surendra Panpaliya

# Regular Backups

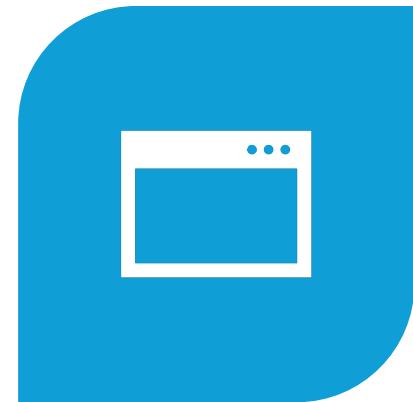
---



SCHEDULE REGULAR  
BACKUPS



BASED ON THE CRITICALITY  
OF THE DATA AND



THE ACCEPTABLE DATA  
LOSS WINDOW.

# Automate Backups

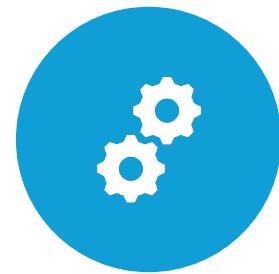
---



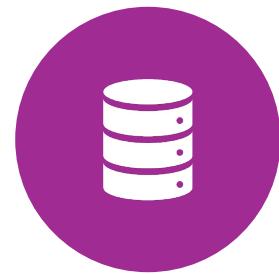
USE CRON  
JOBS



SCHEDULING  
TOOLS



TO AUTOMATE



BACKUP  
PROCESSES

# Verify Backups

---



REGULARLY  
TEST



BACKUPS



TO ENSURE



RESTORE

# Secure Backups

---

Store  
backups in

Secure  
location

with  
restricted  
access

use  
encryption

if necessary

# Keep Multiple Copies

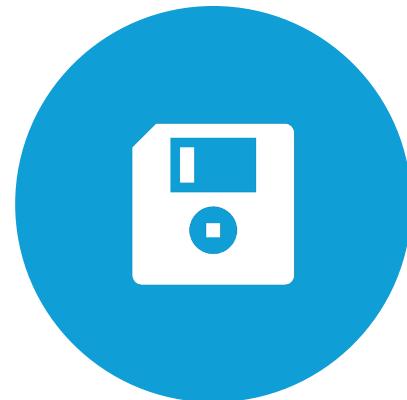
---



STORE MULTIPLE COPIES OF BACKUPS IN



DIFFERENT PHYSICAL OR CLOUD LOCATIONS



TO PROTECT AGAINST DATA LOSS.

# Agenda

---

Advanced SQL Queries

Joins (inner, outer, cross)

Subqueries and Common Table Expressions (CTEs)

Window functions

# Joins (inner, outer, cross)

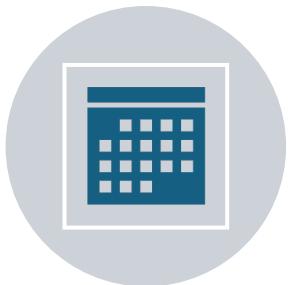
patients:

Stores patient information

doctors:

Stores doctor information

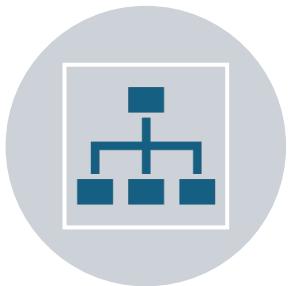
# Joins (inner, outer, cross)



appointments:



Stores information about patient appointments with doctors.



departments:



Stores information about hospital departments.

## Create patients table

---

```
CREATE TABLE patients (
    patient_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    date_of_birth DATE,
    gender CHAR(1)
);
```

# Create doctors table

---

```
CREATE TABLE doctors (
    doctor_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    department_id INTEGER
);
```

# Create departments table

---

```
CREATE TABLE departments (
    department_id SERIAL PRIMARY KEY,
    department_name VARCHAR(50)
);
```

# Create appointments table

---

```
CREATE TABLE appointments (
    appointment_id SERIAL PRIMARY KEY,
    patient_id INTEGER REFERENCES patients(patient_id),
    doctor_id INTEGER REFERENCES doctors(doctor_id),
    appointment_date DATE,
    status VARCHAR(20)
);
```

# Inserting Sample Data

---

-- Insert data into departments

```
INSERT INTO departments (department_name) VALUES  
('Cardiology'), ('Neurology'), ('Pediatrics'), ('Oncology');
```

## -- Insert data into doctors

---

```
INSERT INTO doctors (first_name, last_name, department_id) VALUES  
('John', 'Doe', 1), ('Jane', 'Smith', 2), ('Alice', 'Johnson', 3), ('Bob', 'Brown', 4);
```

# Insert data into patients

---

```
INSERT INTO patients (first_name, last_name, date_of_birth, gender)
VALUES
('Michael', 'Jackson', '1960-08-29', 'M'),
('Elvis', 'Presley', '1935-01-08', 'M'),
('Marilyn', 'Monroe', '1926-06-01', 'F');
```

## -- Insert data into appointments

---

```
INSERT INTO appointments (patient_id, doctor_id, appointment_date, status) VALUES  
(1, 1, '2024-07-01', 'Completed'),  
(1, 2, '2024-07-02', 'Scheduled'),  
(2, 1, '2024-07-03', 'Cancelled'),  
(3, 3, '2024-07-04', 'Completed');
```

# Join Queries

**INNER JOIN**

Fetch  
information

about  
appointments

along with

patient and

doctor details

# INNER JOIN

SELECT

```
p.first_name AS patient_first_name,  
p.last_name AS patient_last_name,  
d.first_name AS doctor_first_name,  
d.last_name AS doctor_last_name,  
a.appointment_date,  
a.status
```

# INNER JOIN

FROM

    appointments a

INNER JOIN patients p ON a.patient\_id = p.patient\_id

INNER JOIN doctors d ON a.doctor\_id = d.doctor\_id;

## 2. LEFT OUTER JOIN

Fetch all patients

appointments

include patients

with no appointments

## 2. LEFT OUTER JOIN

SELECT

```
p.first_name AS patient_first_name,  
p.last_name AS patient_last_name,  
a.appointment_date,  
a.status
```

## 2. LEFT OUTER JOIN

FROM

patients p

LEFT JOIN appointments a ON p.patient\_id = a.patient\_id;

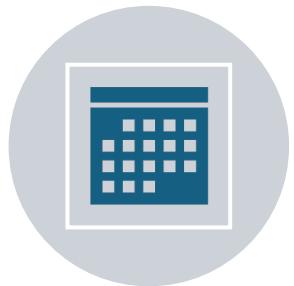
### 3. RIGHT OUTER JOIN



Fetch all appointments



Corresponding patient information



Including appointments



Without patient details

### 3. RIGHT OUTER JOIN

SELECT

```
p.first_name AS patient_first_name,  
p.last_name AS patient_last_name,  
a.appointment_date,  
a.status
```

### 3. RIGHT OUTER JOIN

FROM

patients p

RIGHT JOIN appointments a ON p.patient\_id = a.patient\_id;

## 4. FULL OUTER JOIN

Fetch all  
patients

their  
appointments,

including  
those

with no  
matching  
records

on either side.

## 4. FULL OUTER JOIN

SELECT

```
p.first_name AS patient_first_name,  
p.last_name AS patient_last_name,  
a.appointment_date,  
a.status
```

## 4. FULL OUTER JOIN

FROM

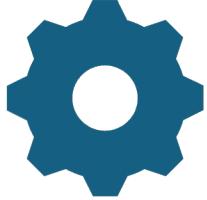
patients p

FULL OUTER JOIN appointments a ON p.patient\_id = a.patient\_id;

# 5. CROSS JOIN



Fetch every possible



combination of



patients and doctors

## 5. CROSS JOIN

Rarely useful

in practice

Can be used

for analysis

# 5. CROSS JOIN

SELECT

```
p.first_name AS patient_first_name,  
p.last_name AS patient_last_name,  
d.first_name AS doctor_first_name,  
d.last_name AS doctor_last_name
```

## 5. CROSS JOIN

FROM

patients p

CROSS JOIN doctors d;

# Summary

Queries showcase

How to combine data

from multiple tables

to retrieve

comprehensive information

# Summary



Essential for



Effective



Database  
management



Reporting



in the healthcare  
domain.

# Subqueries and common table expressions (CTEs)

Can be used

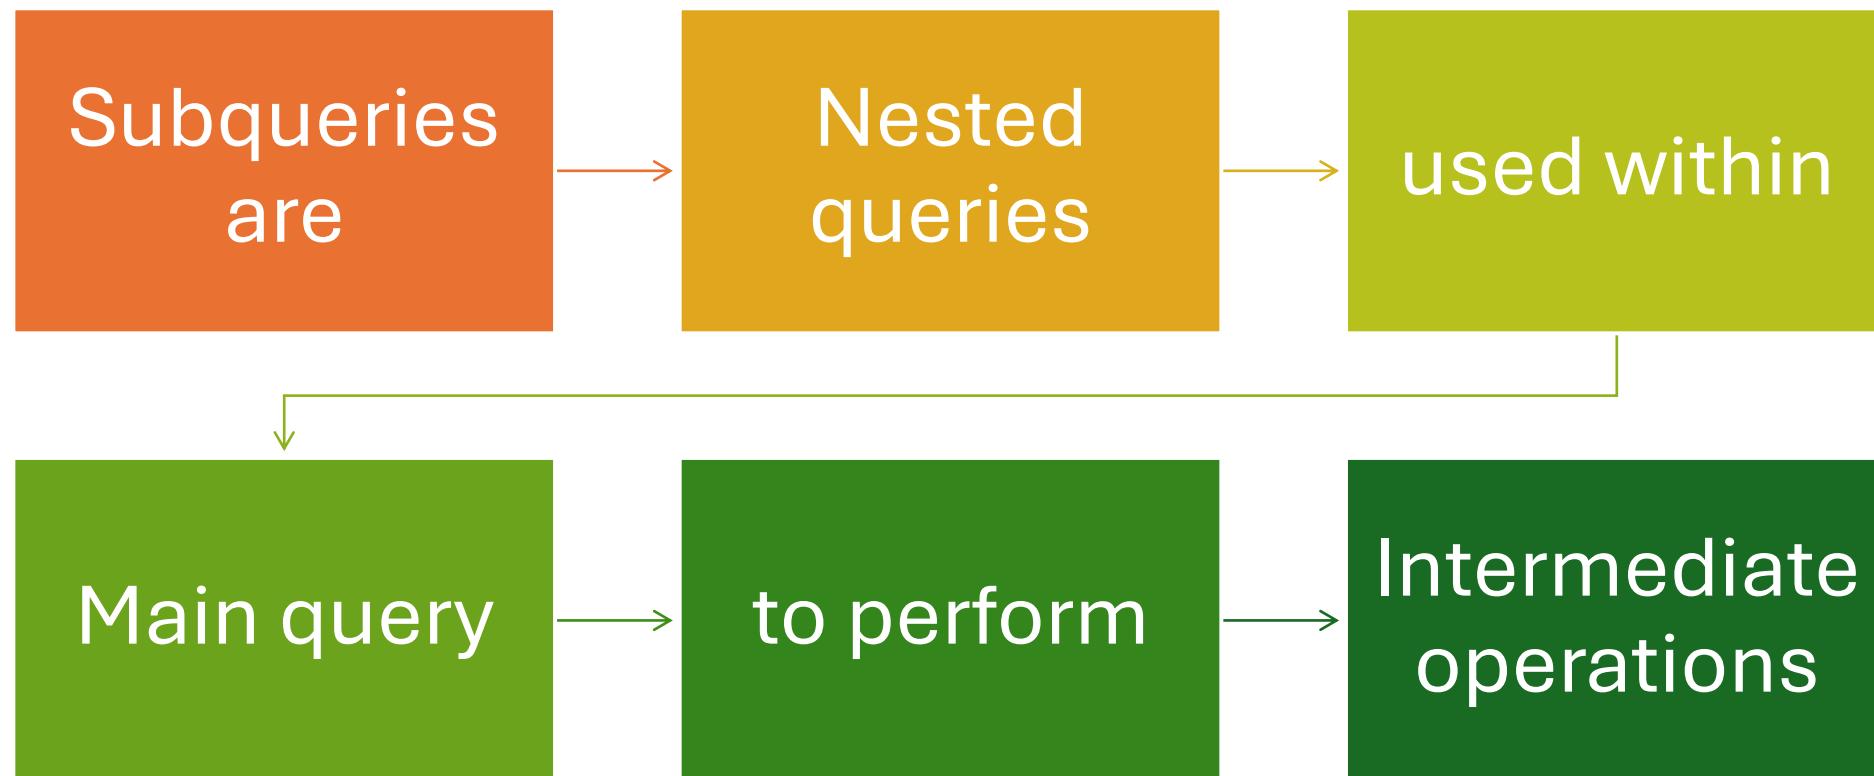
to simplify

complex queries

make them

more readable

# Subqueries



# Example 1: Subquery in SELECT

Fetch

patient names

along with

number of  
appointments

they have.

# Example 1: Subquery in SELECT

```
SELECT
    p.first_name,
    p.last_name,
    (SELECT COUNT(*)
     FROM appointments a
     WHERE a.patient_id = p.patient_id) AS appointment_count
FROM
    patients p;
```

## Example 2: Subquery in WHERE

Fetch  
appointments

of patients

who have

more than

one  
appointment

## Example 2: Subquery in WHERE

```
SELECT  
    a.appointment_id,  
    a.patient_id,  
    a.doctor_id,  
    a.appointment_date,  
    a.status
```

## Example 2: Subquery in WHERE

```
FROM
    appointments a
WHERE
    a.patient_id IN (
        SELECT
            patient_id
        FROM
            appointments
```

## Example 2: Subquery in WHERE

```
GROUP BY
    patient_id
HAVING
    COUNT(*) > 1
);
```

# Common Table Expressions (CTEs)



Provide a way



To break down



Complex  
queries



By defining



Temporary  
result sets

# Common Table Expressions (CTEs)

Can be

Referenced

within

Main query

# Example 1: Simple CTE

Fetch patient names

their total number of

appointments

using a CTE

# Example 1: Simple CTE

```
WITH patient_appointments AS (
    SELECT
        patient_id,
        COUNT(*) AS appointment_count
    FROM
        appointments
    GROUP BY
        patient_id
)
```

# Example 1: Simple CTE

```
SELECT
    p.first_name,
    p.last_name,
    pa.appointment_count
FROM
    patients p
JOIN
    patient_appointments pa ON p.patient_id = pa.patient_id;
```

## Example 2: Recursive CTE

Find all  
patients

who have  
referred

other  
patients

# Example 2: Recursive CTE

---



CREATE



THE REFERRALS  
TABLE



TO STORE



REFERRAL



RELATIONSHIPS

## Example 2: Recursive CTE

```
CREATE TABLE referrals (
   referrer_id INTEGER REFERENCES patients(patient_id),
   referred_id INTEGER REFERENCES patients(patient_id)
);
```

## Example 2: Recursive CTE

-- Insert sample data into referrals table

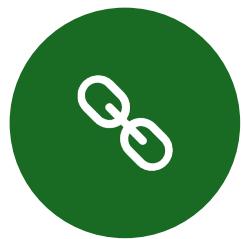
```
INSERT INTO referrals (referrer_id, referred_id) VALUES  
(1, 2), (2, 3), (3, 4);
```

# Example 2: Recursive CTE

---



USE A  
RECURSIVE CTE



TO FIND THE  
CHAIN OF



REFERRALS



STARTING  
FROM



A SPECIFIC  
PATIENT.

## Example 2: Recursive CTE

```
WITH RECURSIVE referral_chain AS (
```

```
    SELECT
```

```
        referrer_id,
```

```
        referred_id,
```

```
        1 AS level
```

```
    FROM
```

```
        referrals
```

# Example 2: Recursive CTE

```
WHERE
    referrer_id = 1
UNION
SELECT
    r.referrer_id,
    r.referred_id,
    rc.level + 1
FROM
    referrals r
INNER JOIN
    referral_chain rc ON r.referrer_id = rc.referred_id
)
```

## Example 2: Recursive CTE

```
SELECT  
    referrer_id,  
    referred_id,  
    level  
FROM  
    referral_chain;
```

# Example 3: CTE for Joining Multiple Tables

---

Fetch patient  
names

Doctor  
names

Appointment  
details

Using CTEs

To simplify  
the query

# Example 3: CTE for Joining Multiple Tables

```
WITH patient_details AS (
    SELECT
        patient_id,
        first_name AS patient_first_name,
        last_name AS patient_last_name
    FROM
        patients
),
```

# Example 3: CTE for Joining Multiple Tables

```
doctor_details AS (
    SELECT
        doctor_id,
        first_name AS doctor_first_name,
        last_name AS doctor_last_name
    FROM
        doctors
)
```

# Example 3: CTE for Joining Multiple Tables

```
SELECT  
    p.patient_first_name,  
    p.patient_last_name,  
    d.doctor_first_name,  
    d.doctor_last_name,  
    a.appointment_date,  
    a.status
```

# Example 3: CTE for Joining Multiple Tables

```
FROM
    appointments a
JOIN
    patient_details p ON a.patient_id = p.patient_id
JOIN
    doctor_details d ON a.doctor_id = d.doctor_id;
```

# Summary

Subqueries

Useful for

performing

intermediate operations

within a main query

# Summary

---

CTEs

Simplify

complex  
queries

by  
defining

temporary

result  
sets

# Summary

Can be

referenced

within

Main query

# Summary

---

Can be

recursive

to handle

hierarchical data

# Window functions

---

Provide a way

to perform calculations

across a set of table rows

that are related

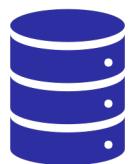
to the current row.

# Window functions

---



Useful for  
performing



Aggregations



Calculations



Over subsets  
of data

# Window Functions Examples

---

Example 1: ROW\_NUMBER()

Assign a unique row number

to each appointment

for a specific patient

# Example 1: ROW\_NUMBER()

```
SELECT
    patient_id,
    doctor_id,
    appointment_date,
    status,
    ROW_NUMBER() OVER (PARTITION BY patient_id ORDER BY
appointment_date) AS row_num
FROM
    appointments;
```

# Example 2: RANK()

---

Rank doctors

based on the  
number of

appointments  
they have.

Ties will  
receive

the same  
rank.

## Example 2: RANK()

```
SELECT
    doctor_id,
    COUNT(*) AS appointment_count,
    RANK() OVER (ORDER BY COUNT(*) DESC) AS rank
FROM
    appointments
GROUP BY
    doctor_id;
```

## Example 3: DENSE\_RANK()

Similar to  
RANK()

But without  
gaps

in ranking  
numbers.

## Example 3: DENSE\_RANK()

```
SELECT
    doctor_id,
    COUNT(*) AS appointment_count,
    DENSE_RANK() OVER (ORDER BY COUNT(*) DESC) AS
dense_rank
FROM
    appointments
GROUP BY
    doctor_id;
```

# Example 4: NTILE()

---

Divide the patients

into four quartiles

based on

their number

of appointments

## Example 4: NTILE()

```
SELECT
    patient_id,
    COUNT(*) AS appointment_count,
    NTILE(4) OVER (ORDER BY COUNT(*) DESC) AS quartile
FROM
    appointments
GROUP BY
    patient_id;
```

## Example 5: LAG()

Fetch

previous  
appointment  
details

for each  
patient

## Example 5: LAG()

```
SELECT
    patient_id,
    appointment_date,
    LAG(appointment_date) OVER (PARTITION BY patient_id ORDER
BY appointment_date) AS previous_appointment_date
FROM
    appointments;
```

## Example 6: LEAD()



FETCH



NEXT APPOINTMENT  
DETAILS



FOR EACH PATIENT.

## Example 6: LEAD()

```
SELECT
    patient_id,
    appointment_date,
    LEAD(appointment_date) OVER (PARTITION BY patient_id ORDER
BY appointment_date) AS next_appointment_date
FROM
    appointments;
```

## Example 7: SUM()

Calculate

the running

total of  
appointments

for each  
patient

over time

## Example 7: SUM()

```
SELECT
    patient_id,
    appointment_date,
    SUM(1) OVER (PARTITION BY patient_id ORDER BY
appointment_date) AS running_total
FROM
    appointments;
```

# Example 8: AVG()

Calculate

the average  
number of  
appointments

per patient

over a rolling  
window

of the last  
three  
appointments.

## Example 8: AVG()

```
SELECT
    patient_id,
    appointment_date,
    AVG(1) OVER (PARTITION BY patient_id ORDER BY
appointment_date ROWS BETWEEN 2 PRECEDING AND CURRENT
ROW) AS rolling_avg
FROM
    appointments;
```

# Summary

Window  
functions

allow for

powerful

flexible

data  
analysis

# Summary

Can be used

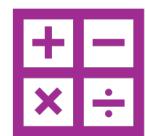
To Assign

Row numbers or

Ranks within

Partitions of data

# Summary



Calculate



Running  
totals,



Moving  
averages



Other  
cumulative

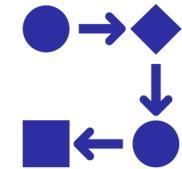


Statistics

# Summary



Fetch previous or



next values in a  
sequence,



useful for time-  
series data



like appointment  
histories.

# Summary



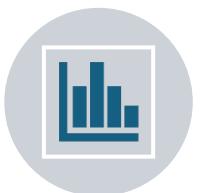
Learned



How to use  
window functions



To perform  
advanced



Data analysis



Reporting in



a healthcare  
database



**Thank you for  
your support and  
patience**

**Surendra Panpaliya**  
**Founder and CEO**  
**GKTCS Innovations**  
<https://www.gktcs.com>