



PostgreSQL

**Advanced Postgres**

Surendra Panpaliya

# Agenda

Views and Materialized Views

Creating and using views

Benefits of materialized views

Refreshing materialized views

# Agenda

01

Performance  
Tuning

02

Understanding  
query plans

03

Indexing  
strategies

04

Analyzing and  
optimizing  
queries

# Views and Materialized Views

**Surendra Panpaliya**

# Views



Virtual table



Representing



Result of



Stored query

# Views



Useful for



Simplifying  
complex queries,



Enforcing  
security



Creating a level  
of abstraction



Over raw data  
tables

# Creating a View

```
CREATE VIEW view_name AS  
SELECT column1, column2, ...  
FROM table_name  
WHERE condition;
```

# Example

```
CREATE TABLE patients (
    patient_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    date_of_birth DATE,
    gender CHAR(1)
);
```

# Example

To create a view that shows only male patients

```
CREATE VIEW male_patients AS  
SELECT patient_id, first_name, last_name, date_of_birth  
FROM patients  
WHERE gender = 'M';
```

# Using the View

```
SELECT * FROM male_patients;
```

# Updating Data Through a View

Views can also be updatable under certain conditions

```
UPDATE male_patients  
SET first_name = 'John'  
WHERE patient_id = 1;
```

# Advantages of Views



Simplifies complex queries.



Enhances security by restricting access



to specific rows or columns.



Provides a consistent, reusable interface.

# Materialized Views



Physical copy



of the result



of a query

# Materialized Views

Unlike regular views,

Materialized views

store data

Must be refreshed

when underlying

data changes.

# Creating a Materialized View

```
CREATE MATERIALIZED VIEW mat_view_name AS  
SELECT column1, column2, ...  
FROM table_name  
WHERE condition  
WITH [ NO ] DATA;
```

# Example

Create a materialized view

that stores the result

of the query

for male patients

# Example

```
CREATE MATERIALIZED VIEW mat_male_patients AS  
SELECT patient_id, first_name, last_name, date_of_birth  
FROM patients  
WHERE gender = 'M';
```

# Refreshing a Materialized View

Materialized  
views

must be  
refreshed

to get  
updated data

# Refreshing a Materialized View

```
REFRESH MATERIALIZED VIEW mat_male_patients;
```

# Using the Materialized View

```
SELECT * FROM mat_male_patients;
```

# Advantages of Materialized Views



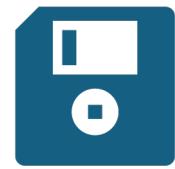
Improves  
performance



by storing query  
results.



Useful for read-  
heavy workloads



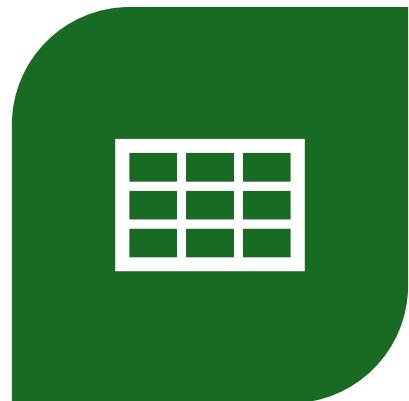
where data does not  
change frequently.

# Differences Between Views and Materialized Views

---



**STORAGE:**



VIEWS DO NOT STORE  
DATA;



MATERIALIZED VIEWS  
STORE THE QUERY RESULT.

# Differences Between Views and Materialized Views

---

**Performance:**

Views execute  
the underlying  
query

each time  
they are  
accessed;

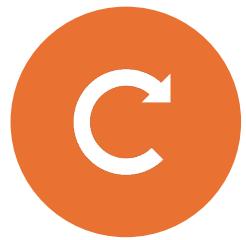
materialized  
views provide

faster access

to  
precomputed  
results.

# Differences Between Views and Materialized Views

---



REFRESH:



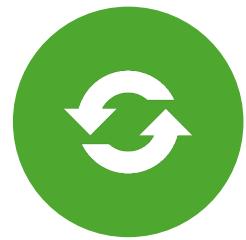
VIEWS ALWAYS  
SHOW



UP-TO-DATE DATA;



MATERIALIZED  
VIEWS



NEED TO BE  
EXPLICITLY  
REFRESHED.

# **Creating and using views**

**Surendra Panpaliya**

# Creating and using views

---



Views can be used



to simplify complex  
queries



improve security



by restricting data  
access

# Creating and using views



Provide



a consistent



Reusable  
interface



for frequently



accessed  
queries

# Creating Tables in the Healthcare

```
CREATE TABLE patients (
    patient_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    date_of_birth DATE,
    gender CHAR(1)
);
```

# Creating Tables in the Healthcare

```
CREATE TABLE appointments (
    appointment_id SERIAL PRIMARY KEY,
    patient_id INT REFERENCES patients(patient_id),
    doctor_name VARCHAR(50),
    appointment_date TIMESTAMP,
    diagnosis VARCHAR(255)
);
```

# Creating Tables in the Healthcare

```
CREATE TABLE prescriptions (
    prescription_id SERIAL PRIMARY KEY,
    appointment_id INT REFERENCES
appointments(appointment_id),
    medication_name VARCHAR(50),
    dosage VARCHAR(50),
    instructions TEXT
);
```

## 2. Creating Views for Simplified Data Access



### Example 1



### View for Patients' Basic Information



A view that provides basic information about patients

# **View for Patients' Basic Information**

```
CREATE VIEW patient_info AS  
SELECT patient_id, first_name, last_name, date_of_birth, gender  
FROM patients;
```

## 2. Creating Views for Simplified Data Access



### Usage

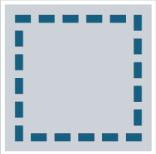


```
SELECT * FROM patient_info;
```

## 2. View for Upcoming Appointments

```
CREATE VIEW upcoming_appointments AS  
SELECT a.appointment_id, p.first_name, p.last_name,  
a.doctor_name, a.appointment_date, a.diagnosis  
FROM appointments a  
JOIN patients p ON a.patient_id = p.patient_id  
WHERE a.appointment_date > CURRENT_TIMESTAMP;
```

# 2 View for Upcoming Appointments



## Usage



```
SELECT * FROM upcoming_appointments;
```

### **3. View for Prescriptions by Patient**

A view that lists

all prescriptions

for a specific patient

# Ex3. View for Prescriptions by Patient

```
CREATE VIEW patient_prescriptions AS  
SELECT p.first_name, p.last_name, a.appointment_date,  
pr.medication_name, pr.dosage, pr.instructions  
FROM prescriptions pr  
JOIN appointments a ON pr.appointment_id = a.appointment_id  
JOIN patients p ON a.patient_id = p.patient_id;
```

### 3. View for Prescriptions by Patient

---

Usage:



```
SELECT * FROM patient_prescriptions WHERE  
first_name = 'John' AND last_name = 'Doe';
```

# 3. Updating Data Through Views



Views can also be used



to update data under certain conditions.



Let's assume we have a view



that includes patients' basic information

### **3. Updating Data Through Views**

```
CREATE VIEW editable_patient_info AS  
SELECT patient_id, first_name, last_name, date_of_birth, gender  
FROM patients  
WITH CHECK OPTION;
```

# 3. Updating Data Through Views

---

## Updating Data Using the View:

---

```
UPDATE editable_patient_info
```

---

```
SET first_name = 'Jonathan'
```

---

```
WHERE patient_id = 1;
```

# 4. Securing Data with Views

Views can restrict access

to sensitive data.

For example

Hide patients' full names in

a public-facing application

# 4. Securing Data with Views

```
CREATE VIEW anonymized_patient_info
AS SELECT patient_id,
CONCAT(SUBSTRING(first_name, 1, 1), '.', last_name)
AS name, date_of_birth, gender
FROM patients;
```

# 4. Securing Data with Views

**Usage:**

```
SELECT * FROM anonymized_patient_info;
```

# **Practical Examples**

**Surendra Panpaliya**

---

# **1. View for All Female Patients**

```
CREATE VIEW female_patients AS  
SELECT patient_id, first_name, last_name, date_of_birth  
FROM patients  
WHERE gender = 'F';
```

## 2. View for Recent Diagnoses

```
CREATE VIEW recent_diagnoses AS  
SELECT a.appointment_id, p.first_name, p.last_name,  
a.doctor_name, a.diagnosis, a.appointment_date  
FROM appointments a  
JOIN patients p ON a.patient_id = p.patient_id  
WHERE a.appointment_date > NOW() - INTERVAL '30 days';
```

### 3. View for Medication Instructions

```
CREATE VIEW medication_instructions AS  
SELECT p.first_name, p.last_name, pr.medication_name,  
pr.dosage, pr.instructions  
FROM prescriptions pr  
JOIN appointments a ON pr.appointment_id = a.appointment_id  
JOIN patients p ON a.patient_id = p.patient_id;
```

## 4. View for Daily Appointments

```
CREATE VIEW daily_appointments AS  
SELECT a.appointment_id, p.first_name, p.last_name,  
a.doctor_name, a.appointment_date  
FROM appointments a  
JOIN patients p ON a.patient_id = p.patient_id  
WHERE DATE(a.appointment_date) = CURRENT_DATE;
```

# Summary

---



VIEWS ARE A  
POWERFUL TOOL



FOR MANAGING



QUERYING DATA  
EFFICIENTLY.

# Summary

---



HELP SIMPLIFY  
COMPLEX  
QUERIES,



ENHANCE  
SECURITY



PROVIDE A LEVEL  
OF ABSTRACTION



THAT CAN MAKE  
THE DATABASE



MORE USER-  
FRIENDLY

# Summary

---



ENSURE THAT



SENSITIVE  
INFORMATION IS  
PROTECTED,



WHILE STILL  
MAKING  
NECESSARY DATA



EASILY  
ACCESSIBLE



TO AUTHORIZED  
USERS.

# **Benefits of materialized views**

**Surendra Panpaliya**

# Benefits of materialized views

Materialized views offer

several benefits in the healthcare domain,

especially in scenarios that

require frequent and

complex queries over large datasets.

# 1. Performance Improvement



**Faster Query Response**



Reduces the time needed



Materialized views



To fetch results



Store the result of a query physically.

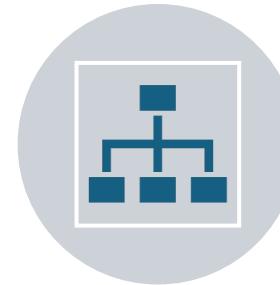
# 1. Performance Improvement



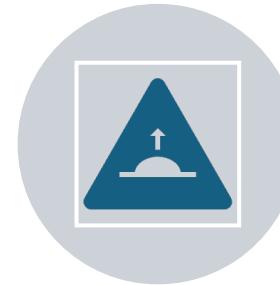
Example



such as patient visit  
summaries,



Frequently accessed  
Reports,



can be generated  
quickly.

# 1. Performance Improvement

```
CREATE MATERIALIZED VIEW patient_visit_summary AS  
SELECT patient_id, COUNT(*) AS visit_count  
FROM appointments  
GROUP BY patient_id;
```

# Optimized Data Aggregation

Aggregated data,

such as daily patient visits or

medication usage statistics,

can be precomputed and stored,

improving query performance.

# Optimized Data Aggregation

```
CREATE MATERIALIZED VIEW daily_patient_visits AS
SELECT DATE(appointment_date) AS visit_date,
COUNT(*) AS visit_count
FROM appointments
GROUP BY DATE(appointment_date);
```

## 2. Reduced Load on Source Tables



**Offloading Queries**



Complex and  
resource-intensive  
queries



are offloaded from  
source tables,



reducing the load  
and



contention on those  
tables.

## 2. Reduced Load on Source Tables

```
CREATE MATERIALIZED VIEW medication_usage AS  
SELECT medication_name, COUNT(*) AS usage_count  
FROM prescriptions  
GROUP BY medication_name;
```

# Improved Performance for Concurrent Users

---



IN ENVIRONMENTS  
WITH MANY  
CONCURRENT USERS,



MATERIALIZED VIEWS  
HELP IN



REDUCING THE  
PERFORMANCE  
IMPACT



ON THE UNDERLYING  
TABLES.

# 3. Simplified Complex Queries

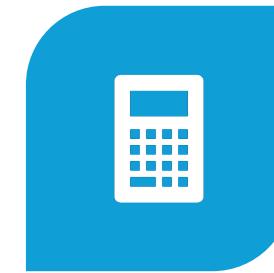
---



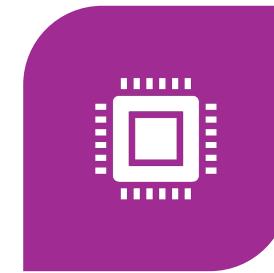
**PRECOMPUTED  
RESULTS**



QUERIES THAT  
INVOLVE COMPLEX  
JOINS AND



CALCULATIONS CAN  
BE PRECOMPUTED  
AND STORED,



MAKING IT EASIER  
FOR END-USERS TO  
RETRIEVE DATA.

### 3. Simplified Complex Queries

```
CREATE MATERIALIZED VIEW patient_medication_details AS  
  
SELECT p.patient_id, p.first_name, p.last_name,  
pr.medication_name, pr.dosage, pr.instructions  
  
FROM patients p  
  
JOIN appointments a ON p.patient_id = a.patient_id  
  
JOIN prescriptions pr ON a.appointment_id = pr.appointment_id;
```

# Consistent Results

---



MATERIALIZED VIEWS  
PROVIDE



CONSISTENT AND  
REPEATABLE  
RESULTS



FOR COMPLEX  
QUERIES,



ENSURING DATA  
INTEGRITY AND  
ACCURACY.

# 4. Historical Data and Snapshots



**Data Archiving**



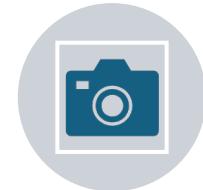
at specific points  
in time.



Materialized views  
can be used



Useful for



to create  
snapshots of data



Historical analysis  
and reporting

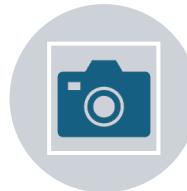
## 4. Historical Data and Snapshots

```
CREATE MATERIALIZED VIEW monthly_patient_visits AS  
  
SELECT DATE_TRUNC('month', appointment_date) AS month,  
COUNT(*) AS visit_count  
  
FROM appointments  
  
GROUP BY DATE_TRUNC('month', appointment_date);
```

# 4. Historical Data and Snapshots



**Trend Analysis**



Historical  
snapshots help



in analyzing trends  
over time,



such as patient  
admission rates or



medication usage  
patterns.

# 5. Improved Data Security and Access Control



**Restricting Access**



restricting direct access



Sensitive data can be filtered and



to the underlying tables



stored in a materialized view,



# 5. Improved Data Security and Access Control

```
CREATE MATERIALIZED VIEW limited_patient_info AS  
SELECT patient_id, first_name, last_name, date_of_birth  
FROM patients  
WHERE gender = 'F';
```

# 5. Improved Data Security and Access Control

## Role-Based Access Control

Materialized views can be used

to provide access to aggregated or

anonymized data,

ensuring compliance with

data protection regulations.

# 6. Reduced Network Traffic



## Local Storage of Data



reducing the need to  
fetch data



Materialized views store  
data locally,



from remote databases  
repeatedly.

# 6. Reduced Network Traffic

## Efficient Data Sharing

Data can be shared efficiently

between different departments or

applications without impacting

the performance of the source tables.

# 7. Data Consistency and Availability

---

**Ensuring  
Data  
Consistency**

Materialized views provide

a consistent state of data,

which is especially useful

in environments where

data consistency is critical.

# 7. Data Consistency and Availability

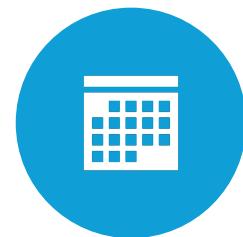
---



HIGH  
AVAILABILITY



MATERIALIZED  
VIEWS CAN BE  
REFRESHED



AT REGULAR  
INTERVALS,



ENSURING THAT  
DATA IS UP-TO-  
DATE AND



AVAILABLE FOR  
REPORTING AND  
ANALYSIS.

# **Examples in Healthcare Domain**

**Surendra Panpaliya**

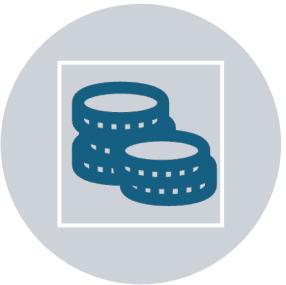
# Example 1: Daily Patient Summary



A materialized view



to summarize daily patient visits,



reducing the need



to query the appointments table frequently.

# Example 1: Daily Patient Summary

```
CREATE MATERIALIZED VIEW daily_patient_summary AS  
SELECT patient_id, DATE(appointment_date) AS visit_date,  
COUNT(*) AS visit_count  
FROM appointments  
GROUP BY patient_id, DATE(appointment_date);
```

# Example 2: Monthly Medication Usage



A materialized view



to track monthly usage  
of medications,



helping in inventory  
management and



procurement planning.

## Example 2: Monthly Medication Usage

```
CREATE MATERIALIZED VIEW monthly_medication_usage AS  
  
SELECT medication_name, DATE_TRUNC('month',  
prescription_date) AS month, COUNT(*) AS usage_count  
  
FROM prescriptions  
  
GROUP BY medication_name, DATE_TRUNC('month',  
prescription_date);
```

# Example 3: High-Risk Patients

A materialized view

to identify high-risk patients

based on certain criteria,

such as multiple visits

in a short period.

# Example 3: High-Risk Patients

```
CREATE MATERIALIZED VIEW high_risk_patients AS  
SELECT patient_id, COUNT(*) AS visit_count  
FROM appointments  
WHERE appointment_date > NOW() - INTERVAL '30 days'  
GROUP BY patient_id  
HAVING COUNT(*) > 5;
```

# Conclusion



Materialized views provide



numerous benefits in the healthcare domain



by improving query performance,



reducing load on source tables,



simplifying complex queries,



enhancing data security.

# Conclusion

---



Are a valuable tool



for efficient data  
management



analysis in healthcare  
databases.

# Refreshing materialized views

Surendra Panpaliya

# Refreshing materialized views



Crucial in a  
healthcare domain



to ensure that the  
data is



up-to-date for  
accurate



reporting and  
analysis.

# Refreshing materialized views

Materialized views can be refreshed

Manually or Automatically

at specified intervals,

depending on the requirements.

# Types of Refresh Methods

## Complete Refresh

This rebuilds

the entire materialized view

from scratch.

# Types of Refresh Methods

**Incremental  
Refresh**

Known as  
"fast refresh,"

Updates only  
the changes  
since

the last  
refresh,

which is  
more  
efficient

than a  
complete  
refresh.

# Types of Refresh Methods

---

## Manual Refresh

---

You can manually refresh a materialized view

---

using the REFRESH MATERIALIZED VIEW command.

---

## Example 1: Complete Refresh

---

```
REFRESH MATERIALIZED VIEW daily_patient_summary;
```

# Example 2: Incremental Refresh

To enable  
incremental  
refresh,

Create the  
materialized  
view

with the WITH  
DATA and

WITH NO  
DATA clauses,

ensure that  
the underlying  
tables

have LOGGED  
updates.

## Example 2: Incremental Refresh

```
CREATE MATERIALIZED VIEW daily_patient_summary AS  
SELECT patient_id, DATE(appointment_date) AS visit_date,  
COUNT(*) AS visit_count  
FROM appointments  
GROUP BY patient_id, DATE(appointment_date)  
WITH NO DATA;
```

# Example 2: Incremental Refresh



Later, you can do an incremental refresh:



```
REFRESH MATERIALIZED VIEW daily_patient_summary WITH  
DATA;
```

# Automatic Refresh

Set up Automatic refresh

using pg\_cron extension or

Similar job scheduler

to refresh the materialized view

at regular intervals.

# Automatic Refresh



## Example: Using pg\_cron



First, install the pg\_cron extension



(if not already installed)



```
CREATE EXTENSION pg_cron;
```

# Automatic Refresh

---

Schedule a job to refresh the materialized view

---

```
SELECT cron.schedule('0 0 * * *', 'REFRESH  
MATERIALIZED VIEW daily_patient_summary');
```

---

Schedules the materialized view to refresh  
daily at midnight.

# Examples in Healthcare Domain

Surendra Panpaliya

# Example 1: Daily Patient Summary

-- A materialized view to summarize daily patient visits, refreshed daily

```
CREATE MATERIALIZED VIEW daily_patient_summary
AS SELECT patient_id, DATE(appointment_date)
AS visit_date, COUNT(*) AS visit_count
FROM appointments
GROUP BY patient_id, DATE(appointment_date);
```

# Example 1: Daily Patient Summary

-- Schedule a daily refresh

```
SELECT cron.schedule('0 0 * * *', 'REFRESH  
MATERIALIZED VIEW daily_patient_summary');
```

# Example 2: Monthly Medication Usage



A materialized view



to track monthly usage of medications,



refreshed monthly

# Example 2: Monthly Medication Usage

```
CREATE MATERIALIZED VIEW monthly_medication_usage  
AS SELECT medication_name, DATE_TRUNC('month',  
prescription_date) AS month, COUNT(*) AS usage_count  
FROM prescriptions  
GROUP BY medication_name, DATE_TRUNC('month',  
prescription_date);
```

# Example 2: Monthly Medication Usage



-- Schedule a monthly refresh



```
SELECT cron.schedule('0 0 1 * *', 'REFRESH MATERIALIZED  
VIEW monthly_medication_usage');
```

# Example 3: High-Risk Patients

A materialized view

to identify high-risk patients

based on certain criteria,

refreshed hourly

# Example 3: High-Risk Patients

```
CREATE MATERIALIZED VIEW high_risk_patients AS
SELECT patient_id, COUNT(*) AS visit_count
FROM appointments
WHERE appointment_date > NOW() - INTERVAL '30 days'
GROUP BY patient_id
HAVING COUNT(*) > 5;
```

# Example 3: High-Risk Patients



-- Schedule an hourly refresh



```
SELECT cron.schedule('0 * * * *',
```



```
'REFRESH MATERIALIZED VIEW high_risk_patients');
```

# Summary



Refreshing  
materialized views



in a healthcare  
domain is



essential for  
maintaining



up-to-date and



accurate data.

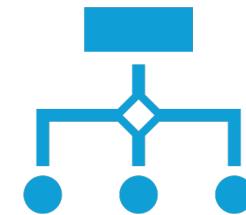
# Summary



Depending on the requirements,



materialized views can be refreshed



manually or automatically.

# Summary

Automating  
the refresh  
process

ensures  
timely  
updates

without  
manual  
intervention

# Summary

---



Improving the efficiency  
and



reliability of



data management in



healthcare databases.

# Agenda

---

1

**Performance Tuning**

2

Understanding query plans

3

Indexing strategies

4

Analyzing and optimizing queries

# Performance Tuning

---

Surendra Panpaliya

# Performance Tuning

---

Involves

Various  
strategies

Techniques to  
optimize

The database  
system

for faster and

more efficient  
operation

# Key Areas and Techniques

---

Surendra Panpaliya

# 1. Indexing

---

Indexes are critical for improving

the performance of queries.

However, they should be used

judiciously because

they can also impact

write performance.

# Types of Indexes

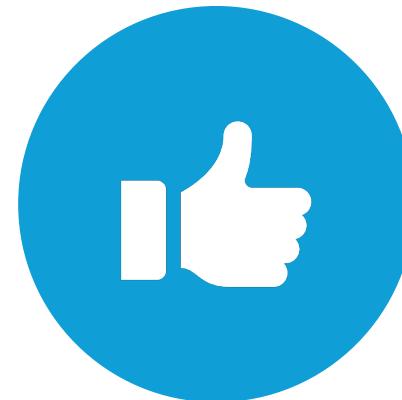
---



B-TREE INDEXES



DEFAULT INDEX TYPE,



GOOD FOR MOST  
QUERIES.

# Types of Indexes

Hash  
Indexes

Useful for

equality

comparisons.

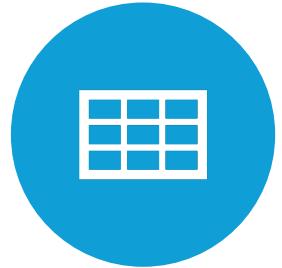
# Types of Indexes



**GIN/GIST INDEXES**



USEFUL FOR FULL-  
TEXT SEARCH,



JSONB DATA TYPES,  
AND



OTHER SPECIALIZED  
QUERIES.

# Types of Indexes

**BRIN Indexes**

Useful for very large tables

where columns have

a natural ordering.

# Creating an Index

```
CREATE INDEX idx_patient_last_name ON patients (last_name);
```

## 2. Query Optimization

Optimize SQL queries

for better performance.

Use EXPLAIN command

to understand

the query execution plan.

## 2. Query Optimization Example

```
EXPLAIN ANALYZE SELECT * FROM patients WHERE  
last_name = 'Smith';
```

# 3. Configuration Tuning



Adjust



PostgreSQL  
configuration



settings in  
`postgresql.conf`



for better  
performance.

# Important Settings



**shared\_buffers**



Amount of memory



the database server



uses for



shared memory buffers.

# Important Settings

`work_mem`

Amount of memory used

for internal sort operations and

hash tables.

# Important Settings



**Maintenance\_work\_mem:**



Memory allocated



for maintenance



tasks such as VACUUM.

# Important Settings



**effective\_cache\_size:**



An estimate of how much memory



is available for disk caching



by the operating system and



within the database itself.

# Important Settings

**max\_connections:**

The maximum number

of concurrent connections

to the database.

# Example Configuration

shared\_buffers = 4GB

work\_mem = 64MB

maintenance\_work\_mem = 512MB

effective\_cache\_size = 12GB

max\_connections = 100

# 4. Vacuuming and Analyzing



Regularly  
vacuum and



analyze the  
database



to maintain  
performance and



recover disk  
space.

# VACUUM Command

---

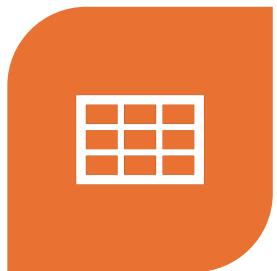
VACUUM FULL;

# **ANALYZE Command**

**ANALYZE;**

# 5. Partitioning

---



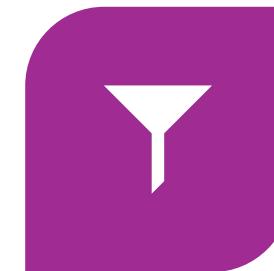
PARTITION LARGE  
TABLES



TO IMPROVE QUERY  
PERFORMANCE AND  
MANAGEABILITY.



POSTGRESQL  
SUPPORTS RANGE,



LIST, AND HASH  
PARTITIONING.

# Creating a Partitioned Table

```
CREATE TABLE patient_visits (
    visit_id SERIAL PRIMARY KEY,
    patient_id INT,
    visit_date DATE,
    details TEXT
) PARTITION BY RANGE (visit_date);
```

# 5. Partitioning



CREATE TABLE



patient\_visits\_2023



PARTITION OF patient\_visits



FOR VALUES



FROM ('2023-01-01') TO ('2023-12-31');

# Connection Pooling

Use connection pooling

to manage

database connections

efficiently.

# Connection Pooling



Tools like PgBouncer



can help



reduce the overhead



of establishing  
connections.

# Connection Pooling

## Example PgBouncer Configuration:

[databases]

mydb = host=127.0.0.1 port=5432 dbname=mydb

[pgbouncer]

listen\_addr = 127.0.0.1

listen\_port = 6432

# Connection Pooling

auth\_type = md5

auth\_file = /etc/pgbouncer/userlist.txt

pool\_mode = session

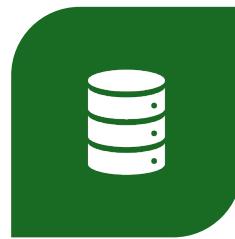
max\_client\_conn = 100

default\_pool\_size = 20

# 7. Caching



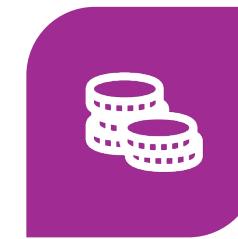
USE CACHING



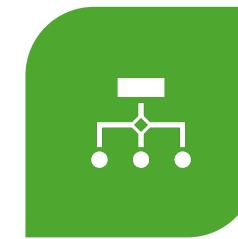
TO STORE  
FREQUENTLY  
ACCESSED DATA



IN MEMORY,



REDUCING THE  
NEED



FOR REPEATED  
DATABASE  
QUERIES.

# Using Redis for Caching

```
import redis
```

```
r = redis.Redis(host='localhost', port=6379, db=0)
```

```
# Set a value in the cache
```

```
r.set('patient_1234', 'John Doe')
```

```
# Get a value from the cache
```

```
patient_name = r.get('patient_1234')
```

# 8. Monitoring and Logging



MONITOR



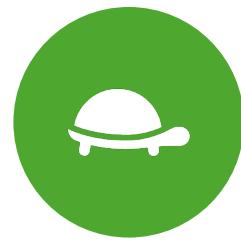
DATABASE  
PERFORMANCE



SET UP LOGGING



TO IDENTIFY  
BOTTLENECKS  
AND



SLOW QUERIES.

# 8. Monitoring and Logging



## Enabling Logging



`log_statement = 'all'`



`log_min_duration_statement = 500`

# Using Monitoring Tools

`pg_stat_statements:`

Provides execution statistics

for all SQL statements executed.

`CREATE EXTENSION pg_stat_statements;`

# Using Monitoring Tools

---



**pgAdmin:**



A web-based PostgreSQL



management tool.



**Prometheus and  
Grafana:**



For advanced  
monitoring and  
visualization.

# 9. Hardware and System Tuning

---



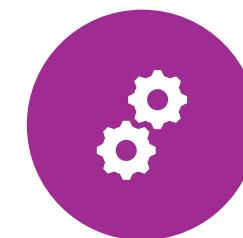
ENSURE THAT



THE UNDERLYING  
HARDWARE AND



OPERATING  
SYSTEM ARE



OPTIMIZED FOR



DATABASE  
PERFORMANCE.

# 9. Hardware and System Tuning

---



**Key  
Considerations:**



**Disk I/O:**



Use fast disks  
(e.g., SSDs)



for better  
performance

# 9. Hardware and System Tuning

**Key Considerations:**

**Memory:**

Ensure enough RAM is available

for the database and the operating system.

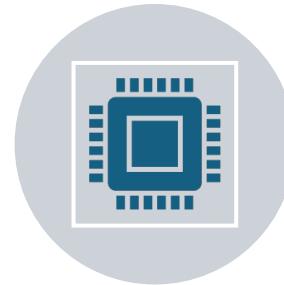
# 9. Hardware and System Tuning



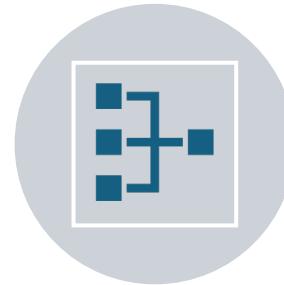
**Key Considerations:**



More CPU cores can help handle



**CPU:**



more parallel processes.

# 9. Hardware and System Tuning



**Key Considerations:**



Optimize network settings



**Network:**

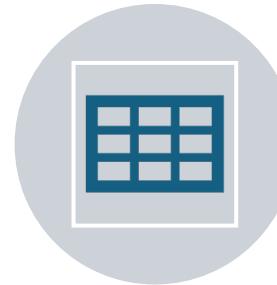


if the database is accessed over a network.

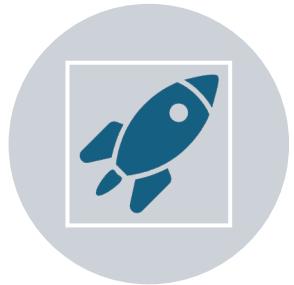
# 10. Use Appropriate Data Types



Choose the most appropriate



data types for your columns



to save space and



improve performance.

# 10. Use Appropriate Data Types

**Example:**

Use INTEGER instead of BIGINT

if the values will fit within the range of INTEGER.

Use TEXT or VARCHAR

for variable-length strings.

# Conclusion

Performance tuning in

PostgreSQL

is a multi-faceted

process

# Conclusion

Involves

optimizing queries,

indexing strategies

Configuration settings,

Hardware resources

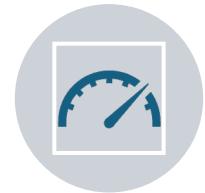
# Conclusion



Regular monitoring,



Maintenance, and  
adjustments based



on workload and  
performance  
metrics



are essential to  
ensure



optimal database  
performance.

# Understanding query plans

**Surendra Panpaliya**

# Understanding query plans

---



Essential for



optimizing  
database  
performance,



particularly in a  
healthcare  
domain



where timely and  
accurate data



retrieval is critical

# Understanding query plans



PostgreSQL  
provides tools



EXPLAIN and



EXPLAIN  
ANALYZE



to understand  
how queries



are executed.

# Why Query Plans are Important?

---



Query plans show



how PostgreSQL  
executes



a query.

# Why Query Plans are Important?

The order of operations (scans, joins, sorts, etc.)

Estimated and actual costs

Estimated and actual row counts

Use of indexes

Memory usage

# Why Query Plans are Important?



Can identify



inefficiencies



Make informed  
decisions



to optimize queries.

# Using EXPLAIN

Provides

query  
execution  
plan

without  
running the  
query.

# EXPLAIN ANALYZE

Runs the query  
and

provides the  
actual  
execution plan,

including run-  
time statistics.

# Example Scenarios

Retrieving Patient Information

```
SELECT * FROM patients WHERE  
last_name = 'Jackson';
```

# Execution Plan

EXPLAIN ANALYZE

SELECT \* FROM patients WHERE  
last\_name = 'Jackson';

# Sample Output

```
"Seq Scan on patients (cost=0.00..13.62 rows=1 width=252) (actual time=0.016..0.019 rows=1 loops=1)
```

```
" Filter: ((last_name)::text = 'Jackson'::text)"
```

```
" Rows Removed by Filter: 2"
```

```
"Planning Time: 0.061 ms"
```

```
"Execution Time: 0.036 ms"
```

# Sample Output

A sequential scan is used

because there is no index

on the last\_name column.

Adding an index on last\_name

can improve performance.

# Sample Output



```
CREATE INDEX idx_patients_last_name ON patients  
(last_name);
```



```
EXPLAIN ANALYZE SELECT * FROM patients WHERE last_name  
= 'Jackson';
```

# Sample Output After Indexing

```
"Seq Scan on patients (cost=0.00..1.04 rows=1 width=252) (actual time=0.010..0.011 rows=1 loops=1)
```

```
" Filter: ((last_name)::text = 'Jackson'::text)"
```

```
" Rows Removed by Filter: 2"
```

```
"Planning Time: 0.316 ms"
```

```
"Execution Time: 0.026 ms"
```

# Indexing strategies

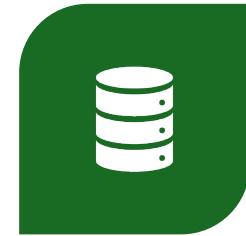
---

Surendra Panpaliya

# Indexing strategies



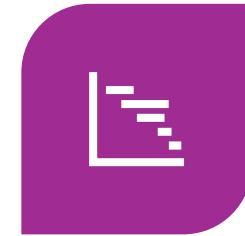
ESSENTIAL TO



ENSURE  
EFFICIENT DATA  
RETRIEVAL,



ESPECIALLY GIVEN



THE LARGE  
VOLUMES



OF DATA.

# Indexing strategies

Proper  
indexing can

significantly  
improve

query  
performance

reduce  
response  
times

optimize  
overall

database  
performance

# Analyzing and optimizing queries

---

**Surendra Panpaliya**

# Analyzing and optimizing queries

---



Crucial to ensure



efficient and timely access



to patient data,



which is vital



for healthcare providers

# Steps for Query Analysis and Optimization



Identify Slow  
Queries



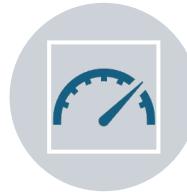
Analyze Query  
Plans



Optimize  
Indexes



Refactor  
Queries



Monitor  
Performance



Regular  
Maintenance

# 1. Identify Slow Queries

Use PostgreSQL's

`pg_stat_statements`  
extension

to track the  
performance of all  
queries:

# 1. Identify Slow Queries

Use PostgreSQL's

`pg_stat_statements`

extension to track

the performance of

all queries

# 1. Identify Slow Queries

```
CREATE EXTENSION pg_stat_statements;
```

Retrieve the most time-consuming queries:

```
SELECT query, total_time, calls, mean_time
```

```
FROM pg_stat_statements
```

```
ORDER BY total_time DESC
```

```
LIMIT 10;
```

## 2. Analyze Query Plans



Use the EXPLAIN



EXPLAIN ANALYZE  
commands



to understand



how PostgreSQL  
executes a query.

## 2. Analyze Query Plans

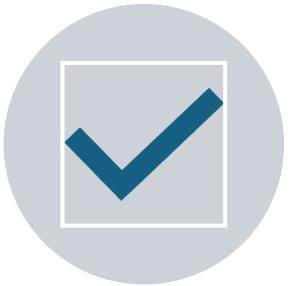
**Example Query:**

```
SELECT * FROM patients WHERE last_name = 'Smith';
```

**Execution Plan:**

```
EXPLAIN ANALYZE SELECT * FROM patients WHERE last_name = 'Smith';
```

# 3. Optimize Indexes



Indexes can significantly



speed up query performance.



Ensure that appropriate



indexes are in place.

# 3. Optimize Indexes



**Example Index:**



```
CREATE INDEX idx_patients_last_name
```



```
ON patients (last_name);
```

### 3. Optimize Indexes

Re-run the query and

check the new execution plan

```
EXPLAIN ANALYZE SELECT * FROM patients WHERE  
last_name = 'Smith';
```

# 4. Refactor Queries

Refactor complex queries

for better performance.

Break down complex queries

into smaller, manageable parts

using Common Table Expressions (CTEs)

or temporary tables.

# Example with CTE:

```
WITH RecentVisits AS (
    SELECT patient_id, visit_date
    FROM visits
    WHERE visit_date > '2023-01-01'
)
```

# **Example with CTE:**

```
SELECT p.patient_id, p.first_name, p.last_name, rv.visit_date  
FROM patients p  
JOIN RecentVisits rv ON p.patient_id = rv.patient_id;
```

# 5. Monitor Performance

Continuously monitor

query performance

using tools like pg\_stat\_statements and

built-in PostgreSQL views.

# 5. Monitor Performance



**Example Monitoring Query:**



```
SELECT * FROM pg_stat_activity WHERE state = 'active';
```

# 6. Regular Maintenance



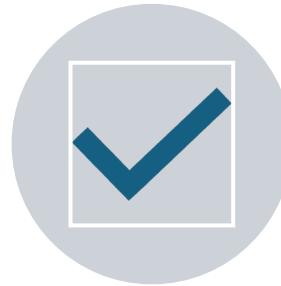
Perform regular maintenance tasks



such as vacuuming, analyzing, and



reindexing to keep the database



performance optimal.

# **6. Regular Maintenance**

**Vacuum and Analyze:**

**VACUUM ANALYZE;**

**Reindexing:**

**REINDEX TABLE patients;**

# Example Scenarios

**Scenario 1:**

Retrieving Patient Records by Last Name

**Initial Query:**

```
SELECT * FROM patients WHERE last_name = 'Smith';
```

# Example Scenarios



**Optimization:**



Add an index on last\_name:



```
CREATE INDEX idx_patients_last_name ON patients  
(last_name);
```

# Example Scenarios

## Revised Execution Plan:

```
EXPLAIN ANALYZE SELECT * FROM patients WHERE  
last_name = 'Smith';
```

Expect an index scan instead of a sequential scan, reducing execution time.

# Scenario 2: Joining Patient Visits with Doctors

```
SELECT p.patient_id, p.first_name, p.last_name,  
v.visit_date, d.doctor_name  
FROM patients p  
JOIN visits v ON p.patient_id = v.patient_id  
JOIN doctors d ON v.doctor_id = d.doctor_id  
WHERE v.visit_date > '2023-01-01';
```

# **Scenario 2: Joining Patient Visits with Doctors**

## **Optimization**

Ensure indexes on visits.patient\_id and doctors.doctor\_id:

```
CREATE INDEX idx_visits_patient_id ON visits (patient_id);
```

```
CREATE INDEX idx_doctors_doctor_id ON doctors (doctor_id);
```

# Revised Execution Plan

```
EXPLAIN ANALYZE SELECT p.patient_id,  
p.first_name, p.last_name, v.visit_date, d.doctor_name  
FROM patients p  
JOIN visits v ON p.patient_id = v.patient_id  
JOIN doctors d ON v.doctor_id = d.doctor_id  
WHERE v.visit_date > '2023-01-01';
```

# Revised Execution Plan

Check for nested loop joins and

index scans,

indicating better performance.



**Thank you for  
your support and  
patience**

**Surendra Panpaliya**  
**Founder and CEO**  
**GKTCS Innovations**

<https://www.gktcs.com>