



PostgreSQL

Advanced Postgres

Surendra Panpaliya

Agenda

Transactions and Concurrency Control

ACID properties

Transaction isolation levels

Handling deadlocks

Agenda

Monitoring and Maintenance

Monitoring tools

Routine maintenance tasks

Logging and troubleshooting

Transactions and Concurrency Control



Ensure data integrity



when multiple
transactions



Consistency



occur simultaneously

Transactions



A Sequence of
operations



performed as



A single logical unit of
work

Transactions

Follow the ACID properties

Atomicity

Consistency

Isolation

Durability

Atomicity

Ensures that all operations

within a transaction are completed successfully

If any operation fails,

the entire transaction is rolled back

Consistency

Ensures that

A transaction brings

the database from

one valid state to another

Isolation



ENSURES THAT



CONCURRENT
TRANSACTIONS



DO NOT AFFECT
EACH OTHER.

Durability



Ensures that



Once a transaction is committed,



It remains so,



even in the event of a system failure.

Transaction Commands



BEGIN



Starts a new transaction.

Transaction Commands

COMMIT

Ends the current transaction and

makes all changes permanent.

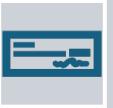
Transaction Commands

ROLLBACK

Ends the current transaction and

discards all changes.

Transaction Example



```
BEGIN;
```



```
INSERT INTO patients (name, ssn) VALUES ('Jane Doe', '123-45-6789');
```



```
UPDATE patients SET name = 'Jane Smith' WHERE ssn = '123-45-6789';
```



```
COMMIT;
```

Concurrency Control

PostgreSQL uses

Multi-Version Concurrency Control (MVCC)

to handle concurrency

Concurrency Control

MVCC allows multiple transactions

to read and write data

without locking each other out,

ensuring high performance and

Isolation Levels

Supports four
isolation
levels

Read
Uncommitted

Read
Committed

Repeatable
Read

Serializable

Read Uncommitted

Lowest isolation level.

Transactions can see

changes made by

other uncommitted transactions.

PostgreSQL treats this level as Read Committed

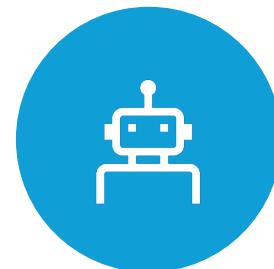
Read Committed



DEFAULT LEVEL.



TRANSACTIONS
ONLY SEE



CHANGES
COMMITTED



BEFORE THE
QUERY BEGAN.

Repeatable Read

Ensures that

if a transaction reads a row,

it will see the same data

if it reads the row again

within the same transaction.

Serializable

Highest isolation level.

Transactions are executed

in a way that their effect is

the same as if

they were executed serially

Example: Setting Isolation Level



```
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;
```



```
BEGIN;
```



```
-- Transaction operations
```



```
COMMIT;
```

Multi-Version Concurrency Control

MVCC allows

Multiple versions

of data to exist

simultaneously

Multi-Version Concurrency Control

Each transaction sees

a snapshot of the database

at a specific point in time,

ensuring consistent reads.

Visibility

Each transaction works

with a snapshot,

making changes invisible

to other transactions

until they are committed.

Version Management

Old versions
of data

are retained
until

no longer
needed

by any
transaction,

then cleaned
up

by the
VACUUM
process.

Locking Mechanisms

While MVCC reduces

the need for locking,

PostgreSQL still uses locks

to manage concurrent access

Locking Mechanisms



Row Locks:



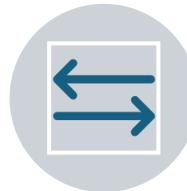
Acquired when



a row is
modified.



Released after



the transaction
ends.

Locking Mechanisms

**Table
Locks:**

Acquired for
operations

affecting an
entire table

Locking Mechanisms



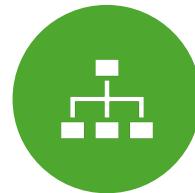
**Advisory
Locks:**



to enforce
complex



Explicit locks
that



business rules.



applications
can use

Example: Row Lock

```
BEGIN;
```

```
SELECT * FROM patients WHERE ssn = '123-45-6789' FOR UPDATE;
```

```
-- Perform some operations
```

```
COMMIT;
```

Deadlock Detection

Deadlocks occur

when two or more transactions

wait for each other

to release locks.

Deadlock Detection

PostgreSQL detects and

resolves deadlocks

by aborting one of

the transactions.

Example: Simulating a Deadlock

```
-- Session 1  
BEGIN;  
UPDATE patients  
SET name = 'John Doe'  
WHERE ssn = '123-45-6789';
```

Example: Simulating a Deadlock

```
-- Session 2  
BEGIN;  
UPDATE patients SET name = 'Jane Doe'  
WHERE ssn = '987-65-4321';  
UPDATE patients SET name = 'Jane Smith'  
WHERE ssn = '123-45-6789'; -- This will wait
```

Example: Simulating a Deadlock

-- Session 1

```
UPDATE patients SET name = 'John Smith'
```

```
WHERE ssn = '987-65-4321';
```

-- This will cause a deadlock

```
COMMIT;
```

ACID properties

Surendra Panpaliya

ACID properties in healthcare domain



CRUCIAL FOR



ENSURING
RELIABLE



CONSISTENT
DATABASE



TRANSACTIONS

1. Atomicity

Atomicity ensures that

a series of operations

within a transaction

are treated as a single unit.

1. Atomicity

Either all operations are

completed successfully, or

none are applied.

Example in Healthcare

Scenario: Updating patient records.

Operations:

A healthcare application updates a

patient's address and

phone number simultaneously.

Importance

Ensures that both the address and phone number are updated together.

If one update fails, both are rolled back to maintain data integrity.

SQL Example

```
BEGIN;  
-- Update patient address  
UPDATE patients  
SET address = '456 Elm St'  
WHERE patient_id = 123;
```

SQL Example

```
-- Update patient phone number
UPDATE patients
SET phone_number = '555-1234'
WHERE patient_id = 123;

COMMIT;
```

2. Consistency

Consistency ensures that

a transaction brings

the database from

one valid state to another,

maintaining database rules,

such as integrity constraints.

Example in Healthcare



Scenario: Adding a new patient record.



Operations:

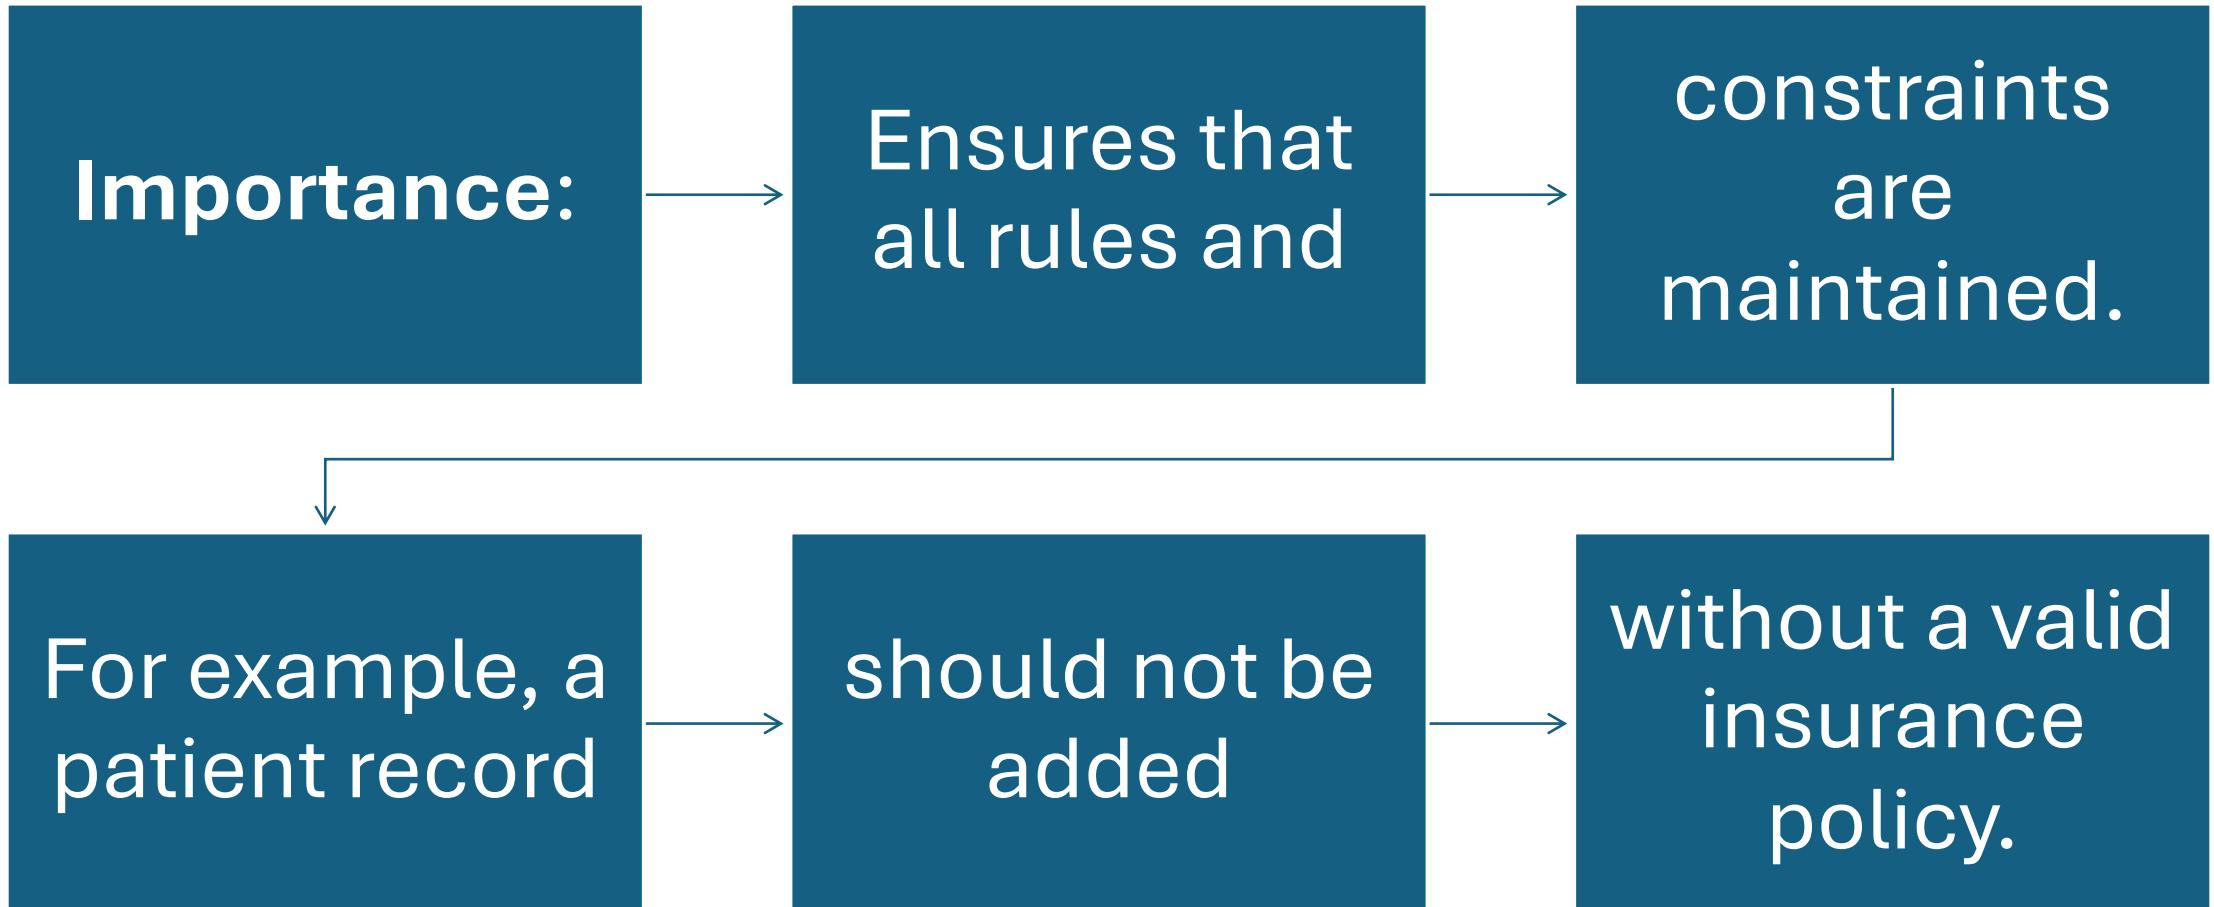


Inserting patient data and



ensuring a valid insurance policy reference.

Example in Healthcare



SQL Example

```
BEGIN;
```

```
-- Insert patient data
```

```
INSERT INTO patients
```

```
(patient_id, name, ssn, insurance_policy_id)
```

```
VALUES (124, 'Alice Brown', '987-65-4321', 789);
```

SQL Example

```
-- Insert insurance policy data
```

```
INSERT INTO insurance_policies (policy_id, patient_id, coverage)
```

```
VALUES (789, 124, 'Full Coverage');
```

```
COMMIT;
```

3. Isolation

Isolation ensures that

transactions are executed independently

without interference.

Each transaction sees

a consistent view of the database.

Example in Healthcare

Scenario:

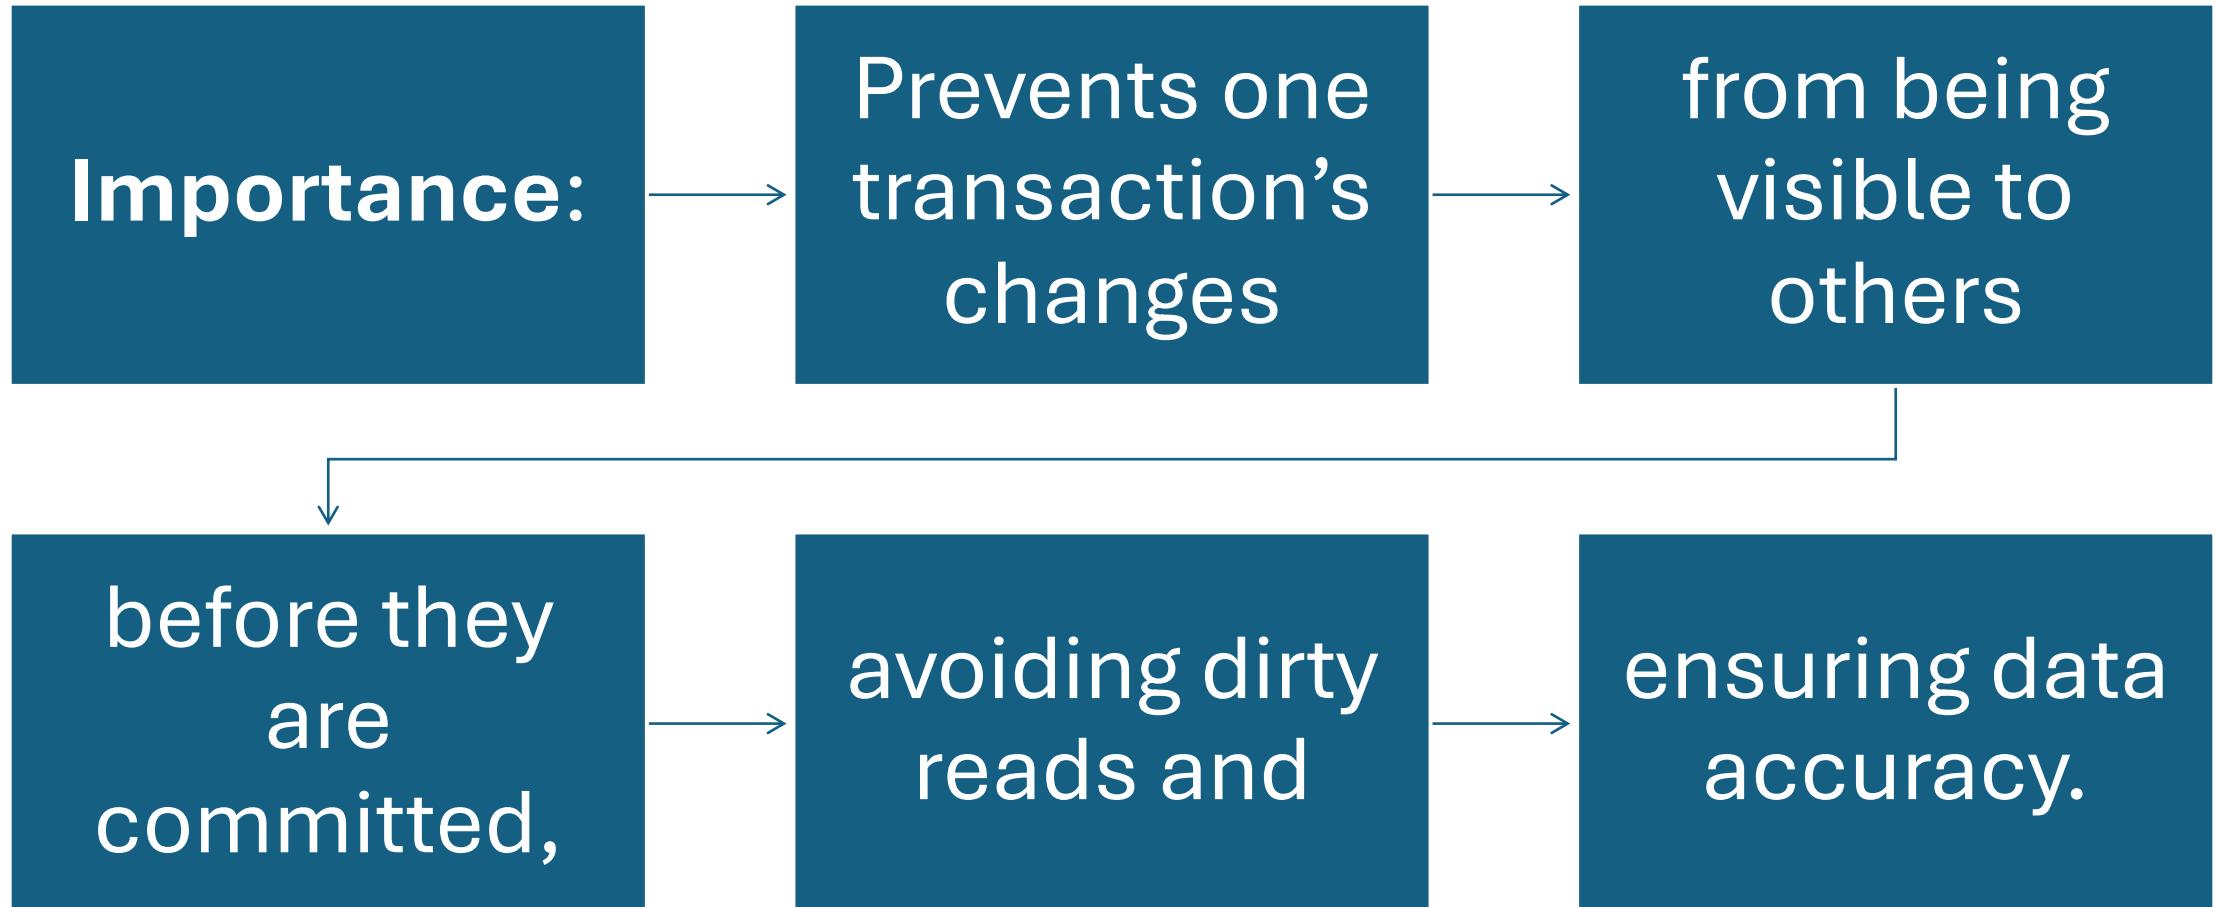
Concurrent access to patient records.

Operations:

One user updates patient data

while another queries the same data.

Example in Healthcare



SQL Example

```
-- Transaction 1  
BEGIN;  
UPDATE patients  
SET address = '789 Maple St'  
WHERE patient_id = 125;
```

SQL Example

-- Transaction 2 (executed concurrently)

```
BEGIN;  
SELECT * FROM patients  
WHERE patient_id = 125;
```

SQL Example

-- This transaction sees the old address

-- until Transaction 1 commits

COMMIT;

-- Transaction 1 commits

COMMIT;

4. Durability

Ensures once a

Transaction is committed,

Changes are permanent,

Even in the event

of a system failure

Example in Healthcare

Scenario:

Saving patient diagnosis records.

Operations:

Storing critical diagnosis information.

Importance



Ensures that once the diagnosis



information is recorded,



it is not lost even if



there is a system crash

SQL Example

```
BEGIN;
```

```
-- Insert diagnosis data
```

```
INSERT INTO diagnoses (diagnosis_id, patient_id,  
diagnosis_date, diagnosis)  
VALUES (2001, 126, '2024-07-01', 'Hypertension');
```

```
COMMIT;
```

Transaction isolation levels



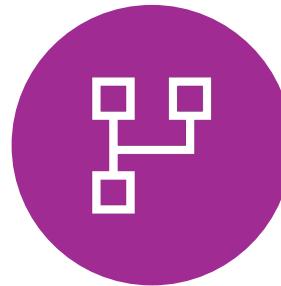
Ensure Data Integrity



Consistency



When multiple
Transactions



occur concurrently

Transaction isolation levels

PostgreSQL
supports

several
isolation
levels that

control the
visibility

of changes
made

by one
transaction

to other
concurrent
transactions.

1. Read Uncommitted

Transactions can

read data changes

made by other

Uncommitted transactions.

1. Read Uncommitted

This level is susceptible to

Dirty reads

Non-repeatable reads

Phantom reads.

Phantom reads

Occur in
database
systems

when a
transaction reads
a set of rows

that satisfy a
certain
condition.

Phantom reads

Upon re-executing

the same read operation,

finds additional rows (phantoms)

that did not exist previously.

Use Case in Healthcare

Example:

Generally not recommended

due to the risk of reading uncommitted or dirty data,

which can lead to inconsistent patient records.

Use Case in Healthcare

Example:

Not used in PostgreSQL as

it treats Read Uncommitted as

Read Committed.

2. Read Committed (Default Level)

A transaction sees only

committed changes

made by other transactions.

It does not see

uncommitted changes.

2. Read Committed (Default Level)

This level
avoids

dirty reads

but can still
encounter

Non-
repeatable
reads and

Phantom
reads.

Use Case in Healthcare



Example:



A doctor updates a patient's record,



while another transaction reads the patient's data.



Reader will see only Committed changes.

Example

-- Session 1: Update patient record

BEGIN;

UPDATE patients

SET address = '456 Elm St'

WHERE patient_id = 101;

Example

-- Session 2: Read patient record

BEGIN;

SELECT * FROM patients WHERE patient_id = 101; -- Sees the old address until Session 1 commits

COMMIT;

-- Session 1 commits

COMMIT;

3. Repeatable Read

Ensures that if a transaction reads a row,

it will see the same data

throughout the transaction,

preventing non-repeatable reads

but still allowing phantom reads.

Use Case in Healthcare



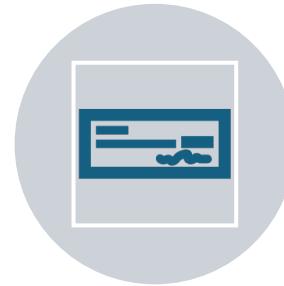
Example:



during a complex query
or report generation.



Ensuring consistent
reads of patient data



Useful for billing or
auditing processes.

Example

-- Session 1: Start transaction

BEGIN;

SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;

-- Session 1: Read patient record

SELECT * FROM patients WHERE patient_id = 102;

Example

-- Session 2: Update patient record

BEGIN;

UPDATE patients

SET address = '789 Maple St'

WHERE patient_id = 102;

COMMIT;

Example

-- Session 1: Read patient record again

```
SELECT * FROM patients
```

```
WHERE patient_id = 102;
```

-- Sees the old address until Session 1 commits

```
COMMIT;
```

4. Serializable

The highest isolation level,

where transactions are executed in a way

that their effect is the same as

if they were executed serially,

one after another.

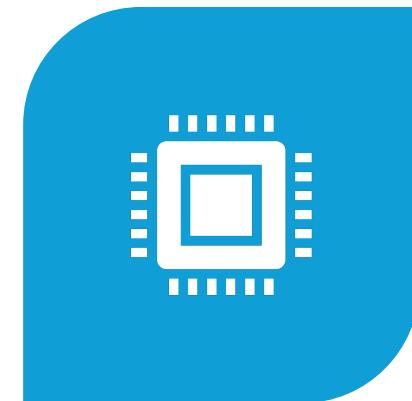
4. Serializable



PREVENTS DIRTY
READS,



NON-REPEATABLE
READS, AND



PHANTOM READS.

Use Case in Healthcare



Example:



Critical operations



where absolute data consistency is required,



such as updating patient medication records or



processing insurance claims.

Example

-- Session 1: Start transaction

BEGIN;

SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;

-- Session 1: Read patient record

SELECT * FROM patients WHERE patient_id = 103;

Example

```
-- Session 2: Attempt to update patient record
```

```
BEGIN;
```

```
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;
```

```
UPDATE patients
```

```
SET address = '1010 Pine St'
```

```
WHERE patient_id = 103;
```

Example

-- This will wait or fail if Session 1 has not yet committed
COMMIT;

-- Session 1: Read patient record again and commit
SELECT * FROM patients WHERE patient_id = 103;
COMMIT;

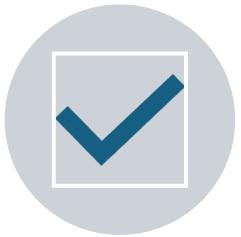
Handling deadlocks

Surendra Panpaliya

Handling deadlocks



INVOLVES
PREVENTING,
DETECTING



RESOLVING
SITUATIONS



WHERE TWO OR
MORE
TRANSACTIONS



ARE WAITING
INDEFINITELY



FOR ONE ANOTHER
TO RELEASE LOCKS.

Handling deadlocks

In healthcare systems,

deadlocks can lead to system downtime or

data inconsistencies, which can

have severe consequences.

1. Understanding Deadlocks



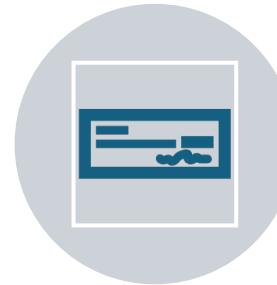
A deadlock occurs when two or more transactions



hold locks that the other transactions need,



creating a cycle of dependencies



that prevents any of the transactions from proceeding.

Example in Healthcare

Scenario:

Two transactions are updating patient records and insurance claims simultaneously, leading to a deadlock.

2. Preventing Deadlocks Strategies



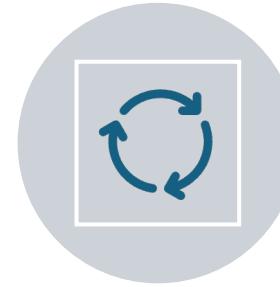
Consistent Lock Ordering:



in a consistent order



Ensure that transactions acquire locks



to prevent cyclical dependencies.

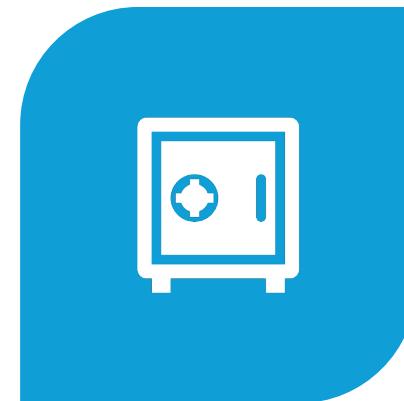
2. Preventing Deadlocks Strategies



**MINIMIZE LOCK
DURATION:**



**KEEP TRANSACTIONS
SHORT AND**



**LOCKS HELD FOR THE
MINIMAL AMOUNT OF TIME.**

2. Preventing Deadlocks Strategies

Use Appropriate Isolation Levels:

Choose isolation levels

that balance consistency

and concurrency.

Scenario



Updating patient records and



insurance claims in a specific order.

Scenario

-- Session 1: Update patient record, then update insurance claim

BEGIN;

UPDATE patients SET address = '456 Elm St'

WHERE patient_id = 201;

UPDATE insurance_claims SET status = 'processed'

WHERE claim_id = 301;

COMMIT;

Scenario

-- Session 2: Update patient record, then update insurance claim

BEGIN;

UPDATE patients SET phone_number = '555-6789'

WHERE patient_id = 202;

UPDATE insurance_claims SET status = 'pending'

WHERE claim_id = 302;

COMMIT;

3. Detecting Deadlocks



PostgreSQL automatically
detects deadlocks and



aborts one of the transactions
to break the cycle.



You can configure logging to
monitor deadlocks.

3. Detecting Deadlocks



Configuration:



Enable Deadlock Logging: Modify postgresql.conf.



`log_min_error_statement = 'error'`



Monitor Logs: Regularly check the PostgreSQL logs for deadlock messages.

Monitoring and Maintenance



Monitoring tools (`pg_stat_activity`, `pg_stat_statements`)

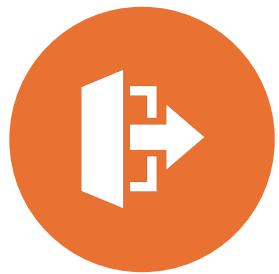


Routine maintenance tasks (`VACUUM`, `ANALYZE`)

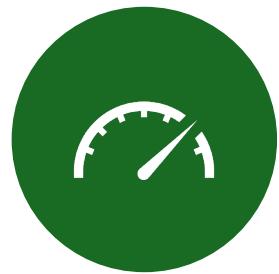


Logging and troubleshooting

Monitoring and Maintenance



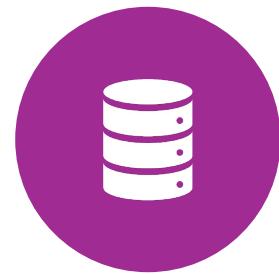
CRUCIAL FOR



ENSURING THE
PERFORMANCE,



RELIABILITY, AND
STABILITY OF A



POSTGRESQL
DATABASE.

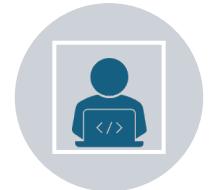
Monitoring and Maintenance



Effective monitoring helps in



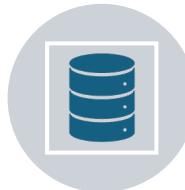
identifying and resolving issues



before they impact the system,



while regular maintenance



keeps the database running efficiently.

1. pg_stat_activity

This view shows the current activity in the database.

Example:

```
SELECT pid, username, application_name,  
client_addr, state, query  
FROM pg_stat_activity;
```

2. pg_stat_database

Provides aggregate statistics for each database in the cluster.

Example:

```
SELECT datname, numbackends, xact_commit, xact_rollback, blks_read, blks_hit  
FROM pg_stat_database;
```

3. pg_stat_user_tables

Shows access and usage statistics for tables.

Example:

```
SELECT relname, seq_scan, seq_tup_read, idx_scan, idx_tup_fetch, n_tup_ins,  
n_tup_upd, n_tup_del  
FROM pg_stat_user_tables;
```

4. pg_stat_user_indexes

Provides statistics on indexes.

Example:

```
SELECT relname, indexrelname, idx_scan, idx_tup_read, idx_tup_fetch  
FROM pg_stat_user_indexes;
```

5. pg_stat_bgwriter

Statistics about the background writer process, which is responsible for writing dirty pages to disk.

Example:

```
SELECT checkpoints_timed, checkpoints_req, buffers_checkpoint, buffers_clean,  
maxwritten_clean  
FROM pg_stat_bgwriter;
```

6. pg_stat_statements

Tracks execution statistics of all SQL statements executed by the server.

Installation:

```
CREATE EXTENSION pg_stat_statements;
```

Example

```
SELECT query, calls, total_time, rows,  
shared_blk_hit, shared_blk_read  
FROM pg_stat_statements  
ORDER BY total_time DESC  
LIMIT 5;
```

Routine maintenance tasks



VACUUM and
ANALYZE are essential



for maintaining the
performance and



reliability of
PostgreSQL databases

Routine maintenance tasks

Ensure that

the database
operates
efficiently

by reclaiming
storage,

updating
statistics for

the query
planner, and

preventing
database
bloat.

1. VACUUM

Used to
clean up

dead
tuples

Rows that
have been

updated
or deleted

1. VACUUM

Helps in reclaiming storage

preventing table bloat,

which can significantly

degrade performance over time.

Types of VACUUM



**Standard
VACUUM:**



Reclaims
storage



occupied by



dead tuples.

Types of VACUUM

VACUUM FULL:

Rewrites the entire table to reclaim space and

shrink the table size.

It locks the table during the operation.

Routine Cleanup of Patient Records



Regular
updates and



deletions of
patient records



can create dead
tuples,



requiring
routine cleanup



to maintain
performance.



VACUUM
patients;

Reclaiming Space After Bulk Deletions



When large amounts of outdated data,



such as old patient records, are deleted,



running VACUUM FULL helps reclaim the space.



VACUUM FULL patients;

Configuration for Autovacuum

Autovacuum is a background process

that automatically vacuums tables

at regular intervals,

reducing the need

for manual intervention.

Configuration for Autovacuum

```
psql -U postgres
```

```
# SHOW config_file;
```

```
C:\Program Files\PostgreSQL\16\data\postgresql.conf
```

Enable and configure autovacuum in postgresql.conf

autovacuum = on

autovacuum_max_workers = 3

autovacuum_naptime = 1min

autovacuum_vacuum_threshold = 50

autovacuum_analyze_threshold = 50

autovacuum_vacuum_scale_factor = 0.2

autovacuum_analyze_scale_factor = 0.1

2. ANALYZE

Command updates

the statistics used by

the PostgreSQL query planner

to determine the most efficient

way to execute queries.

Example Use Cases

Improving Query Performance for Patient Lookups

Frequent lookups of patient data can benefit from updated statistics to ensure optimal query plans.

ANALYZE patients;

Updating Statistics After Large Data Modifications

After bulk insertions, updates, or deletions in tables like appointments or lab_results, running ANALYZE ensures that the query planner has the most recent statistics.

ANALYZE appointments;

Combining VACUUM and ANALYZE

To perform both VACUUM and ANALYZE
in a single command,
use VACUUM ANALYZE.

This reclaims storage and
updates statistics simultaneously.

VACUUM ANALYZE patients;



Logging and troubleshooting

Surendra Panpaliya



Logging



Effective logging helps in identifying issues,



tracking changes, and ensuring compliance



with regulatory requirements.

Troubleshooting



TROUBLESHOOTING
TECHNIQUES ARE



ESSENTIAL FOR
DIAGNOSING AND



RESOLVING
PROBLEMS PROMPTLY



TO MINIMIZE
DISRUPTIONS.

Logging in PostgreSQL

PostgreSQL provides

various
logging
options

to capture
detailed
information

about
database
activities,

errors, and
performance.

Configuration for Logging

Logging settings are configured in the postgresql.conf file.

```
psql -U postgres
```

```
# SHOW config_file;
```

C:\Program Files\PostgreSQL\16\data\postgresql.conf

Basic Logging Settings

```
logging_collector = on          # Enable the collection of logs  
log_directory = 'pg_log'        # Directory where log files are stored  
log_filename = 'postgresql-%Y-%m-%d.log'  # Log file name pattern  
log_rotation_age = 1d           # Rotate log files daily  
log_rotation_size = 100MB       # Rotate log files after 100MB  
log_statement = 'all'           # Log all SQL statements
```

Basic Logging Settings

```
log_duration = on          # Log the duration of each completed statement  
  
log_line_prefix = '%m [%p] %q%u@%d '  
  
# Log line prefix with timestamp, PID, user, and database  
  
log_error_verbosity = default    # Default verbosity for error messages  
  
log_min_duration_statement = 500  
  
# Log statements that take longer than 500ms
```

Additional Logging Options

Auditing User Activities:

log_connections = on # Log each successful connection

log_disconnections = on # Log each disconnection

log_lock_waits = on

Log queries that wait for more than the specified amount of time

Troubleshooting in PostgreSQL



Effective troubleshooting involves



identifying the root cause of issues and



resolving them to ensure



minimal disruption



to healthcare operations.

Common Troubleshooting Scenarios

Identifying and Resolving Deadlocks:

Deadlocks can occur

when two or more transactions

hold locks that the other transactions need.

Detecting Deadlocks



Enable deadlock logging in postgresql.conf:



`log_lock_waits = on`



`deadlock_timeout = 1s`

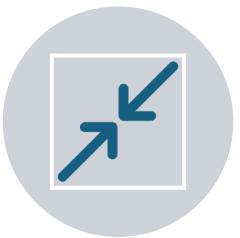
Resolving Deadlocks



ANALYZE THE QUERIES



CAUSING THE DEADLOCK AND



OPTIMIZE THEM TO AVOID CONFLICTS.



ENSURE TRANSACTIONS



ACQUIRE LOCKS IN A CONSISTENT ORDER.

2. Addressing Slow Queries



Slow Queries



Can significantly
impact



Performance of



Healthcare
applications.

Example Query to Identify Slow Queries

```
SELECT query, state, wait_event_type, wait_event,  
now() - query_start AS duration  
FROM pg_stat_activity  
WHERE state <> 'idle' AND  
now() - query_start > interval '1 minute'  
ORDER BY duration DESC;
```

Optimizing Slow Queries

1

Use indexes to speed up query execution.

2

Rewrite inefficient joins or subqueries.

3

Analyze and update statistics using ANALYZE.

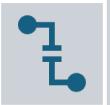
3. Handling Connection Issues



Connection issues can prevent users



from accessing critical healthcare data.



Example Log Analysis for Connection Issues:



```
grep 'connection' /var/lib/pgsql/pg_log/postgresql-*.log
```

Resolving Connection Issues



Check for network problems or firewall issues.



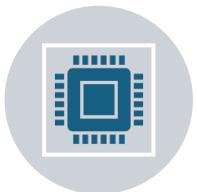
Increase the maximum number of connections



if the limit is reached.



Ensure the database server



has sufficient resources (CPU, memory).

4. Monitoring Database Performance

Regular monitoring helps

identify and address performance issues

before they affect users.

pg_stat_statements to Monitor Performance

```
SELECT query, calls, total_time, mean_time, rows,  
shared_blks_hit, shared_blks_read  
FROM pg_stat_statements  
ORDER BY total_time DESC  
LIMIT 10;
```



**Thank you for
your support and
patience**

Surendra Panpaliya
Founder and CEO
GKTCS Innovations

<https://www.gktcs.com>